


Deep Neural network

Last time

- Difference between linear regression and logistic regression
- How to optimize logistic regression

Preliminaries: Digital image representation



=


11	14	45	36	26	13	14	24	66	...
13	46	36	25	24	23	32	23	52	52
21	64	80	82	104	33	101	140	33	101
45	68	77	107	111	120	187	100	120	187
13	55	101	140	121	33	101	140	50	41
32	86	33	112	140	120	187	100	104	100
23	85	120	187	100	34	77	107	111	116
...	86	33	101	140	33	101	140	121	90
...	...	120	187	100	120	187	100	140	10

Picture element: pixel

Typically 8-bits per channel [0-255] (UINT)

Preliminaries: Convolution / image filtering

Linear filtering




$$w(x, y) \star f(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x - s, y - t)$$

$w(-1, -1)$	$w(-1, 0)$	$w(-1, 1)$
$w(0, -1)$	$w(0, 0)$	$w(0, 1)$
$w(1, -1)$	$w(1, 0)$	$w(1, 1)$

Filter kernel

Preliminaries: Convolution

Linear filtering



How to deal with pixels at the border?

Preliminaries: Convolution

Flip mask w.r.t. signal

Image f

0	0	0	0	0
0	0	0	0	0
0	0	1	0	0
0	0	0	0	0
0	0	0	0	0

Kernel w

1	2	3
4	5	6
7	8	9

$$w(x, y) \star f(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x-s, y-t)$$

Zero padded image

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0




Zero padded image

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Cropped result

0	0	0	0	0
0	1	2	3	0
0	4	5	6	0
0	7	8	9	0
0	0	0	0	0

Preliminaries: Convolution



try your self


```
import cv2
import numpy as np

# Create a dummy input image.
canvas = np.zeros((100, 100), dtype=np.uint8)
canvas = cv2.circle(canvas, (50, 50), 20, (255,), -1)


kernel = np.array([[-1, -1, -1],
                   [-1, 4, -1],
                   [-1, -1, -1]])

dst = cv2.filter2D(canvas, -1, kernel)
cv2.imwrite("./filtered.png", dst)
```





Preliminaries: Convolution




Low-pass filter




Preliminaries: Convolution






Preliminaries: Convolution






High-pass filter






Preliminaries: Convolution



Preliminaries: Convolution






Preliminaries: Convolution






Preliminaries: Convolution

Single filters to find specific color changes



Preliminaries: Convolution

Single filters to find specific color changes



Preliminaries: Convolution







Image features Co-occurrence matrix

- Given a grey-level image I , co-occurrence matrix computes how often pairs of pixels with a specific value and offset occur in the image.
- The offset, $(\Delta x, \Delta y)$, is a position operator that can be applied to any pixel in the image (ignoring edge effects): for instance, $(1,2)$ could indicate "one down, two right".
- An image with p different pixel values will produce a $p \times p$ co-occurrence matrix, for the given offset.
- The $(i,j)^{\text{th}}$ value of the co-occurrence matrix gives the number of times in the image that the i^{th} and j^{th} pixel values occur in the relation given by the offset.

$$C_{\Delta x, \Delta y}(i, j) = \sum_{x=1}^n \sum_{y=1}^m \begin{cases} 1, & \text{if } I(x, y) = i \text{ and } I(x + \Delta x, y + \Delta y) = j \\ 0, & \text{otherwise} \end{cases}$$

Co-occurrence matrix



Statistical measures

Homogeneity	$\sum_i \sum_j \frac{P(i,j)}{1+ i-j }$
Contrast	$\sum_i \sum_j (i-j)^2 P(i,j)$
Energy	$\sum_i \sum_j P(i,j)^2$
Dissimilarity	$\sum_i \sum_j P(i,j) i-j $
Entropy	$-\sum_i \sum_j P(i,j) \log(P(i,j) + \varepsilon)$
Correlation	$\sum_i \sum_j \frac{(i-\mu_x)(i-\mu_y)P(i,j)}{\sigma_x \sigma_y}$

$$C_{\Delta x, \Delta y}(i, j) = \sum_{x=1}^n \sum_{y=1}^m \begin{cases} 1, & \text{if } I(x, y) = i \text{ and } I(x + \Delta x, y + \Delta y) = j \\ 0, & \text{otherwise} \end{cases}$$

Example


https://scikit-image.org/docs/stable/auto_examples/features_detection/plot_glcmm.html

Preliminaries: Image features

- Generally not optimal to work on the raw data
 - Very sensitive to changes in viewpoint, illumination, scaling, rotation, etc.
 - Super high dimensional: curse of dimensionality
- Better idea : use a low dimensional mapping of the original data
 - Summarize the image into a set of descriptive features (lines, corners, colors, texture, ...)
 - Enables training relatively simple and robust classification models
- Concept also used in neural networks
 - Use an encoding scheme to obtain a representation in a latent (feature) space
 - Similar to image compression!

Preliminaries: Classification

- Linear classification example: separate lemons from oranges






Color:
orange
Shape:
sphere
Diameter:
Diameter: ± 8 cm
Weight:
 ± 0.1 kg



Color:
yellow
Shape:
elipsoid
Diameter:
Diameter: ± 8 cm
Weight:
 ± 0.1 kg

→ Use “color” and “shape” as features



Pre-deep learning era




Feature representation

Compute
hand-crafted
image features

$$\begin{bmatrix} 0.24 \\ 1 \\ - \\ 2 \\ - \\ 0 \end{bmatrix} \begin{bmatrix} 3.32 \\ 2.23 \\ 1.21 \\ \vdots \\ -2.12 \\ 0.32 \\ 0.21 \\ \ddots \end{bmatrix}$$


Train machine
learning model




How to get such a model?

Neural Network

The “one learning algorithm” hypothesis



The “one learning algorithm” hypothesis




Sensor representations in the brain



Seeing with your tongue



Human echolocation (sonar)



Haptic belt: Direction sense




Implanting a 3rd eye

[BrainPort; Welsh & Blasch, 1997; Nagel et al., 2005; Constantine-Paton & Law, 2009]

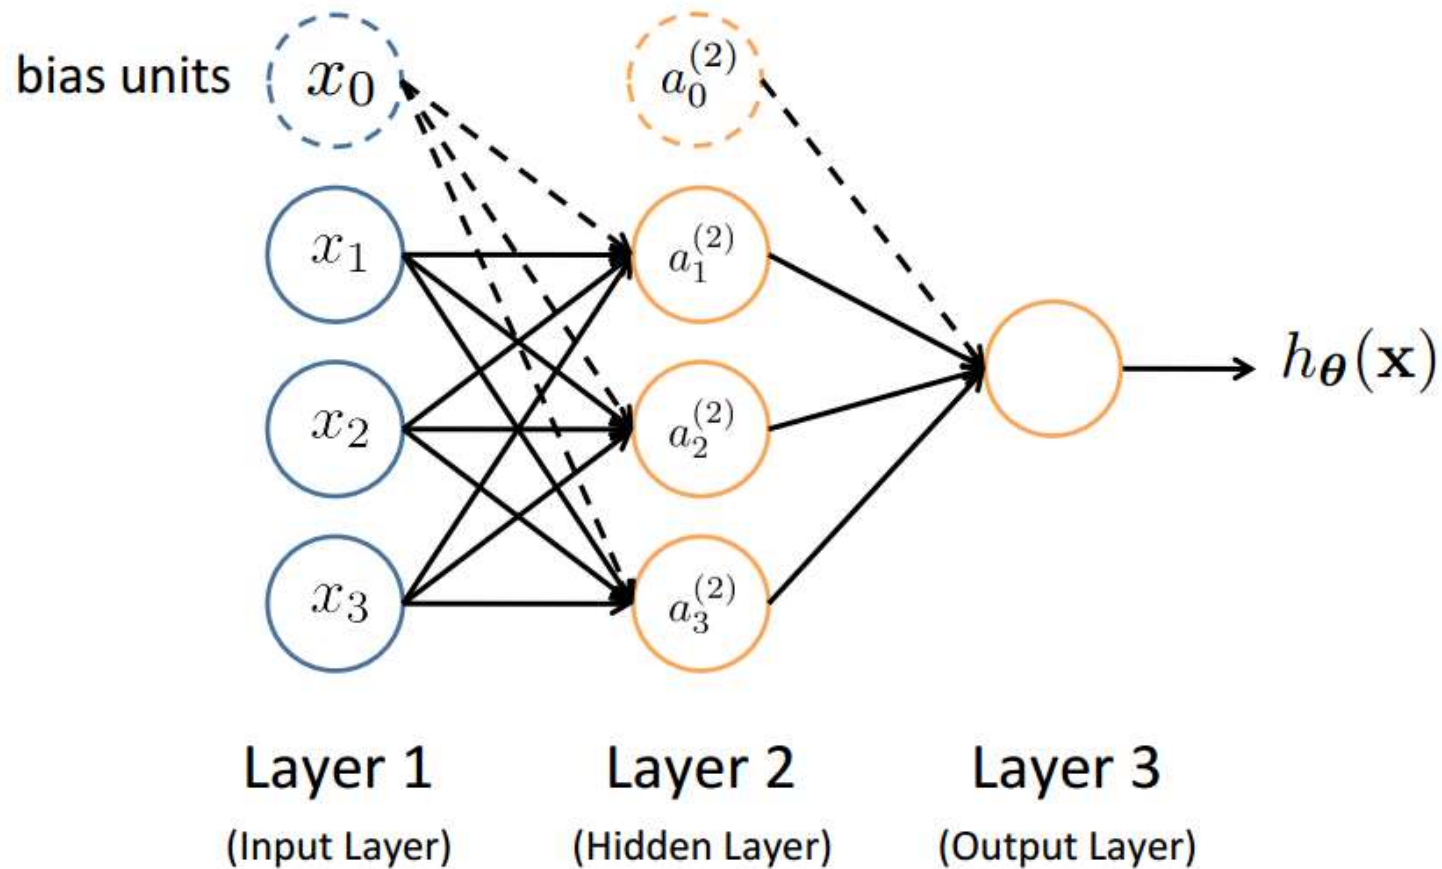
Slide credit: Andrew Ng

Neural Network




Sigmoid (logistic) activation function: $g(z) = \frac{1}{1 + e^{-z}}$

Neural Network (feed forward)



Feed-Forward Process

- Input layer units are features
- Working forward through the network, the **input function** is applied to compute the input value
 - E.g., weighted sum of the input
- The **activation function** transforms this input function into a final value
 - Typically a **nonlinear** function (e.g, **sigmoid**)



$a_i^{(j)}$ = “activation” of unit i in layer j
 $\Theta^{(j)}$ = weight matrix controlling function mapping from layer j to layer $j + 1$

$$a_1^{(2)} = g(\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3)$$

$$a_2^{(2)} = g(\Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3)$$

$$a_3^{(2)} = g(\Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3)$$

$$h_{\theta}(x) = a_1^{(3)} = g(\Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)})$$

If network has s_j units in layer j *and* s_{j+1} units in layer $j+1$,
 then $\Theta^{(j)}$ has dimension $s_{j+1} \times (s_j + 1)$.

$$\Theta^{(1)} \in \mathbb{R}^{3 \times 4} \quad \Theta^{(2)} \in \mathbb{R}^{1 \times 4}$$


Vector Representation

$$a_1^{(2)} = g \left(\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3 \right) = g \left(z_1^{(2)} \right)$$

$$a_2^{(2)} = g \left(\Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3 \right) = g \left(z_2^{(2)} \right)$$

$$a_3^{(2)} = g \left(\Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3 \right) = g \left(z_3^{(2)} \right)$$

$$h_{\Theta}(\mathbf{x}) = g \left(\Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)} \right) = g \left(z_1^{(3)} \right)$$



Feed-Forward Steps:

$$\mathbf{z}^{(2)} = \Theta^{(1)} \mathbf{x}$$

$$\mathbf{a}^{(2)} = g(\mathbf{z}^{(2)})$$

$$\text{Add } a_0^{(2)} = 1$$

$$\mathbf{z}^{(3)} = \Theta^{(2)} \mathbf{a}^{(2)}$$

$$h_{\Theta}(\mathbf{x}) = \mathbf{a}^{(3)} = g(\mathbf{z}^{(3)})$$

Can extend to multi-class



Pedestrian




Car



Motorcycle



Truck



$$h_{\Theta}(\mathbf{x}) \in \mathbb{R}^K$$

We want:

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

when pedestrian

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

when car

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

when motorcycle

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$


when truck

Why staged predictions?


Simple example: AND

$$x_1, x_2 \in \{0, 1\}$$

$$y = x_1 \text{ AND } x_2$$







$$h_{\theta}(\mathbf{x}) = g(-30 + 20x_1 + 20x_2)$$

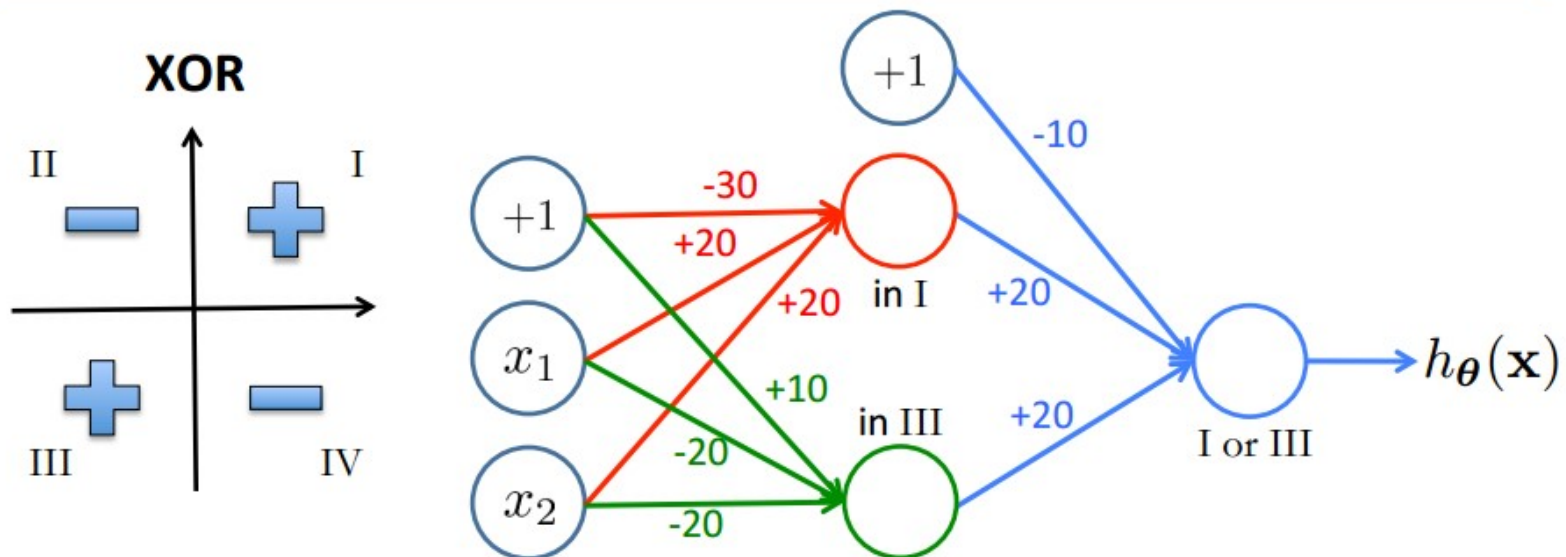



x_1	x_2	$h_{\theta}(\mathbf{x})$
0	0	$g(-30) \approx 0$
0	1	$g(-10) \approx 0$
1	0	$g(-10) \approx 0$
1	1	$g(10) \approx 1$


Representing Boolean Functions



Combining Representations to Create Non-Linear Functions




Layering Representations



20×20 pixel images

$d = 400$ 10 classes



$x_1 \dots x_{20}$

$x_{21} \dots x_{40}$

$x_{41} \dots x_{60}$



•
•
•

$x_{381} \dots x_{400}$

Each image is “unrolled” into a vector x of pixel intensities

Layering Representations

7	9	6	5	8	7	4	4	1	0
0	7	3	3	2	4	8	4	5	7
6	6	3	2	9	2	3	3	2	6
1	3	7	1	5	6	5	2	4	4
7	0	9	0	7	5	8	9	5	4
4	6	6	5	0	2	1	3	6	9
8	5	1	8	9	7	8	7	3	6
1	0	2	8	2	3	0	5	1	5
6	7	8	2	5	3	9	7	0	0
7	9	3	9	8	5	7	2	9	8



Visualization of Hidden Layer

Stochastic Sub-gradient Descent

Given a training set $\mathcal{D} = \{(\mathbf{x}, y)\}$

Initialize $\mathbf{w} \leftarrow \mathbf{0} \in \mathbb{R}^n$

For epoch $1 \dots T$:

For (\mathbf{x}, y) in \mathcal{D} :

Update $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla f(\theta)$

• Return θ

Recap: Logistic regression

$$\min_{\boldsymbol{\theta}} \quad \frac{\lambda}{2n} \boldsymbol{\theta}^T \boldsymbol{\theta} + \frac{1}{n} \sum_i \log (1 + e^{-y_i(\boldsymbol{\theta}^T \mathbf{x}_i)})$$

Let $h_{\boldsymbol{\theta}}(x_i) = 1/(1 + e^{-\boldsymbol{\theta}^T x_i})$ (probability $y = 1$ given x_i)

$$\frac{\lambda}{2n} \boldsymbol{\theta}^T \boldsymbol{\theta} + \frac{1}{n} \sum_i y_i \log (h_{\boldsymbol{\theta}}(x_i)) + (1 - y_i) (\log (1 - h_{\boldsymbol{\theta}}(x_i)))$$

Cost Function

$$f(\theta) = J(\theta) + g(\theta), \quad g(\theta) = \gamma \theta^T \theta$$

Logistic Regression:

$$J(\theta) = -\frac{1}{n} \sum_{i=1}^n [y_i \log h_{\theta}(\mathbf{x}_i) + (1 - y_i) \log (1 - h_{\theta}(\mathbf{x}_i))] + \frac{\lambda}{2n} \sum_{j=1}^d \theta_j^2$$

Neural Network:

$$h_{\Theta} \in \mathbb{R}^K \quad (h_{\Theta}(\mathbf{x}))_i = i^{th} \text{ output}$$

$$J(\Theta) = -\frac{1}{n} \left[\sum_{i=1}^n \sum_{k=1}^K y_{ik} \log (h_{\Theta}(\mathbf{x}_i))_k + (1 - y_{ik}) \log (1 - (h_{\Theta}(\mathbf{x}_i))_k) \right] \\ + \frac{\lambda}{2n} \sum_{l=1}^{L-1} \sum_{i=1}^{s_{l-1}} \sum_{j=1}^{s_l} (\Theta_{ji}^{(l)})^2$$

k^{th} class:	true, predicted
not k^{th} class:	true, predicted

Optimizing the Neural Network

$$J(\Theta) = -\frac{1}{n} \left[\sum_{i=1}^n \sum_{k=1}^K y_{ik} \log(h_{\Theta}(\mathbf{x}_i))_k + (1 - y_{ik}) \log(1 - (h_{\Theta}(\mathbf{x}_i))_k) \right] \\ + \frac{\lambda}{2n} \sum_{l=1}^{L-1} \sum_{i=1}^{s_{l-1}} \sum_{j=1}^{s_l} \left(\Theta_{ji}^{(l)} \right)^2$$

Solve via: $\min_{\Theta} J(\Theta)$

$J(\Theta)$ is not convex, so GD on a neural net yields a local optimum

- But, tends to work well in practice

Need code to compute:


- $J(\Theta)$
- $\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$

Forward Propagation

- Given one labeled training instance (\mathbf{x}, y) :

Forward Propagation

- $\mathbf{a}^{(1)} = \mathbf{x}$
- $\mathbf{z}^{(2)} = \Theta^{(1)}\mathbf{a}^{(1)}$
- $\mathbf{a}^{(2)} = g(\mathbf{z}^{(2)})$ [add $a_0^{(2)}$]
- $\mathbf{z}^{(3)} = \Theta^{(2)}\mathbf{a}^{(2)}$
- $\mathbf{a}^{(3)} = g(\mathbf{z}^{(3)})$ [add $a_0^{(3)}$]
- $\mathbf{z}^{(4)} = \Theta^{(3)}\mathbf{a}^{(3)}$
- $\mathbf{a}^{(4)} = h_{\Theta}(\mathbf{x}) = g(\mathbf{z}^{(4)})$



Online examples

- https://www.w3schools.com/ai/ai_perceptrons.asp


Next time

- Can we make a model for image classification?
- Can we measure the quality of a certain model?
- How can we improve this by learning from data?

- Middle test: 18th Oct 2022

- Final project

Backpropagation: Compute Gradient



$$\frac{d}{dt}f(g(t)) = f'(g(t))g'(t) = \frac{df}{dg} \cdot \frac{dg}{dt}$$

$\delta_j^{(l)}$ = “error” of node j in layer l

Formally, $\delta_j^{(l)} = \frac{\partial}{\partial z_j^{(l)}} \text{cost}(\mathbf{x}_i)$

where $\text{cost}(\mathbf{x}_i) = y_i \log h_{\Theta}(\mathbf{x}_i) + (1 - y_i) \log(1 - h_{\Theta}(\mathbf{x}_i))$