

Recurrent neural networks

Last time

- Transfer learning
- Weight initialization
- Dropout
- Prevent over-training

This time

- Sequential modeling
- How can we model sequential relations using deep neural networks?
- Recurrent Neural Networks (RNNs)
- How can we optimize gradient flow over the sequential dimension ?
- Long Short Term Memory (LSTM)

Till now


- CNNs to do
 - Image classification; Object detection/segmentation;
 - Regression
- Mostly fixed size input → fixed size output
- How would we do
 - Image captioning [one image to variable number of words]
 - Sentiment classification [variable number of words to one sentiment]
 - Translation [variable number of words to variable number of words]
- Recurrent Neural Networks (RNNs)

Examples of sequence data

- Speech recognition
- Music generation
- Sentiment classification
- DNA sequence analysis
- Machine translation
- Video activity recognition

Recurrent neural networks

one to one



one to many






Image
captioning

many to one




Text sentiment
classification

many to many




Machine translation
many to many
Video

many to many



Video frame
classification

Recurrent neural networks



- Recurrence relation
- For a certain time t
- Straightforward option:
- Two linear relations:
 - old state \rightarrow new state
 - Input \rightarrow new state
- Linear output relation:
 - New state \rightarrow output

Function with
parameter Φ

Input

$$h_t = f_w(h_{t-1}, x_t)$$

New state

Old state

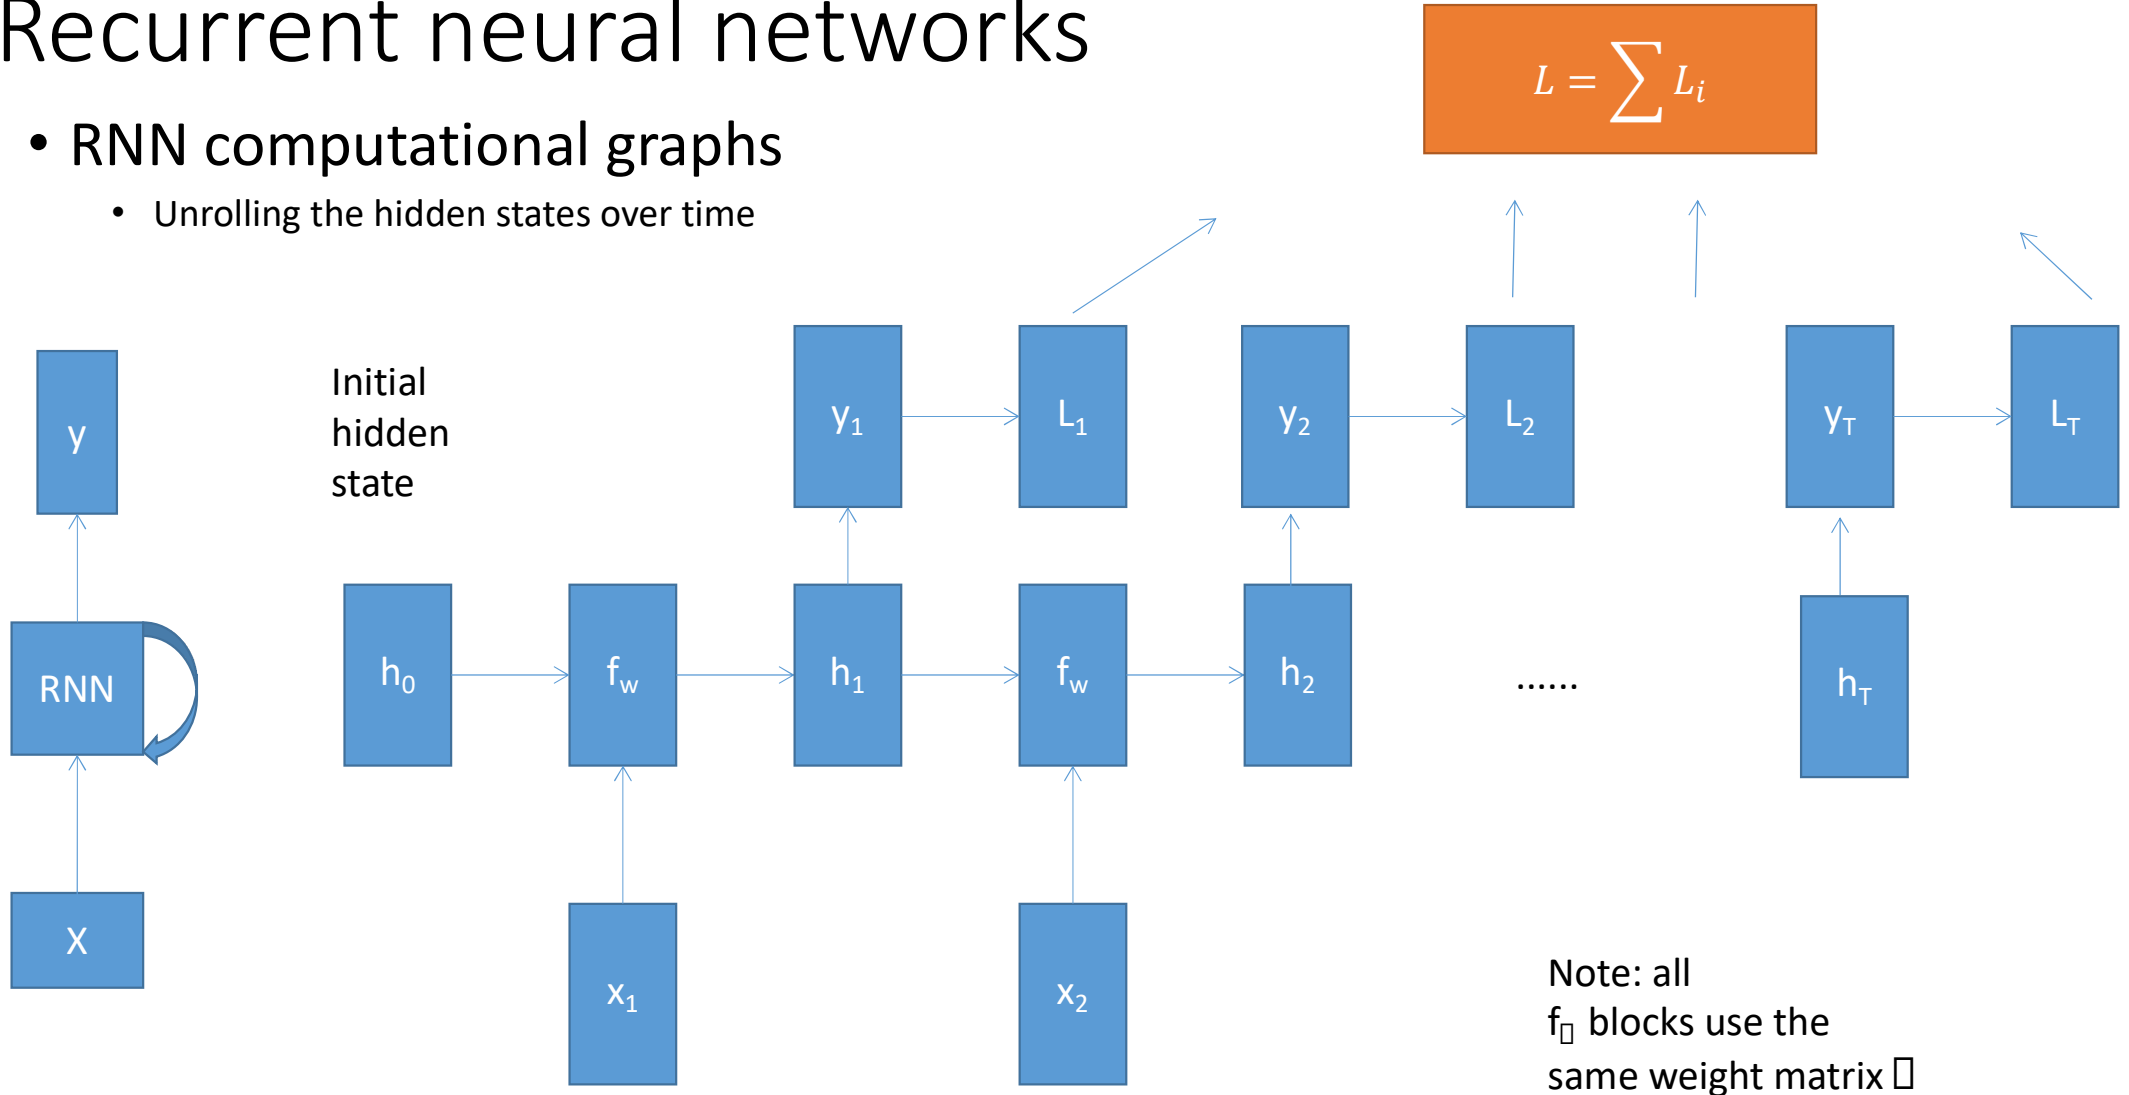
$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t + b_y$$

Note:
parameters W
fixed over time


Recurrent neural networks

- RNN computational graphs
 - Unrolling the hidden states over time




Recurrent neural networks

RNN computational graphs




Many
to
many



Use single input to initialize
the hidden state of the model

One
to
many

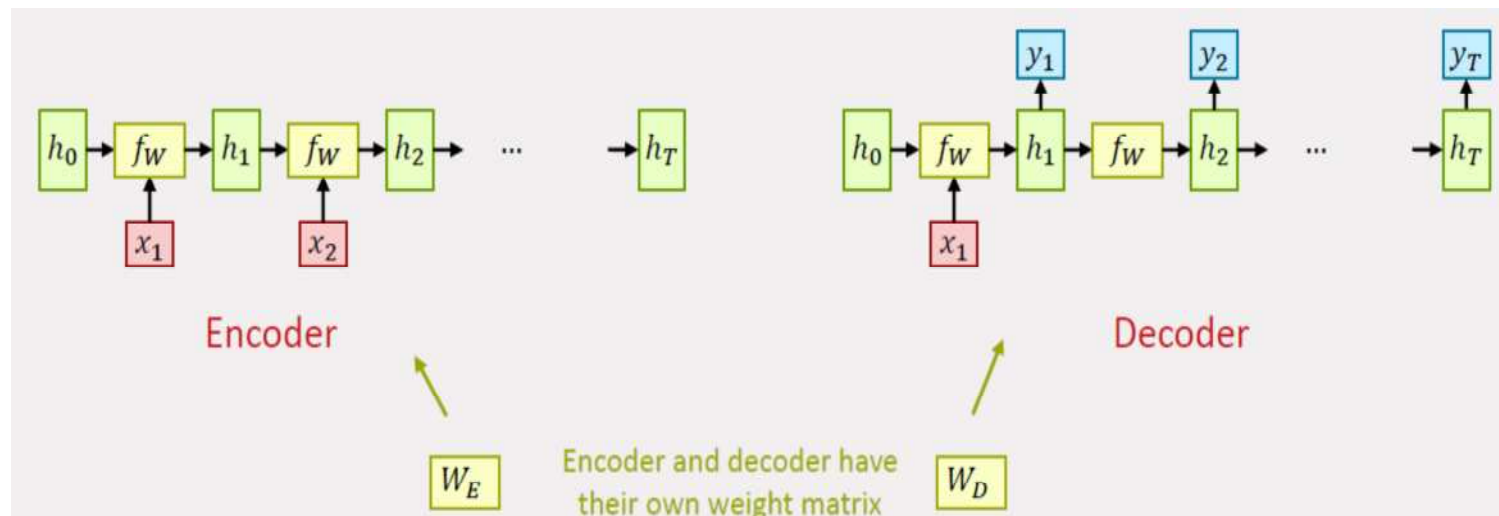
Many
to
one



Final hidden state
summarized all of the context
of the entire sequence

Recurrent neural networks

- RNN computational graphs: Sequence to sequence
- Many to one followed by one to many



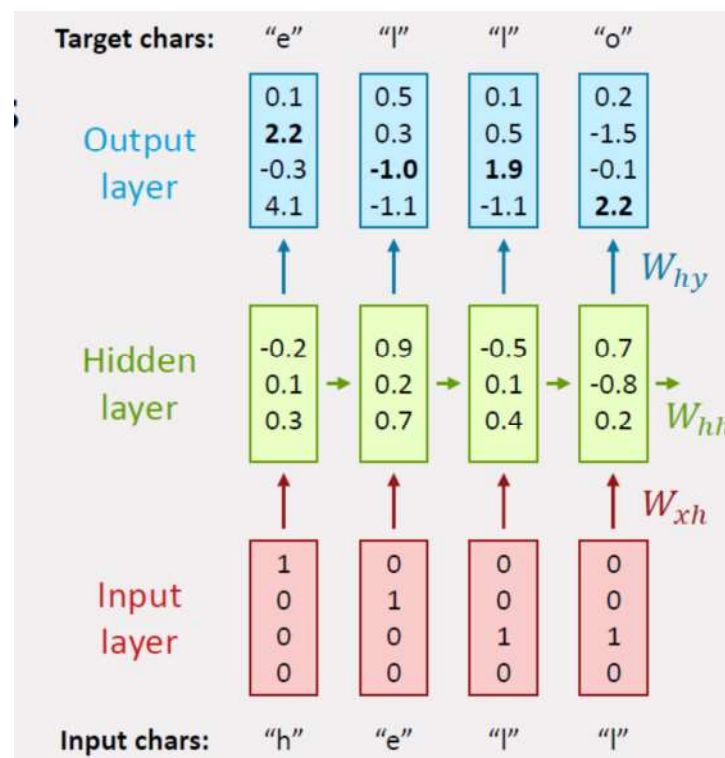
Recurrent neural networks

- RNN computational graphs: Softmax


- Sequence to sequence
- Many to one followed by one to many

- Example

- character level language model
- Predict next letter
- Vocabulary: [h, e, l, o]
- Training sequence: "hello"



Recurrent neural networks- applications



https://www.researchgate.net/figure/An-illustration-of-the-RNN-based-neural-network-model-for-Chinese-to-English-machine_fig1_306093825

Recurrent neural networks- applications

- This License refers to version of the GNU General Public License.
Copyright also means copyright-bick,
- Remade me any thing to his sword
- To his salt and most hidden loose to be so for sings, but not in a libutt of his matter than that shall be sure as will be soldye
- As master compary, do not live in traitor.
- Bless thy five wits!
- Using AI to mix Shakespearean and Modern English

<https://www.knime.com/blog/use-deep-learning-to-write-like-shakespeare>

What you can do with RNNs









<p>A young boy is playing basketball.</p>  A young boy in a white jersey and red shorts is dribbling a basketball on a wooden court. Other players are visible in the background.	<p>Two dogs play in the grass.</p>  Two dogs, one black and white and one brown, are playing in a green grassy field under a blue sky.	<p>A dog swims in the water.</p>  A dog is swimming in a body of water, creating a large splash.
<p>A group of people walking down a street.</p>  A group of people are walking down a street lined with shops and buildings.	<p>A group of women dressed in formal attire.</p>  A group of women are standing together, dressed in formal evening wear.	<p>Two children play in the water.</p>  Two children are playing in a shallow pool of water, splashing around.
<p>A skier is skiing down a snowy hill.</p>  A skier is descending a snowy slope, with trees and a clear sky in the background.	<p>A little girl in a pink shirt is swinging.</p>  A young girl in a pink shirt is swinging happily on a swing set.	<p>A dog jumps over a hurdle.</p>  A dog is jumping over a hurdle on a grassy field.

Image
captioning

What you can do with RNNs

- Produce code

```
static void do_command(struct seq_file *m, void *v)
{
    int column = 32 << (cmd[2] & 0x80);
    if (state)
        cmd = (int)(int_state ^ (in_8(&ch->ch_flags) & Cmd) ? 2 : 1);
    else
        seq = 1;
    for (i = 0; i < 16; i++) {
        if (k & (1 << 1))
            pipe = (in_use & UMXTHREAD_UNCCA) +
                ((count & 0x00000000ffffffff) & 0x000000f) << 8;
        if (count == 0)
            sub(pid, ppc_md.kexec_handle, 0x20000000);
        pipe_set_bytes(i, 0);
    }
    /* Free our user pages pointer to place camera if all dash */
    subsystem_info = &of_changes[PAGE_SIZE];
    rek_controls(offset, idx, &soffset);
    /* Now we want to deliberately put it to device */
    control_check_polarity(&context, val, 0);
    for (i = 0; i < COUNTER; i++)
        seq_puts(s, "policy ");
}
```

RNN practice

Follow use case:


<https://www.simplilearn.com/tutorials/deep-learning-tutorial/rnn>

RNN gradient flow

- What happens during RNN training?

$$\begin{aligned} h_t &= \tanh(W_{hh}h_{t-1} + W_{xh}x_t) \\ &= \tanh\left((W_{hh} \quad W_{xh})\begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}\right) \\ &= \tanh\left(W\begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}\right) \end{aligned}$$

During backwards pass
need to multiply by $\frac{1}{\sigma}$



RNN gradient flow

$$\frac{\partial E_3}{\partial W} = \sum_{k=0}^3 \frac{\partial E_3}{\partial \hat{y}_3} \frac{\partial \hat{y}_3}{\partial s_3} \frac{\partial s_3}{\partial s_k} \frac{\partial s_k}{\partial W}$$

the gradient of the error with respect to the input weight W can be expressed as

$$\frac{\partial s_3}{\partial s_1} = \frac{\partial s_3}{\partial s_2} \frac{\partial s_2}{\partial s_1}$$


We note that this expression includes the derivative of the current hidden state s_3 with respect to each of the k other hidden states — which in turn must be evaluated via the chain rule

$$\frac{\partial E_3}{\partial W} = \sum_{k=0}^3 \frac{\partial E_3}{\partial \hat{y}_3} \frac{\partial \hat{y}_3}{\partial s_3} \left(\prod_{j=k+1}^3 \frac{\partial s_j}{\partial s_{j-1}} \right) \frac{\partial s_k}{\partial W}$$

Long Short Term Memory (LSTM)

- Introduction of a “cell state”

i	Input gate
f	Forget gate
o	Output gate
g	Gain gate



Cell state $c_t = f \odot c_{t-1} + i \odot g$
Hidden state $h_t = o \odot \tanh(c_t)$

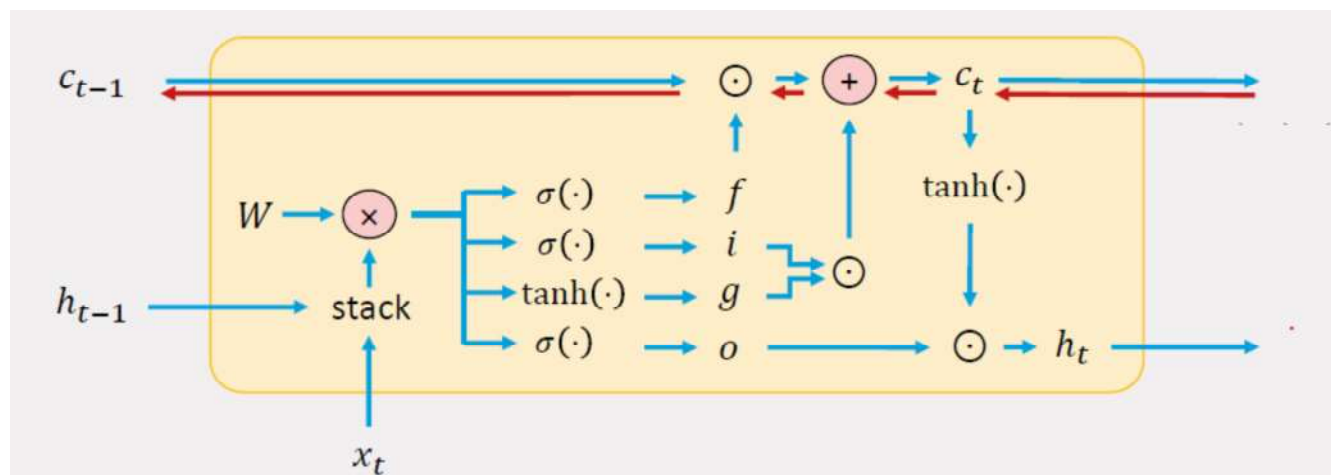
$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma(\cdot) \\ \sigma(\cdot) \\ \sigma(\cdot) \\ \tanh(\cdot) \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

Long Short Term Memory (LSTM)

$$c_t = f \odot c_{t-1} + i \odot g$$

$$h_t = o \odot \tanh(c_t)$$


$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma(\cdot) \\ \sigma(\cdot) \\ \sigma(\cdot) \\ \tanh(\cdot) \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$



During backprop
from c_t to c_{t-1} :

- only elementwise multiplication by $\frac{\partial c_t}{\partial c_{t-1}}$
- no matrix multiplication by $\frac{\partial c_t}{\partial c_{t-1}}$


Calculating New Values to Add to Cell State



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$


$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Updating Cell State



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

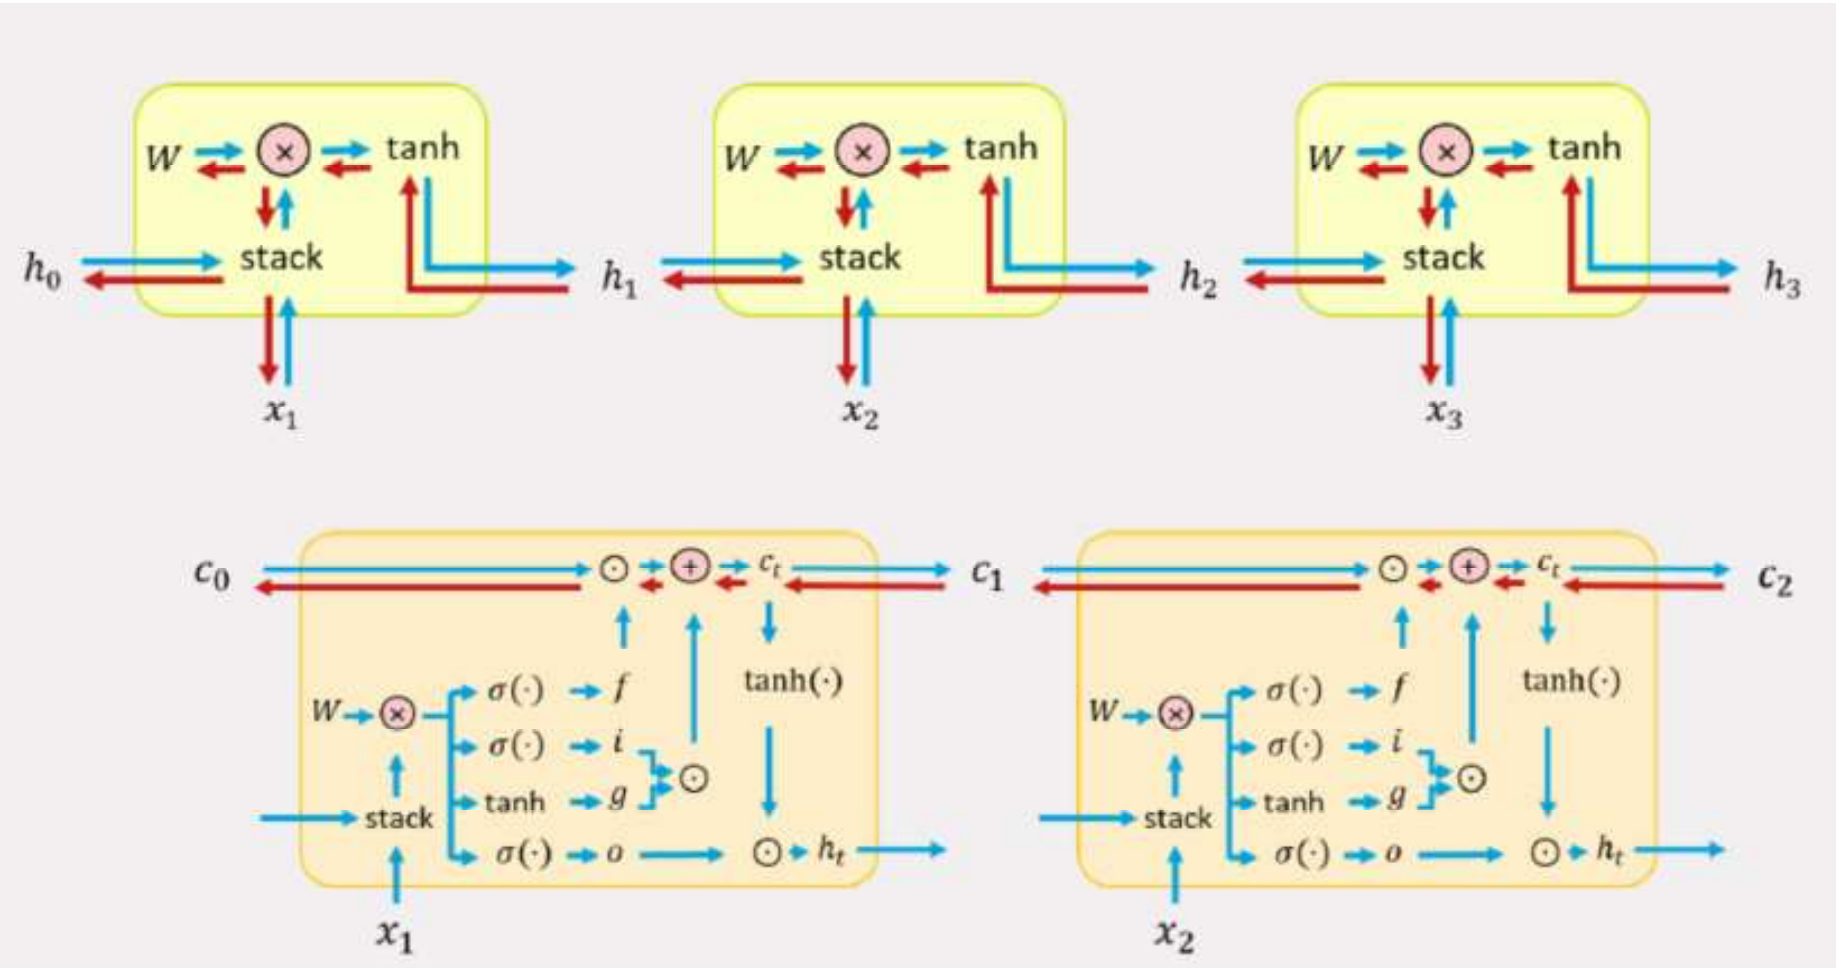
Calculating Hidden State from Cell State



$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

Vanilla RNN



LSTM

Temporal modeling summary

- Recurrent Neural Networks (RNNs)
 - Add recurrent relation in the network to model sequential relations
 - Unroll the network over time and apply standard backpropagation to train
 - In practice: vanilla RNNs don't work that well due to vanishing/exploding gradients
- Long Short Term Memory (LSTM)
 - Introduce an additional cell state, governing the input and output of the hidden state
 - Input --, output --, forget --, "gating" gates
 - Uninterrupted gradient flow

Case study

- LSTM case study
 - See Canvas

- Cifar case study

<https://ermlab.com/en/blog/nlp/cifar-10-classification-using-keras-tutorial/>