

15 Binary Search Trees


Instructor : Ke Wei (柯韋)

➡ A319 ☎ Ext. 6452 ✉ wke@ipm.edu.mo

<http://brouwer.ipm.edu.mo/COMP122/19/>

Bachelor of Science in Computing, School of Public Administration, Macao Polytechnic Institute

March 18, 2019



Outline

1 Associative Arrays

2 Binary Search Trees

- Searching
- Insertion
- In-order Traversals
- Deletion

3 Balance of Binary Search Trees

Associative Arrays

- Given an index i of an array-based list a , $a[i]$ is the value stored at location i .
- If we abstract the location away, an array associates a value ($a[i]$) with each integer i in the range from 0 to $\text{len}(a)$. For example,

$a = [\text{'John'}, \text{'Mary'}, \text{'Mary'}, \text{'Susan'}]$

can be regarded as a set of associations:

$\{0 \mapsto \text{'John'}, 1 \mapsto \text{'Mary'}, 2 \mapsto \text{'Mary'}, 3 \mapsto \text{'Susan'}\}.$

- The indices of a list are the keys, each of which is unique. If we generalize the integer keys to any type of data, we have an associative array. For example,

$\{\text{'North'} \mapsto \text{'John'}, \text{'South'} \mapsto \text{'Mary'}, \text{'West'} \mapsto \text{'Mary'}, \text{'East'} \mapsto \text{'Susan'}\}.$

- An associative array is a map from keys to values, where each of the keys is unique. An associative array is also called a *map*, a *table*, or a *dictionary*.

Associative Array Operations

The operations on associative arrays varies from system to system, depending on applications. The following are a few common operations.

- `__getitem__(self, key)` — returns the value associated with the *key*. If the *key* is not in the associative array, it raises *KeyError*.

For an associative array *m*, `__getitem__(self, key)` is called by `m[key]`.

- `__setitem__(self, key, value)` — inserts the *key* \mapsto *value* association into the associative array if the *key* is new. If the *key* exists in the associative array, it associates the new *value* with the *key*.

For an associative array *m*, `__setitem__(self, key, value)` is called by `m[key] = value`.



- `__delitem__(self, key)` — removes the *key* and its associated value from the the associative array. If the *key* is not in the associative array, it raises *KeyError*.

For an associative array *m*, `__delitem__(self, key)` is called by `del m[key]`.

- `__iter__(self)` — iterates over all the keys in the associative array.

Binary Search Trees

- A binary tree can hold a collection of $\langle \text{key}, \text{value} \rangle$ pairs, each of them is stored in a node. Each *key* in the tree is *unique*, and it is also the *key* of the node.
- There must be an order `__lt__` (*self*, *other*) ($<$) defined between the keys.
- Such a binary tree is a binary search tree if for every node in the tree,
 - 1 all the keys in its left subtree are less than the key of the node, and
 - 2 all the keys in its right subtree are greater than the key of the node.



Searching

- The main application of binary search trees is to search for a node with a given *key*.
- According to the properties of binary search trees, we can find a node very quickly:
 - 1 If the tree is empty, we return **None**, indicating that the key is not found.
 - 2 If the root node has the *key*, we simply return the root.
 - 3 If the *key* is less than the root key, the node to find must be in the left subtree. We can recursively search for it.
 - 4 If the *key* is greater than the root key, we recursively search for it in the right subtree.
- The searching advances a level in each step, so the number of steps cannot be more than the depth of the tree.

Finding the Node with a Given Key

The tree node is defined as usual, except that we split *itm* into *key* and *value*.

```
1 class Node:
2     def __init__(self, key, value):
3         self.key, self.value = key, value
4         self.left = self.right = None
```

We search for a given key by the *find* function.

```
1 def find(root, key):
2     if root is None or key == root.key:
3         return root
4     elif key < root.key:
5         return find(root.left, key)
6     else:
7         return find(root.right, key)
```

Finding the Node with a Given Key (2)

The tail-recursion on the previous slide can be transformed to a loop.


```
1 def find(root, key):
2     while root is not None and key != root.key:
3         root = root.left if key < root.key else root.right
4     return root
```

The node with the least key is the leftmost node.

```
1 def find_min(root):
2     if root is not None:
3         while root.left:
4             root = root.left
5     return root
```

Inserting a Pair

- Insertion of a pair *key* and *value* can also be recursively performed.
 - 1 If the tree is empty, we make a tree of a single node with the *key* and the *value*.
 - 2 If the root has the *key*, we associate it with the new *value*.
 - 3 If the *key* is less than the root key, we insert the pair to the left subtree.
 - 4 If the *key* is greater than the root key, we insert the pair to the right subtree.
- The new node is always a leaf.




Inserting a Pair — Code

```
1 def insert(root, key, value):
2     if root is None:
3         return Node(key, value)
4     else:
5         if key == root.key:
6             root.value = value
7         elif key < root.key:
8             root.left = insert(root.left, key, value)
9         else:
10            root.right = insert(root.right, key, value)
11    return root
```

Traversals of Binary Search Trees

- For a binary tree, a parent node may be considered between its left and right children. Thus, we can define the *in-order* traversal as follows:


- 1 recursively traverse the left subtree;
- 2 visit the root node;
- 3 recursively traverse the right subtree.



- The in-order traversal of a binary search tree results an ordered sequence.
- We can rebuild a binary search tree from the pre-order traversal sequence.
 - Just keep fetching items from the sequence and inserting them to a tree (initially empty), using the previously described insertion method.
- For a general binary tree, we can rebuild it from the sequences of pre-order and in-order traversals.

Deleting a Node by a Key


- Deletion must keep the ordering property of a binary search tree.



- We must move the nearest left or nearest right node in place of the deleted one.
 - First, find the node p to delete.
 - If one of p 's subtrees is empty, move the other one in place of it.
 - Find the right-most node q of its left subtree (or the other way).
 - Since q has no right subtree, move q 's left subtree in place of q .
 - Move q in place of p .

Deleting a Node by a Key


- Deletion must keep the ordering property of a binary search tree.



- We must move the nearest left or nearest right node in place of the deleted one.
 - First, find the node p to delete.
 - If one of p 's subtrees is empty, move the other one in place of it.
 - Find the right-most node q of its left subtree (or the other way).
 - Since q has no right subtree, move q 's left subtree in place of q .
 - Move q in place of p .

Deleting a Node by a Key


- Deletion must keep the ordering property of a binary search tree.



- We must move the nearest left or nearest right node in place of the deleted one.
 - First, find the node p to delete.
 - If one of p 's subtrees is empty, move the other one in place of it.
 - Find the right-most node q of its left subtree (or the other way).
 - Since q has no right subtree, move q 's left subtree in place of q .
 - Move q in place of p .

Deleting a Node by a Key


- Deletion must keep the ordering property of a binary search tree.



- We must move the nearest left or nearest right node in place of the deleted one.
 - First, find the node p to delete.
 - If one of p 's subtrees is empty, move the other one in place of it.
 - Find the right-most node q of its left subtree (or the other way).
 - Since q has no right subtree, move q 's left subtree in place of q .
 - Move q in place of p .

Deleting a Node by a Key


- Deletion must keep the ordering property of a binary search tree.



- We must move the nearest left or nearest right node in place of the deleted one.
 - First, find the node p to delete.
 - If one of p 's subtrees is empty, move the other one in place of it.
 - Find the right-most node q of its left subtree (or the other way).
 - Since q has no right subtree, move q 's left subtree in place of q .
 - Move q in place of p .

Deleting a Node by a Key


- Deletion must keep the ordering property of a binary search tree.



- We must move the nearest left or nearest right node in place of the deleted one.
 - First, find the node p to delete.
 - If one of p 's subtrees is empty, move the other one in place of it.
 - Find the right-most node q of its left subtree (or the other way).
 - Since q has no right subtree, move q 's left subtree in place of q .
 - Move q in place of p .

Deleting a Node by a Key

- Deletion must keep the ordering property of a binary search tree.



- We must move the nearest left or nearest right node in place of the deleted one.
 - First, find the node p to delete.
 - If one of p 's subtrees is empty, move the other one in place of it.
 - Find the right-most node q of its left subtree (or the other way).
 - Since q has no right subtree, move q 's left subtree in place of q .
 - Move q in place of p .

Deleting the Rightmost Node

- Similar to the deletion of the leftmost leaf in a perfectly balanced tree, except that the rightmost node of a search tree may have a nonempty left subtree.
- The recursion stops when the root is the rightmost node, the root will change to *root.left* after the deletion.
- The *delete_rightmost* function returns a pair — the new root and the deleted rightmost node.

```
1 def delete_rightmost(root):
2     if not root.right:
3         return (root.left, root)
4     else:
5         root.right, rightmost = delete_rightmost(root.right)
6         return (root, rightmost)
```

Deleting a Key — Code


```
1 def delete(root, key):
2     if root is None: return (root, None)
3     elif key == root.key:
4         if not root.left: return (root.right, root)
5         elif not root.right: return (root.left, root)
6         else:
7             sub, rightmost = delete_rightmost(root.left)
8             rightmost.left, rightmost.right = sub, root.right
9             return (rightmost, root)
10    elif key < root.key:
11        root.left, deleted = delete(root.left, key)
12        return (root, deleted)
13    else:
14        root.right, deleted = delete(root.right, key)
15        return (root, deleted)
```

Implementing Associative Arrays Based-on Binary Search Trees

```
1 class BSTAssocArray:
2     def __init__(self):
3         self.root = None
4     def __getitem__(self, key):
5         p = find(self.root, key)
6         if p is None: raise KeyError
7         return p.value
8     def __setitem__(self, key, value):
9         self.root = insert(self.root, key, value)
10    def __delitem__(self, key):
11        self.root, deleted = delete(self.root, key)
12        if deleted is None: raise KeyError
13    def __iter__(self):
14        yield from (p.key for p in inorder_nodes(self.root))
```

Rebalancing

- The best binary search tree of n nodes has the minimum depth of $\log n$. This results the quickest searching.
- The worst has the maximum depth of $n - 1$. In this case, the tree degenerates to a linked list.
- Since the in-order traversal results an ordered sequence, we may convert the tree to such a sequence and rebuild a balanced tree from it.
- To build a balanced tree from an ordered sequence:
 - 1 if the sequence is empty, return an empty tree; otherwise,
 - 2 take the middle item as the root, and
 - 3 recursively build the left and right subtrees from the front and rear halves of the sequence.



Rebalancing — Code

This method takes a list of nodes and connects the nodes ranging from index i to index $j-1$ into a balanced binary tree, whose in-order traversal sequence is the same as the list $a[i:j]$. The method returns the root node of the balanced tree.

```

1 def build_bal(a, i, j):
2     if i < j:
3         m = (i+j)//2
4         root = a[m]
5         root.left = build_bal(a, i, m)
6         root.right = build_bal(a, m+1, j)
7         return root
8     else:
9         return None

```

Rebalancing — Proof

We prove that when $i \leq j$, $build_bal(a, i, j)$ returns a perfectly balanced tree Δ of size $j - i$, and with the in-order traversal sequence equals $a[i:j]$, denoted as $\Delta_{j-i}^{a[i:j]}$. We induct on $j - i$.

- Base case: when $j - i = 0$, $a[i:j] = []$ and $build_bal(a, i, j)$ returns $\Delta_0^[]$.
- Induction step: when $j - i > 0$, we have $j > i$. Thus, $m = \lfloor \frac{i+j}{2} \rfloor$, we have 1) $i + j = 2m$, or 2) $i + j = 2m + 1$. In case 1) we have

$$m - i = \frac{2m - 2i}{2} = \frac{i + j - 2i}{2} = \frac{j - i}{2} \quad \text{and} \quad j - (m + 1) = \frac{2j - 2m}{2} - 1 = \frac{2j - i - j}{2} - 1 = \frac{j - i}{2} - 1,$$

and in case 2) we have

$$m - i = \frac{2m - 2i}{2} = \frac{i + j - 1 - 2i}{2} = \frac{j - i - 1}{2} \quad \text{and} \quad j - (m + 1) = \frac{2j - 2m - 2}{2} = \frac{2j - i - j + 1 - 2}{2} = \frac{j - i - 1}{2}.$$

In either case, $0 \leq$ the sizes $< j - i$ and the size difference is no more than 1. By induction hypothesis, $build_bal(a, i, m)$ returns $\Delta_{m-i}^{a[i:m]}$, and $build_bal(a, m + 1, j)$ returns $\Delta_{j-(m+1)}^{a[m+1:j]}$.

Therefore, $build_bal(a, i, j)$ returns $\Delta_{m-i+1+j-(m+1)}^{a[i:m]+a[m]+a[m+1:j]} = \Delta_{j-i}^{a[i:j]}$.

