

22 Dijkstra's Shortest Path Algorithm


Instructor : Ke Wei (柯韋)

➡ A319 ☎ Ext. 6452 ✉ wke@ipm.edu.mo

<http://brouwer.ipm.edu.mo/COMP122/19/>

Bachelor of Science in Computing, School of Public Administration, Macao Polytechnic Institute

April 24, 2019




Outline

- 1 Edge-Weighted Graphs
- 2 Shortest Paths
- 3 Dijkstra's Algorithm
- 4 Analysis

Edge-Weighted Graphs

- An *edge-weighted* graph is a graph having a weight, or number, associated with each edge.
- Some algorithms require all weights to be non-negative.
- Edge weights may represent distances, costs, delays, etc.

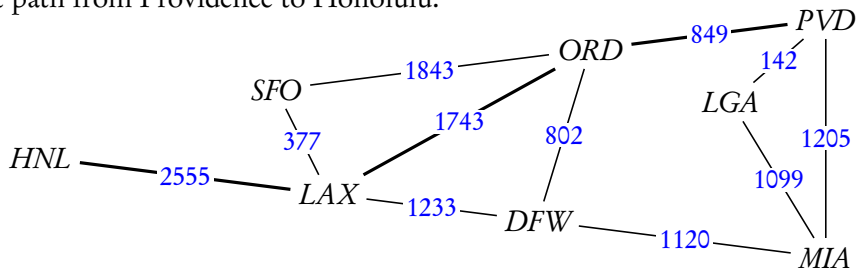
A flight route graph:



Shortest Paths

- Given an edge-weighted graph and two vertices u and v , we want to find a path with the minimum total weight between u and v . This is the *shortest path problem*.
- Some applications:
 - Flight reservations
 - Driving directions
 - Network packet routing


The shortest path from Providence to Honolulu:



Properties of Shortest Paths

- A sub-path of a shortest path is itself a shortest path.
- The shortest paths from a starting vertex s to all the other vertices form a tree rooted at s .

The tree of shortest paths from Providence:




Dijkstra (1930–2002)'s Algorithm

- The distance of a vertex v from a vertex s is the total weight of a path between s and v .
- Dijkstra's algorithm computes the shortest distances of all the vertices from a given starting vertex s .
- Assumptions:
 - The graph is connected. Edges of infinite weight can be introduced to apply the algorithm to a general graph.
 - The edge weights are *non-negative*. A path can not be shortened by appending more edges.

Dijkstra's Algorithm (2)

- We grow a set of “known” vertices,
 - whose shortest distances are already known and can not be shortened further,
 - beginning with s and eventually containing all the vertices.
- We store with each vertex v a field $dist(v)$, called the *distance* of v , representing the shortest distance of v from s in the *subgraph* consisting of
 - the set of “known” vertices, often called the “cloud”, and
 - their adjacent vertices, with only the edges from the “known” vertices (the cloud).
- At each step:
 - we add to the set the vertex u outside the set with the shortest distance field, then
 - we update the distance fields of the vertices adjacent to u , if the fields can be shortened.





Edge Relaxation







- Consider an edge $e = (u, z)$ such that
 - u is the vertex most recently added to the “known” set.
 - z is not in the “known” set.
- The relaxation of the edge e updates $dist(z)$, the distance of z , as follows:

$$dist(z) \leftarrow \min(dist(z), dist(u) + weight(e)).$$


- We also record u as the parent of z in a field $parent(z)$. We can then use the field to track back the path from s to z .



Dijkstra's Algorithm — Illustrated



Dijkstra's Algorithm — Illustrated (2)



Step	Next Vertex	Current Distance (Parent)						
		<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>	<i>g</i>
—	—	0 _(—)	∞ _(—)	∞ _(—)	∞ _(—)	∞ _(—)	∞ _(—)	∞ _(—)
1.	<i>a</i>	✓	2 _(<i>a</i>)		1 _(<i>a</i>)			
2.	<i>d</i>			3 _(<i>d</i>)	✓	3 _(<i>d</i>)	9 _(<i>d</i>)	5 _(<i>d</i>)
3.	<i>b</i>		✓					
4.	<i>e</i>					✓		
5.	<i>c</i>			✓			8 _(<i>c</i>)	
6.	<i>g</i>						6 _(<i>g</i>)	✓
7.	<i>f</i>						✓	

Revisiting the Array-based Heap

- We store the vertices in a heap, and compare vertices based on their distances.
- There's a need to decrease a specified vertex in the heap during a distance update.
- In an array-based heap, an element can be sifted up and down. Therefore, modifying an element is possible.
- We need a vertex to have a place to store its position in the heap.
- When a vertex has been modified, we get its position and notify the heap to sift the vertex at that position up or down.
- When we move a vertex in the heap, we also update its position.

Analysis

- Each vertex is enqueued and dequeued once: $\mathcal{O}(|V| \log |V|)$.
- Each edge is visited once and each visit may call *priorityDecreased* on the destination vertex:

$$\mathcal{O}(|E| \log |V|).$$

- The total running time is $\mathcal{O}((|V| + |E|) \log |V|)$.
- or, since the graph is (weakly) connected ($|E| \geq |V| - 1$), the overall running time is

$$\mathcal{O}(|E| \log |V|).$$

- If the graph is a DAG, we can select the vertices according to their topological order.
- This selection rule works because when a vertex v is selected, its distance, $dist(v)$, can no longer be shortened, since by the topological order it has no incoming edges from unknown vertices.
- Since the selection takes constant time, the running time is $\mathcal{O}(|V| + |E|)$.