



# Notes #1: Computer System Review

---

COMP 213 (211/212)  
Operating Systems

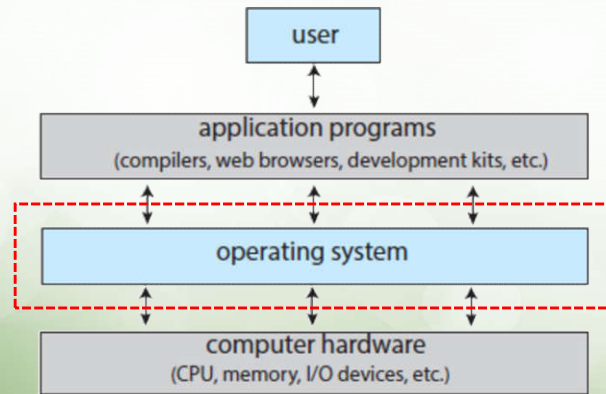
2019-2020 1st Semester

1

## Topics to Cover

- History and development of computing systems
- Evolution of operating systems
- Textbook
  - Stallings's : Chapter 1.1-1.7
- References
  - Tanenbaum's : Chapter 1
  - Silberschatz's : Chapter 1

## What this Course is about



Eddie Law 3

Q: Which Ones Below is/are Not Operating System(s)?




Eddie Law 4


## Short History of Computers

- Dawn of Time (1945 – 55):  
ENIAC (Electronic Numerical Integrator And Computer)
  - Harwell Computer at National Museum of Computing (TNMOC) at Bletchley Park (1949)
- Core Memories (1950s - 60s)
- Multics System: MIT & Bell (1975)
- Disk drives (early 1970s)


ENIAC: weighted 30 tons,  
consumed 200 kilowatts



Harwell: can  
still run today



Magnetic core memory



Multics  
System



Model 3340 hard disk  
1973

1.7

140

Model 3370  
1979

7.7

2,300

7.7 Mbit/sq. in  
2,300 MBytes

1.7 Mbit/sq. in  
140 MBytes



5

## Products and Solutions

- Nowadays, the world is a large parallel system
  - Microprocessors in everything
  - Vast infrastructure behind them

Internet  
Connectivity



Scalable,  
reliable,  
Secure  
services

Databases  
Information Collection  
Remote Storage  
Online Games  
Commerce  
...




MEMS sensor


Eddie Law 6

# Moore's Law

- Gordon Moore (co-founder of Intel) predicted in 1965 that the **transistor density** of semiconductor chips would double roughly every 18 months



- $2 \times$  transistors/chip every 1.5 years - called "**Moore's Law**"
- i.e., microprocessors have become smaller, denser, and more powerful



Eddie Law 7


Joy's Law (1984) - peak computer speed doubled every year

## Death of the Moore's Law


Then how to further improve the system performance in future?

Parallelism, e.g., multi-core, discuss later

### MOORE'S LAW IN DRAMS



### UNIPROCESSOR PERFORMANCE (SINGLE CORE)



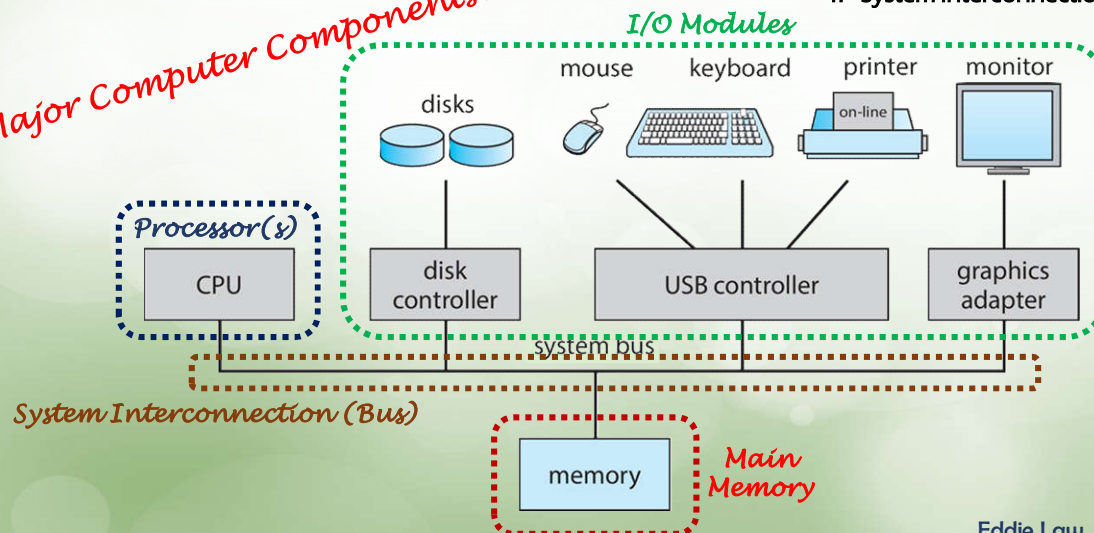
Eddie Law 8

## Components of a Typical PC

*Major Computer Components??*

Typical components:

1. Processor(s)
2. Main memory
3. I/O modules
4. System interconnection (bus)




Eddie Law 9

## Roles of Processor(s)


Controls the operation of a computer

Performs data processing functions

The chip is called  
**Central Processing Unit (CPU)**



*The heart and soul of a computer*



Eddie<sup>19</sup>Law

## Main Memory

- Holds data and instructions temporarily which the CPU will process and execute
- Volatile (RAM: Random Access Memory)
- Contents of the memory is lost when the computer is shut down
- Also referred to as real, or primary memory


RAM in market:

- 3200 MT/s DDR4

- 6000+ MT/s DDR5 (2019)

DDR (double data rate)

MT/s (mega transfers per second)



Eddie Law 11

## I/O Modules

Moves data between computer and external devices, for example:

Storage (e.g., hard drive)

Communications equipment (e.g., wifi)

Terminals (i.e., monitor)

Eddie Law 12




## System Bus

- Provides for communications among processors, main memory, and I/O modules

*Examples of buses in PC:*

1. ISA (industry standard architecture)
2. PCI (peripheral component interconnect)
3. PCIe (PCI express)



Eddie Law 13



## Microprocessor(s)

- Enabled desktop, handheld computing, small board computer
- Processor(s) on a single chip
- Fastest general purpose processor
- Multiprocessors
  - SMP (Symmetric), AMP (Asymmetric)
- Multi-core: a chip (socket) contains multiple processors (cores)
  - “ManyCore” for GPU refers to many processors/chip, e.g., 128 cores
  - Parallelism must be exploited - OS facilitates it

*Examples  
CPU: Intel, AMD,  
ARM, Atmel, etc.  
GPU: ATI, NVidia*

### AMD Ryzen 9 3950X:

16 cores, 32 threads, 3.5-4.7GHz  
L1 64KB/core, L2 512KB/core, L3 64MB  
Thermal design power (TDP) 105W




Eddie Law 14

## Graphical Processing Units (GPUs)

- Arrays of **small computing cores**
- Provide efficient computation on arrays of data
- Good for general numerical processing, such as
  - Physics simulations for games
  - Computations on large spreadsheets

GTX 750 Ti GPU Engine Specs:	
CUDA Cores	640
Base Clock (MHz)	1020
Boost Clock (MHz)	1085
GTX 750 Ti Memory Specs:	
Memory Clock	5.4 Gbps
Standard Memory Config	2048 MB
Memory Interface	GDDR5
Memory Interface Width	128-bit
Memory Bandwidth (GB/sec)	86.4
GTX 750 Ti Support:	
OpenGL	4.4
Bus Support	PCI Express 3.0
Certified for Windows 7, Windows 8, Windows Vista or Windows XP	Yes



RADEON™ RX 460 GRAPHICS	
GEN ARCHITECTURE	4th Generation
COMPUTE UNITS	14 CUs
STREAM PROCESSORS	896
CLOCK SPEEDS (BOOST / BASE)	1200 / 1090 MHz
PEAK PERFORMANCE	Up to 2.2 TFLOPS
MEMORY CLOCK SPEED (MHz)	1750
MEMORY BANDWIDTH	112 GB/s
MEMORY INTERFACE	128 bit
MEMORY TYPE	GDDR5

**CUDA (Compute Unified Device Architecture), the Nvidia's proprietary compute platform on GPU**

**OpenCL (Open Computing Language) On ATI's Radeon GPU**

Eddie Law 15

## Digital Signal Processors (DSPs)

- Used to be embedded in devices like modems
- Effective for **streaming signals** such as audio or video
- **Encoding/decoding** speech and video (codecs)
- Support for **encryption** and **security**

1.25GHz, 8-core, TMS320C6678 DSP





Eddie Law 16



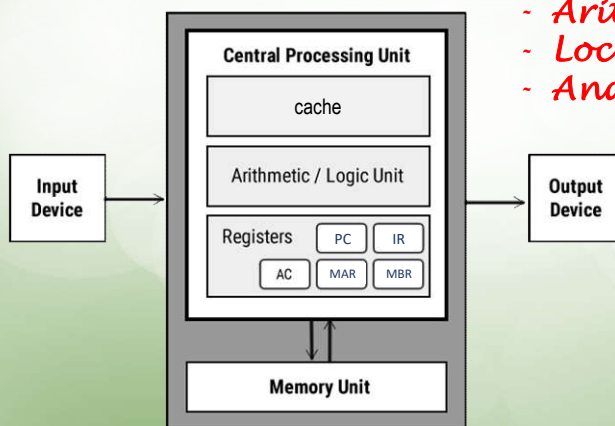
## System-on-Chip (SoC)

- To satisfy the requirements of handheld devices, the microprocessor gives way to the SoC
- Components such as DSPs, GPUs, codecs and main memory, in addition to the CPUs and caches, **are on the same chip**



Eddie Law 17

## PC and the CPU



*What inside a CPU?*

- *Arithmetic and logic unit (ALU)*
- *Local cache memory (L1)*
- *And some registers*

**Some shown registers**

**PC: program counter**

**IR: instruction register**

**AC: accumulator in ALU**

**MAR: memory address register**

**MBR: memory buffer register**

Eddie Law 18

## What is a register?

- Small memory set inside the processor
- Termed as “general registers” and/or “special registers”
- Each **usually holds only single value**, e.g., instruction, data, or address
- Types of registers
  - Data (instruction is a type of data)
  - Address: base and index registers; segment pointer and an offset register; stack pointer

Eddie Law 19

## User-Invisible Registers (1)

- **Program Counter (PC)** - contains address of next instruction to be fetched
- **Instruction Register (IR)** - contains the most recently fetched instruction for next execution
- **Processor Status Word (PSW)**
  - Condition codes (or flags) - bits set by processor hardware as results of operations, e.g., positive result, negative result, divided by zero, overflow
  - Some control bits: Interrupt enable / disable; kernel (supervisor) / user mode

Eddie Law 20

## User-Visible Registers (2)

- Enable programmer to minimize main-memory references by optimizing register use
- Available to all applications and system programmes
- **Data registers**
  - For calculation
  - Minimize references to main memory
- **Address registers**
  - Contain main memory address of data and instructions
  - May contain a portion of an address used to calculate the complete address
- Settings of registers would be vendor-specific and chip-specific

Eddie Law 21

## Control and Status Registers

- Used by processor to control the operations of a processor
- Used by privileged operating-system routines to control the execution of programmes
- E.g., the PSW

Eddie Law 22

## An Instruction Cycle (1)

- In its simplest form, instruction processing consists of **two** steps:
  1. The processor **reads** (*fetches*) instructions from memory one at a time
  2. It then **executes** the instruction
- Program execution consists of repeating the process of instruction fetch and instruction execution
- Instruction execution may involve several operations and depends on the nature of the instruction



Eddie Law 23

## Registers and Instruction Cycles


- In an instruction cycle
  - PC holds the address of next instruction to be fetched
  - Fetched instruction is placed in IR
  - Then PC is incremented after the fetch
  - Then execution of instruction

Eddie Law

## An Instruction Cycle (2)


(a) Some books called it a three-stage pipeline

- i.e., “fetch-and-decode” for the “fetch” stage



(a)

(b) A superscalar CPU



(b)

## Short Notes on Superscalar and Multicore (1)

- Super-scalar processors
  - Multiple instructions can be dispatched during a single clock cycle
  - Keeping track of multiple instructions in-flight, but all instructions are from a single program; that is, it is still just **one process**
  - Hence, there is only one instruction counter (different from multi-core)
  - For only "one instruction counter," and it is technically true that the code will run and experience no disparity unless using some branch prediction schemes (speculative execution: simultaneously execute both branches and throw away the "wrong" prediction's result)

## Short Notes on Superscalar and Multicore (2)

- Multi-core
  - Multiple instruction streams can execute simultaneously
  - The important part is that each core (executing with its own instruction counter) can also be super-scalar in order to execute each single process more quickly!

Eddie Law 27

## A Typical Set of Registers

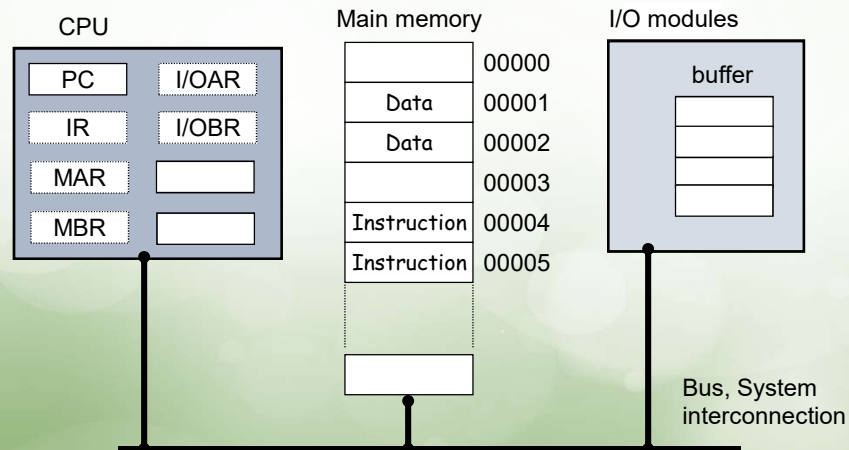
- PC - Program Counter
- IR - Instruction Register
- MAR - Memory Address Register
- MBR - Memory Buffer Register
- I/OAR - I/O Address Register
- I/OBR - I/O Buffer Register

Usually, some of these registers are directly visible to system programmers, but most of them are **not**

Eddie Law 28




## Computer Components: Top Level View



Eddie Law 29

## Reading/Writing Memory



### A side notes on addressing:

1. If using decimal addressing, then 5-digit address gives 10000 address space
2. If using hexadecimal addressing, i.e., 0...9,a,b,c,d,e,f, then 5-digit address =  $2^{20} = 1\text{ M} = 1,048,576$  address space

If each address points to 1 byte, then case (2) above has 1 MBytes RAM

If each address points to 2 bytes, then case (2) above has 2 MBytes RAM

A normal PC usually points to only 1 byte per address space

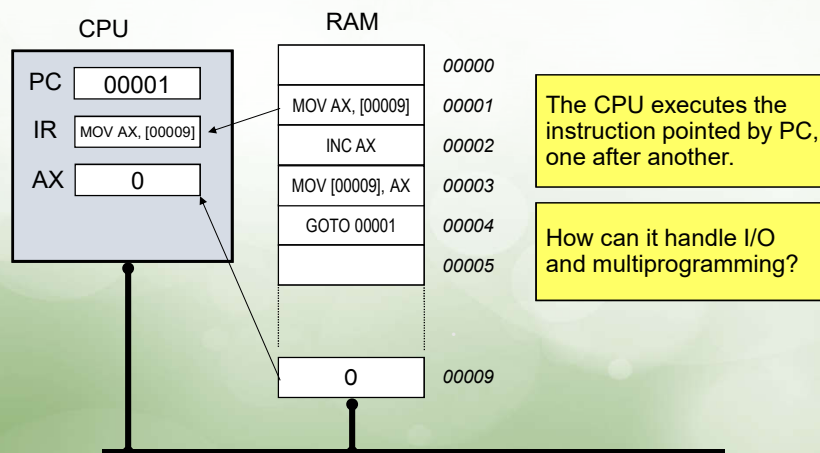
Eddie Law 30

## Categories of Instructions

- **Processor-memory:** transfer data between processor and memory
- **Processor-I/O:** transfer data between processor and I/O module
- **Data processing:** perform arithmetic or logic operation on data
- **Control:** change the sequence of execution

Eddie Law

## Instruction Execution



Eddie Law 32

## Interrupts

- An interruption of the normal processing of processor (e.g. from I/O, memory)
- A way to improve processor utilization

Eddie Law 33

## Classes of Interrupts

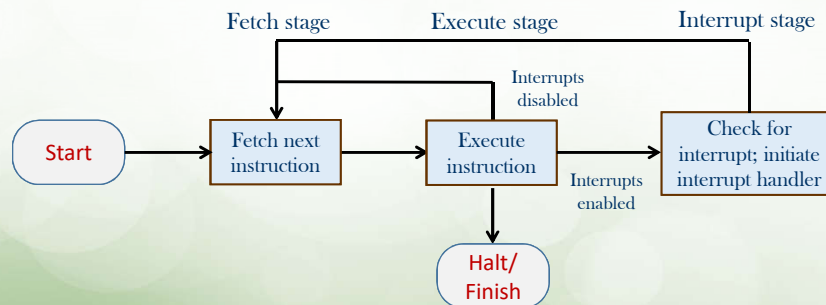
- Program
  - Due to results of instruction executions
  - Examples: arithmetic overflow, division by zero, execute an illegal machine instruction, and reference outside a user's allowed memory space (fault segmentation)

Why interrupt? • Timer

- I/O
  - Generated by a timer within the processor
  - Generated by I/O controller, to signal normal completion of an operation or a variety of error conditions
- Hardware failure
  - Generated by a failure: power failure or memory parity error

Eddie Law 34

## Interrupt Cycles with Interrupts



Eddie Law 35

## Interrupt Stage

- Processor checks for interrupts
- If no interrupts, fetch the next instruction from the current process
- If an interrupt is pending, suspend execution of the current process, and execute the interrupt handler

Eddie Law 36


## Interrupt Handler or Interrupt Service Routine

- Usually a small process which determines the nature of an interrupt and performs whatever actions are needed
- When the CPU is interrupted, control is transferred to this interrupt handler process
- Upon finishing, control is transferred back to the interrupted process
- It could generally be a part of an operating system, e.g., mouse click

Eddie Law 37

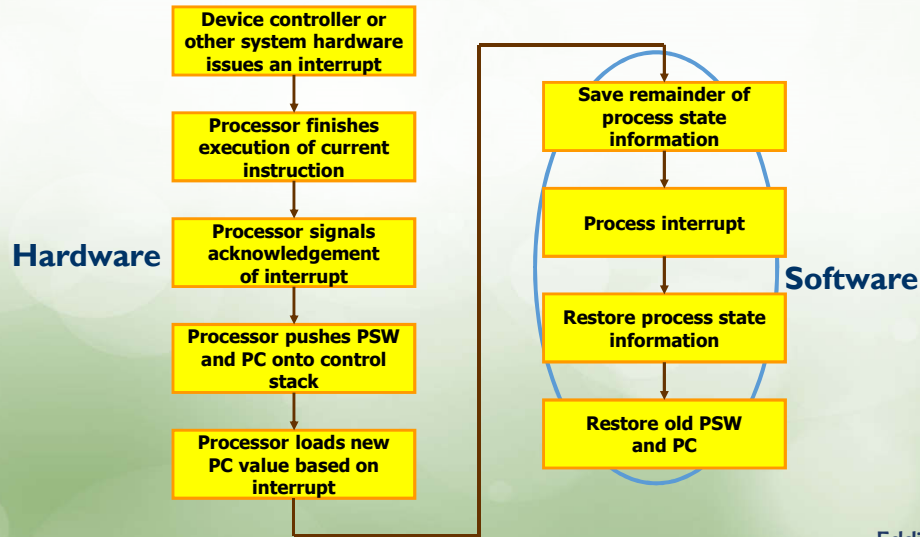
## Transfer of Control via Interrupts

- The processor and the OS are responsible for suspending a user program, and then resuming it at the same point



Eddie Law 38


## Procedure of a Simple Interrupt



Eddie Law 39

## A Program Flow without Interrupts


- Short I/O waiting time



Eddie Law 40

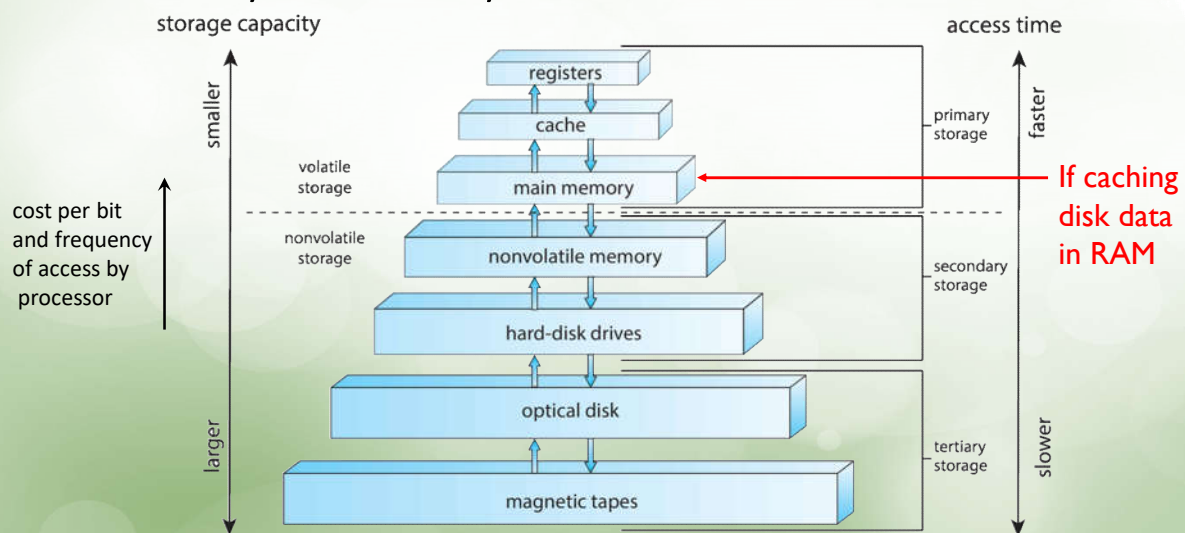


## A Program Flow with Interrupts



Eddie Law 41

## Memory Hierarchy



Eddie Law 42


## Disk Cache

- Could be some proprietary implementations
- A portion of main memory used as a buffer to temporarily hold data for the disk
- Disk writes are clustered
- Some data written out may be referenced again
- The data are retrieved rapidly from the software cache instead of slowly from disk

Eddie Law 43

## Cache Memory

- Invisible to operating system
- Processor speed is much faster than memory access speed
- Shorten the data retrieval time, in general
- Virtually increase the access speed of memory on average
- Small caches have a significant impact on performance



Eddie Law 44


## Cache Memory: How It Works

- Operation of cache: saves a portion of data from main memory
- How to do it?
- When data is requested, processor first checks cache(s)
- If not found in cache(s), the block of memory containing the needed information is moved to the cache


Eddie Law 45

## Caches and Main Memory

- Single cache



- Three-level cache




Eddie Law 46

## Basis of Validity


- Based on the principle known as the **locality of reference**

Eddie Law 47

## Cache/Main Memory Example




- Number of blocks in main memory =  $M = 2^n/K$
- Number of slots in cache =  $C \ll M$
- Size of a tag is small to allow quick search
  - But a tag maps to a number of blocks



Eddie Law 48

## I/O Modules

- Example: 2<sup>nd</sup> memory, such as hard drive
- Hardware structure: mechanical motions are slow




Eddie Law 49

## I/O Communication Techniques

- Programmed I/O
- Interrupt-driven I/O
- Direct Memory Access (DMA)

Eddie Law 50

## I/O Communication Techniques



Eddie Law 51

## Programmed I/O

- No interrupts occur
- Processor is kept busy checking status
- Polling

I/O access speeds are always much slower than the CPU. It is inefficient for the CPU to wait for I/O completion in a tight loop. (busy waiting). More on this in later chapters.

Eddie Law 52



## Programmed I/O

```
procedure readString(var s);  
repeat  
    Send I/O command "go read a word"  
    repeat  
        Read I/O status  
    until I/O done  
    Read word from I/O module  
    Write word into memory  
until finished reading  
.....
```

Pseudo-code

Eddie Law 53

## Interrupt-Driven I/O

- No busy waiting.
- Processor can proceed to do other things when I/O is in progress
- When I/O is done, the CPU is interrupted

Still consumes a lot of processor time  
because every word read or written  
passes through the processor

Eddie Law 54

## Interrupt-Driven I/O

### Install interrupt handler

```

...
procedure readString(var s);
repeat
  Send I/O command "go read a word"
  ...
  Now CPU does something else
  No need to check I/O status
  ...
until finished reading
.....

```

```

/* interrupt handler */
Read word from I/O module
Write word into memory
return

```

When the I/O module finished reading the word, it sends an interrupt to the CPU. The CPU runs an interrupt handler which moves a character(word) to memory

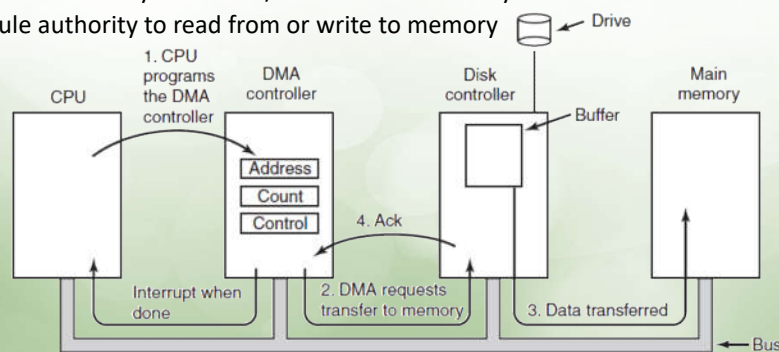
Interrupt-driven I/O is still inefficient in data transfer of large amount because the CPU has to transfer the data word by word between I/O module and memory

Afterwards, control is returned to the program which continues to read the next character/word

Eddie Law 55

## Direct Memory Access


- Goal: relieves CPU's responsibility for data exchange with main memory
- DMA
  - Control exchanges of data directly between I/O device and memory
  - CPU grants I/O module authority to read from or write to memory



Eddie Law 56

## Direct Memory Access: Operations

- Operations
  - Processor is only involved at the beginning and end of the transfer, continues with other work
  - Data block transfers to or from memory directly
  - An interrupt is sent when the transfer of an entire block is complete
- Any limitations?



Eddie Law 57

## DMA: The Example

```

procedure readString(var s);
  [CPU requests DMA module to
   read some data]
  ...

```

...  
**Now, CPU does *something else***  
**No checking on I/O status**  
 ...

The DMA module starts reading each word of the data and save them in the memory; transparent to the process

DMA I/O is more efficient in large data transfer because the interaction with the I/O module and data transfer between I/O module and memory are performed by the DMA module.

After the DMA has transferred all the data requested to the memory, it notifies that it has finished the I/O by sending the CPU an interrupt

Eddie Law 58

## Summary

- Processor, memory, I/O, bus
- How processor runs program
- How processor interacts with memory and I/O
- Interrupt
- Cache
- Programmed I/O, Interrupt-driven I/O and DMA

Eddie Law 59

## Next Topic

- Operating System Overview
- Read Chapter 2

Eddie Law 60