



澳門理工學院
Instituto Politécnico de Macau
Macao Polytechnic Institute

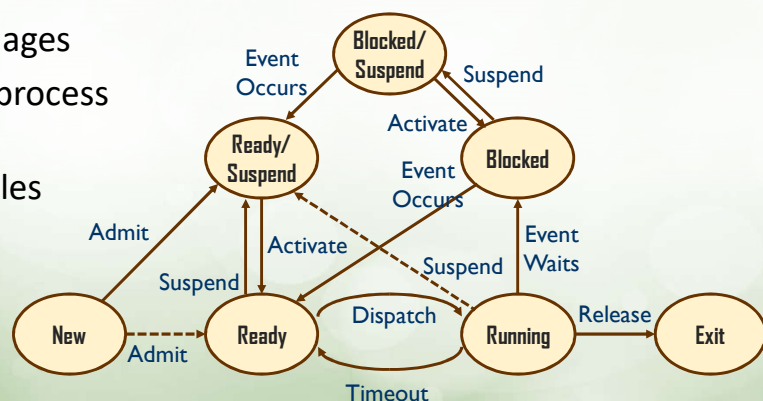
School of Applied Sciences (B.Sc. in Computing)

Notes #4: Processes and Threads

COMP 213
(21121/21221)
Operating Systems
2019-2020 1st Semester

Review

- PCBs and process images
- 5-state, and 7-state process models
- Different control tables



Eddie Law

Topics

- Parent-child processes
- Process vs. Thread
- Multithreading
- Process image in multithreaded process
 - Thread Control Block (TCB)
- Thread switching
- Benefits of threads
- Symmetric multiprocessing (SMP) and multi-core
- Chapters 4.2, 4.1, + this note

Eddie Law 3

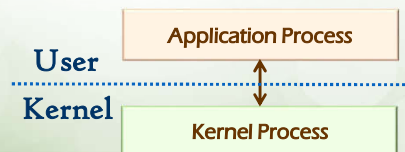
Processes

- Have been studying about “processes”
- 7-state process models for multiprogramming
 - Resource ownership – a virtual address space to hold process image
 - Scheduling/execution – interleaving execution path with other processes
- But no discussion on user-level process and kernel-level process relationship yet

Eddie Law 4

Linking User-level and Kernel-level Processes

- It's natural to have a kernel-level process working with a user-level process
- One-to-one mapping between user-level and kernel-level process?
 - E.g., MS-DOS



Eddie Law 5

Process: Concurrency

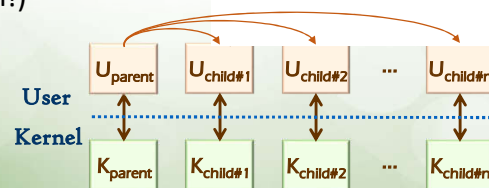
- Process
 - An “abstract” model
 - Exchanges of control should be done automatically and implicitly on multiprogramming paradigm
- Multiprogramming is natural on multicore systems
- But can a process be run in parallel on multiprocessor? Even on a uniprocessor system??

Eddie Law

Process and Thread: Concurrency

- How can we run **ONE** process in parallel?

- Parent-child processes?
- Multi-threading?
(discuss soon!)



Eddie Law 7

Start with Parent-Child Process Design

- For example, in Linux, BSD...
- `fork()`: for creation of a child process with an image copy of the parent process
 - `fork()` returns 0 for child process, returns child PID for parent


```
pid = fork()
if (pid == 0) cout << "child process";
else cout << "parent process";
```
- `wait()`: for synchronization of processes
 - A parent waits for the termination of child process
 - Parent process is awakened; parent can determine which of its child processes terminates

Eddie Law 8

How Expensive to Create Process?

- Must construct new PCB
 - Inexpensive
- Must set up new tables for address space
 - More expensive
- Copy I/O state (file handles, etc)
 - Medium expense
- Copy data from parent process? (e.g., `fork()`)
 - Semantics of `fork()` are that the child process **gets a complete copy of the parent memory and I/O state**
 - Originally very expensive
 - Much less expensive with “copy on write”

Eddie Law

Parallelism: An Example

- Consider calculating roots of a quadratic equation

$$\frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

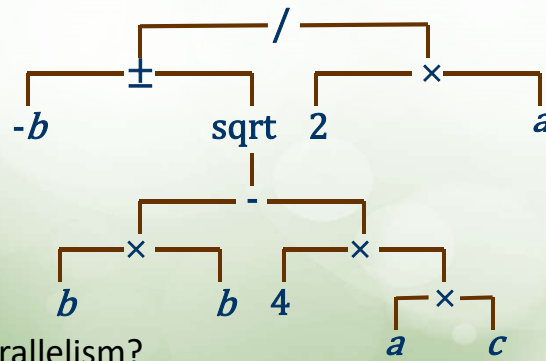
- To program it in **sequential statements**:

$$\begin{aligned} t_1 &= b \times b \\ t_2 &= 4 \times a \\ t_3 &= t_2 \times c \\ t_4 &= t_1 - t_3 \\ t_5 &= \text{sqrt}(t_4) \\ t_6 &= -b + t_5 \\ t_7 &= -b - t_5 \\ t_8 &= 2 \times a \\ t_9 &= t_6 \div t_8(\text{ans1}) \\ t_{10} &= t_7 \div t_8(\text{ans2}) \end{aligned}$$

Eddie Law 10

If Breaking Down the Equation

- Decompose the equation from observation (outside in)

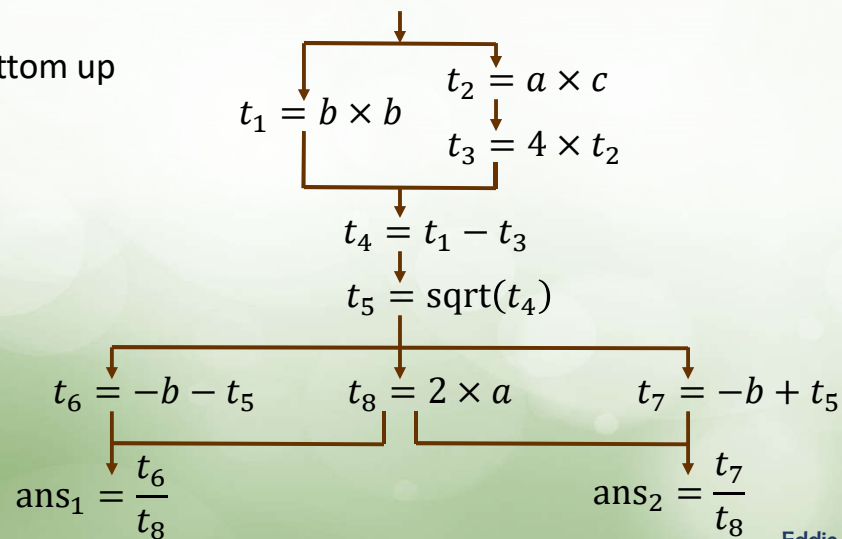


- Do you see the parallelism?

Eddie Law 11

Parallelized the Computations

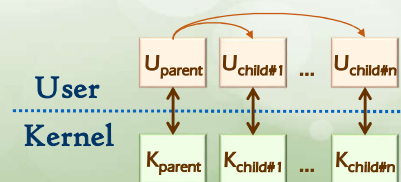
- From bottom up



Eddie Law 12

Fork and Join Process Notation

- (Beware: following not the Unix's `fork()` call)
- `fork()` primitive
 - Creates a named parallel process that executes a designated code fragment
- `join()` primitive
 - Waits for a parallel process created by a fork to terminate
- `quit()`
 - Terminates a forked process



Eddie Law

Fork/Join: Example

process Main

```

fork A;  $t_1 = b * b$ ; join A
 $t_4 = t_1 - t_3$ ;  $t_5 = \text{sqrt}(t_4)$ 
fork B; fork C;  $t_6 = -b - t_5$ ; join B; join C
fork D;  $t_{10} = t_6 / t_8$ ; join D
  
```

end process Main

process A

```

 $t_2 = 4 * a$ ;  $t_3 = t_2 * c$ 
  
```

end process A

process B

```

 $t_8 = 2 * a$ 
  
```

end process B

process C

```

 $t_7 = -b + t_5$ 
  
```

end process C

process D

```

 $t_9 = t_7 / t_8$ 
  
```

end process D

Eddie Law

Co-begin, Co-end

- One way to restrict programs to make code less obscure
 - Is to make the control statements of the program properly nested
- **Co-begin co-end** pairs
 - Enforce the proper nesting of parallel processes.
- Statements enclosed by a **co-begin co-end** pair are executed in parallel

Eddie Law

Co-begin, Co-end

- Unlike **fork** and **join**, non-nested dependencies between processes cannot be expressed
- The **co-begin, co-end** primitives also in form of **par-begin, par-end**

Eddie Law

Co-begin, Co-end: Example

begin Main

cobegin

$t_1 = b * b;$

begin

$t_2 = 4 * a;$

$t_3 = t_2 * c$

end

coend

$t_4 = t_1 - t_3$

$t_5 = \text{sqrt}(t_4)$

cobegin

$t_6 = -b - t_5;$

$t_8 = 2 * a;$

$t_7 = -b + t_5;$

coend

cobegin

$t_{10} = t_6 / t_8;$

$t_9 = t_7 / t_8$

coend

end Main

Eddie Law

Process and Thread

- Faster performance through parallelization though multiple processes
- Any problems?
 - For each process, resources allocated by OS, including memory to hold the process image
 - All duplicated code base for duplicated processes, too much redundancy, waste resources
 - Further, follows an execution path that may be interleaved with other processes → bad case would be that some processes may have to wait
- A newer design → modern OS runs “thread” in kernel
 - Multithreading design
 - What is a thread?

Eddie Law 18

Process Image

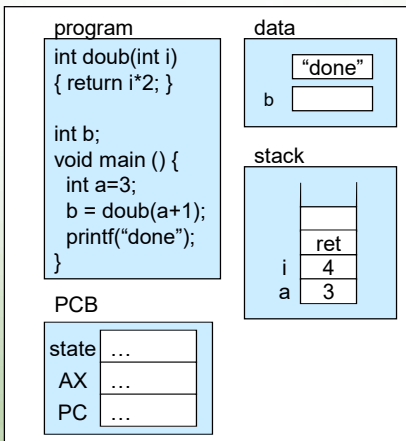
- This example: 1 process (1 thread)

Each process is allocated some resources, including the memory to hold the process image.

For two processes A and B:

- Process A **cannot** access the memory of process B.
- Process A and B **cannot** share data in memory.

Process image in RAM



19

Eddie Law

Process/Thread Execution (1)

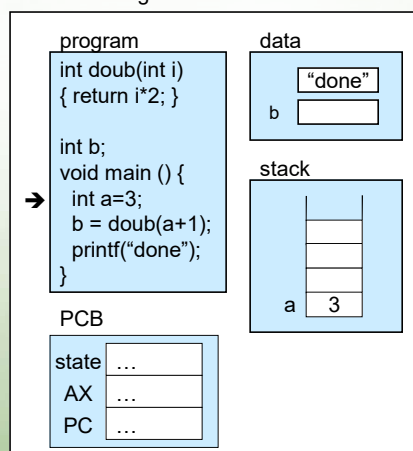
Each process has **one** thread of execution. At any time, only one instruction is being executed.

The arrow indicates the current instruction being executed.

Its relation with the register PC and the PC value in the PCB.

Also notice how the stack changes in the course of execution.

Process image in RAM



Eddie Law 20

Process/Thread Execution (2)

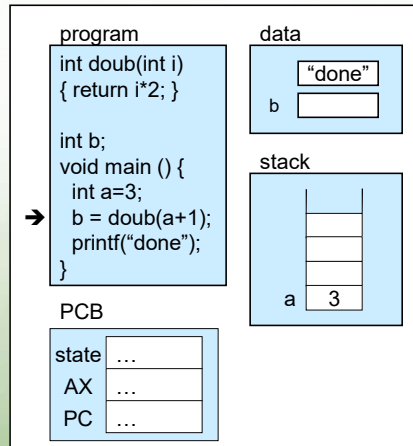
Each process has **one** thread of execution. At any time, only one instruction is being executed.

The arrow indicates the current instruction being executed.

Its relation with the register PC and the PC value in the PCB.

Also notice how the stack changes in the course of execution.

Process image in RAM



Eddie Law 21

Process/Thread Execution (3)

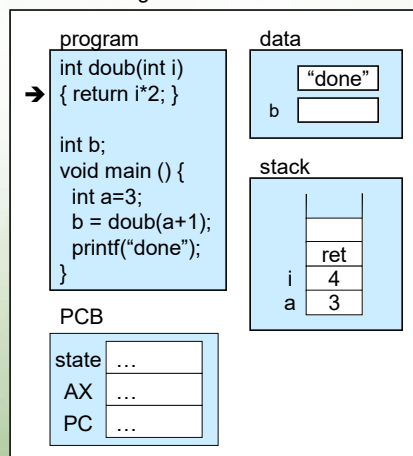
Each process has **one** thread of execution. At any time, only one instruction is being executed.

The arrow indicates the current instruction being executed.

Its relation with the register PC and the PC value in the PCB.

Also notice how the stack changes in the course of execution.

Process image in RAM



Eddie Law 22

Process/Thread Execution (4)

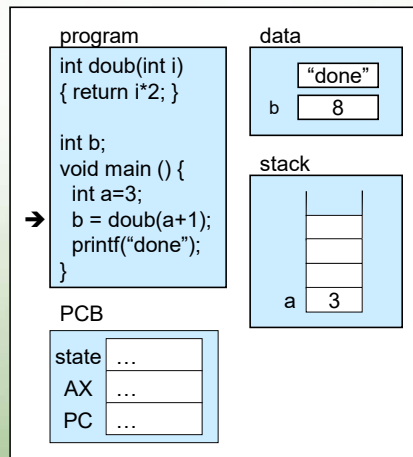
Each process has **one** thread of execution. At any time, only one instruction is being executed.

The arrow indicates the current instruction being executed.

Its relation with the register PC and the PC value in the PCB.

Also notice how the stack changes in the course of execution.

Process image in RAM



Eddie Law 23

Process/Thread Execution (5)

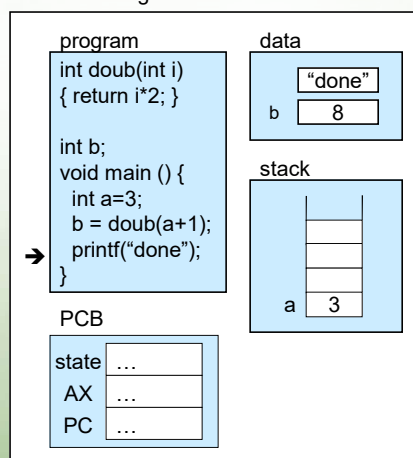
Each process has **one** thread of execution. At any time, only one instruction is being executed.

The arrow indicates the current instruction being executed.

Its relation with the register PC and the PC value in the PCB.

Also notice how the stack changes in the course of execution.

Process image in RAM



Eddie Law 24

Process/Thread Execution (6)

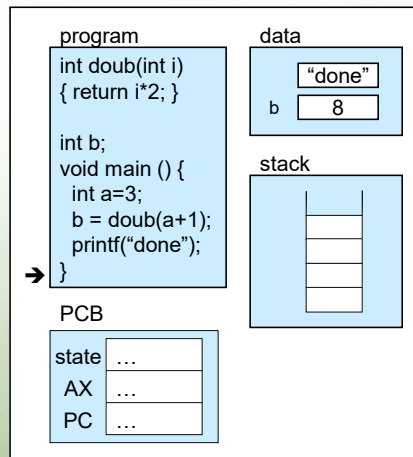
Each process has **one** thread of execution. At any time, only one instruction is being executed.

The arrow indicates the current instruction being executed.

Its relation with the register PC and the PC value in the PCB.

Also notice how the stack changes in the course of execution.

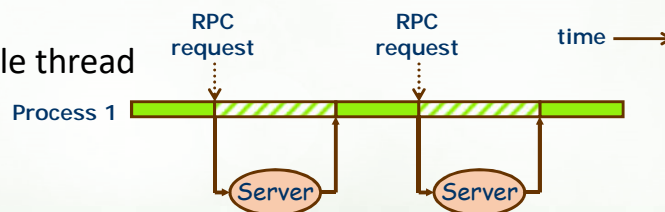
Process image in RAM



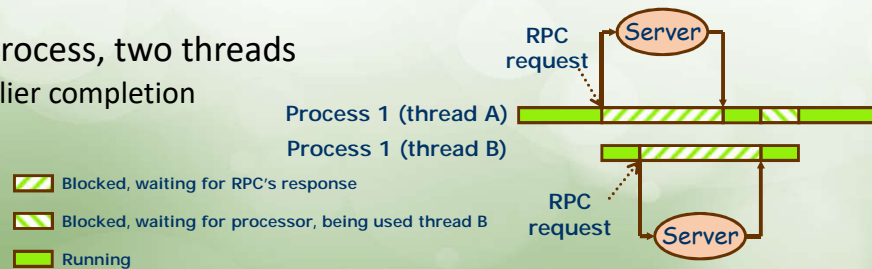
Eddie Law 25

Example: Remote Procedure Call (RPC)

- One process, single thread



- One process, two threads
 - Earlier completion

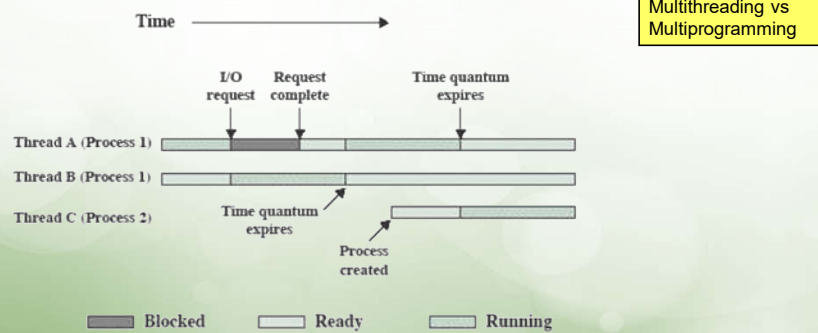


- Blocked, waiting for RPC's response
 - Blocked, waiting for processor, being used thread B
 - Running

Eddie Law

Multithreading

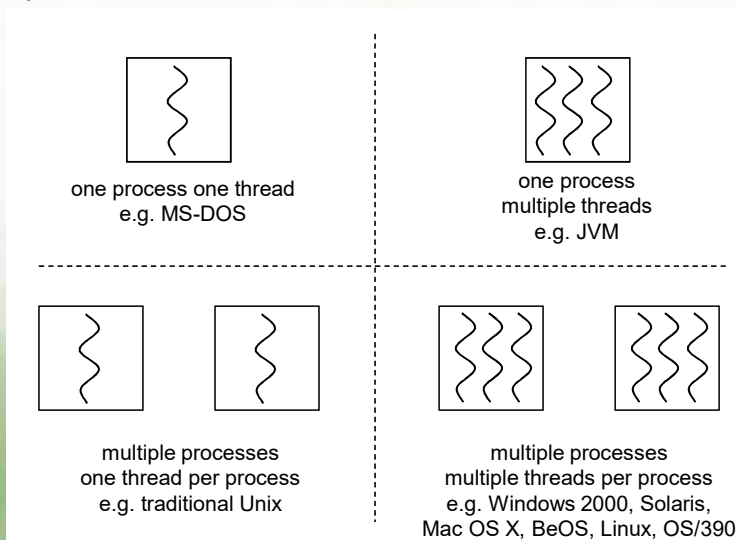
- Multiprogramming, multiple threads of execution



Multithreading example on uniprocessor

Eddie Law 27

OS Support for Threads and Processes



Eddie Law 28

Process vs. Thread

- A process has one or more threads
- Process – Unit of resource ownership
 - Some resources allocated by OS, including memory, open files
- Thread – Unit of dispatching
 - An execution path
 - Execution may be interleaved with other threads / processes

Eddie Law 29

Process and Threads

- In a multithreaded environment, a process is defined as the unit of resource allocation and a unit of protection
- The following are associated with processes:
 - A virtual address space that holds the process image
 - Protected access to processors, other processes (for inter-process communication), files, and I/O resources (devices and channels)

Eddie Law 30

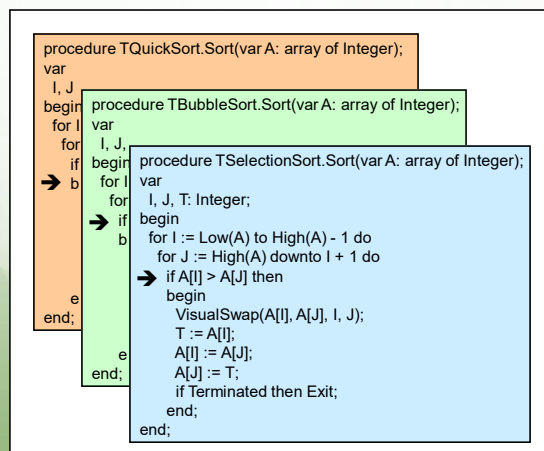
Process and Threads (cont'd)

- Within a process, there may be one or more threads, each with the following:
 - A thread execution state (Running, Ready, etc.)
 - A saved thread context when not running; one way to view a thread is as an independent program counter operating within a process
 - An execution stack
 - Some per-thread static storage for local variables
 - Access to the memory and resources of its process, shared with all other threads in that process

Eddie Law 31

Multithreading, Example

Process ThdDemo

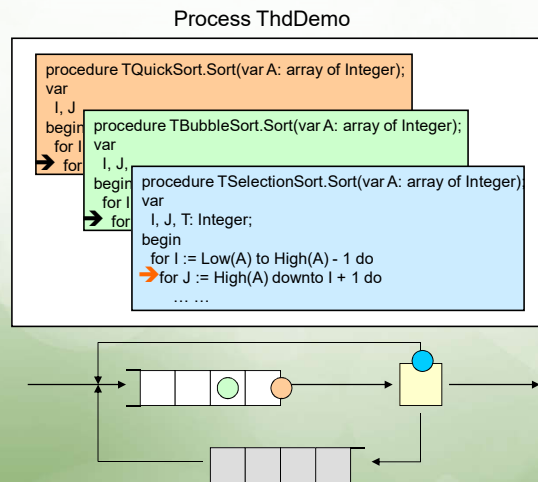


When we run the *program thddemo.exe*, we create a *process* which consists of three *threads*: one thread for each sort procedure.

Actually, there is one more thread, the primary thread, in the process thddemo.

Eddie Law 32

Thread as Unit of Dispatching

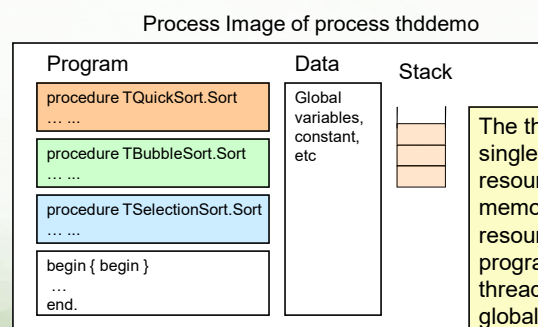


The OS schedules the three threads. Each thread is dispatched in turn, one at a time. The three threads of the process appear to be running 'at the same time'.

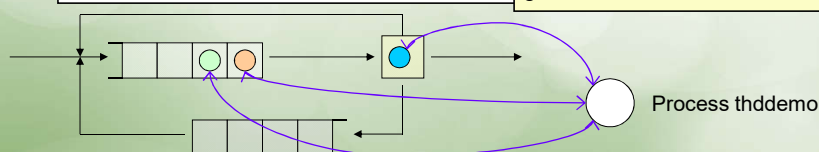
In Chap 3, the "process ball" that represents a process is in fact the single thread within the processes. And process switching is in fact **thread switching**.

Eddie Law 33

Process as Unit of Resource Ownership

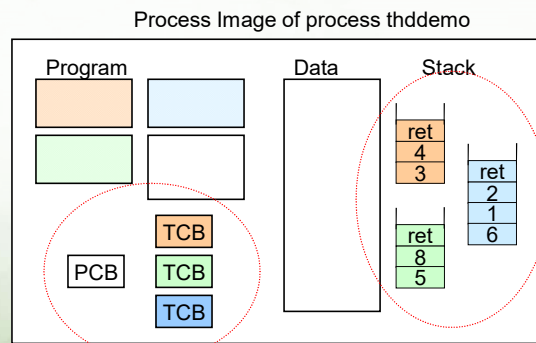


The three threads belong to a single process, which owns resources. They share memory, opened files and I/O resources. Only a copy of the program is stored. The three threads use the same copy of global variables.



Eddie Law 34

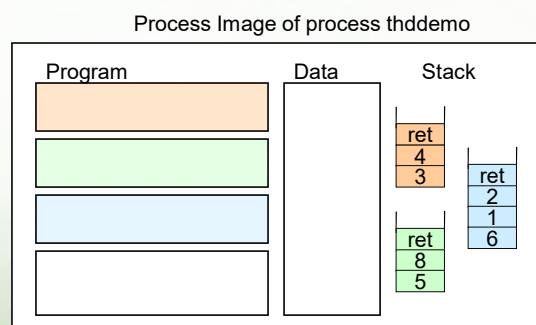
Stack and PCB in Multithreaded Process



We see that a single copy of program and data are shared by the multiple threads of a multithreaded process. However, a single stack and a PCB is not sufficient.

Eddie Law 35

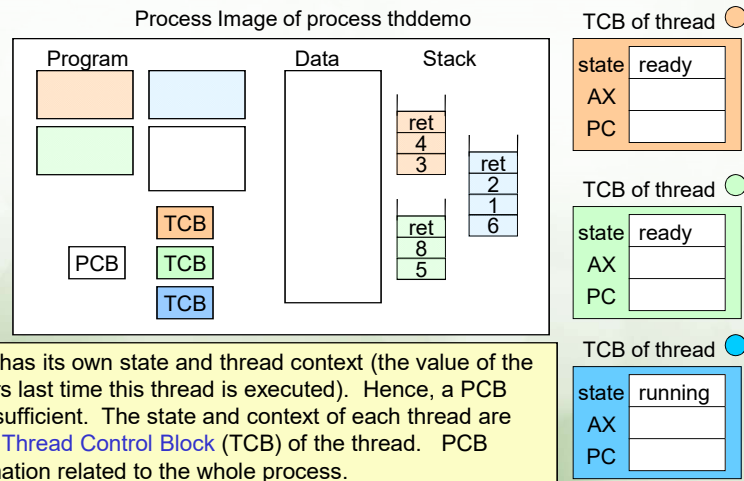
Each thread needs a user stack



Stack stores local variables in function call. Since each thread can run independently, their local variables appear and disappear in any order. We cannot use a single stack for multiple threads. Instead, each thread has its own stack.

Eddie Law 36

PCB and TCB (Thread Control Block)



Eddie Law 37

Splitting Information on PCB

- Process Identification:
 - Process ID, User ID [P], Thread ID [T]
- Processor State Information [T]
- Process Control Information
 - Scheduling [T] [P] and state information [T] [P]
 - Data structuring [T] [P]
 - Inter-process communication [P]
 - Process privileges [P]
 - Memory management [P]
 - Resource ownership and utilization [P]

Eddie Law 38

Threads in Process: Resource Sharing

- Items shared by all threads in a process
- Items private to each thread

Per process items

Address space
Global variables
Open files
Child processes
Pending alarms
Signals and
signal handlers
Accounting
information

Per thread items

Program counter
Registers
Stack
State

Eddie Law

Creation and Termination

- Usually, a process is created with a single thread (called primary thread in Windows) This thread may create more threads.
- A process terminates when
 - in Windows: the primary thread terminates
 - in Unix: all threads of the process terminate

Eddie Law 40

Activities similar to Processes

- Threads have execution states and may synchronize with one another.
 - Similar to processes
- We look at these two aspects of thread functionality in turn.
 - States
 - Synchronisation (in Ch. 5)

Eddie Law 41

Thread Execution States

- Key **states** for a thread
 - Running
 - Ready
 - Blocked
- **Operations** associated with a change in thread state
 - Spawn (another thread)
 - Block
 - Issue: will blocking a thread block other, or all, threads
 - Unblock
 - Finish (thread)
 - De-allocate register context and stacks

What about 'Suspended' state?

Eddie Law 42

Thread Suspend State?

- Suspending a process involves suspending all threads of the process

since all threads share the same
address space

Why?

Eddie Law 43

Thread Synchronization

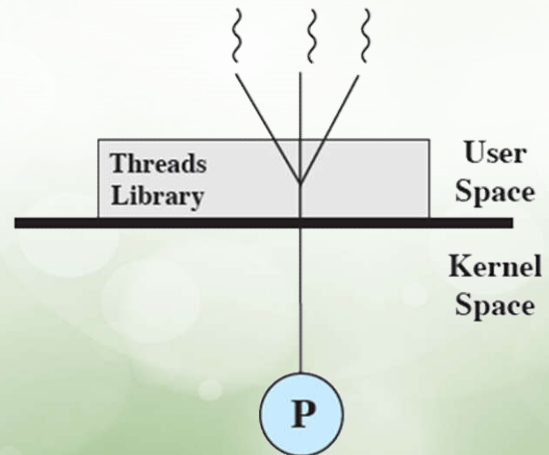
More on this in
Chapter 5

- It is necessary to synchronize the activities of the various threads
 - All threads of a process share the same address space and other resources
 - Any alteration of a resource by one thread affects the other threads in the same process

Eddie Law 44

User-Level Threads (ULTs)

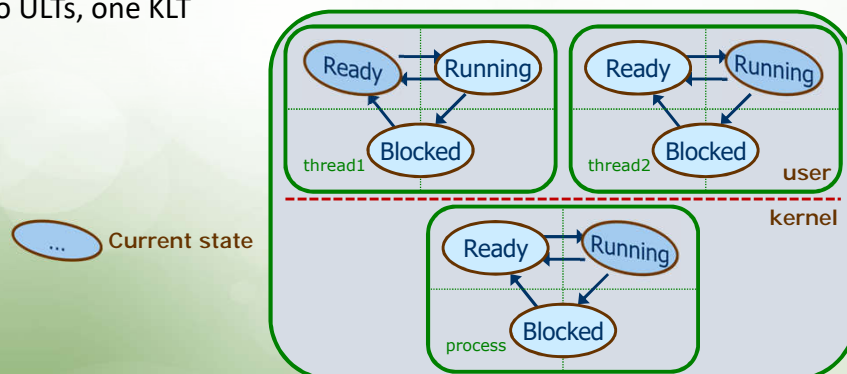
- One of the two types of threads
- All thread management is done by the application
- The kernel is not aware of the existence of thread



Eddie Law 45

Example on ULT

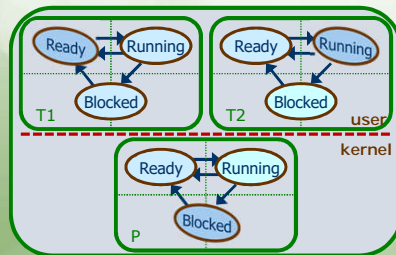
- Consider state 1 of a process
 - Two ULTs, one KLT



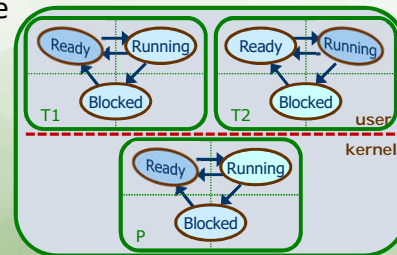
Eddie Law

E.G. (cont')

- From state 1
- Thread 2 (T2) makes a blocking I/O call → kernel invokes I/O → places Process (P) to blocked state
- The T library maintains T2 in running state



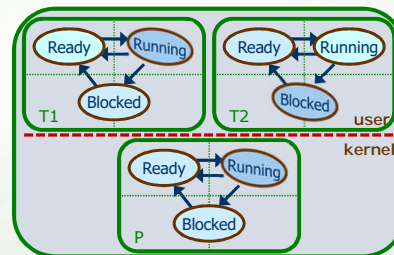
- From state 1
- Clock interrupt → moves to kernel → time slot for P is over → places P to ready state
- The T library maintains T2 in running state



Eddie Law

E.G. (cont')

- From state 1
- T2 is waiting data from T1 → T2 is blocked and T1 becomes ready
- Kernel is not invoked → state of P is unchanged



Eddie Law

Advantages of ULTs

- Thread switching does not require kernel mode privileges
- Scheduling can be application specific
- ULTs can run on any OS

Eddie Law 49

Disadvantages of ULTs

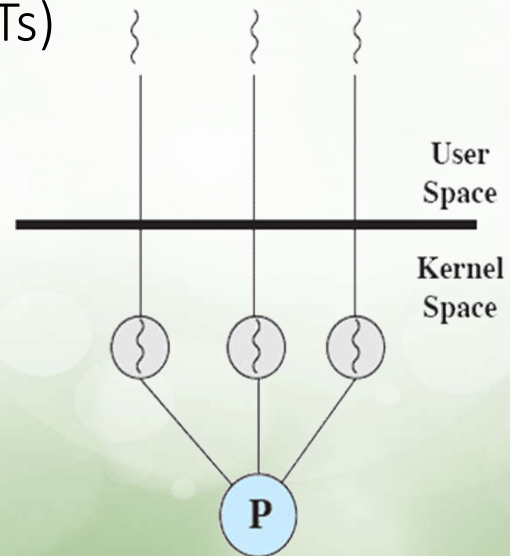
- In a typical OS, many system calls are blocking → all of the threads within the process are blocked
- In a pure ULT strategy, a multithreaded application cannot take advantage of multiprocessing



Eddie Law 50

Kernel-Level Threads (KLTs)

- Thread management is done by the kernel
- no thread management is done by the application
- Windows is an example of this approach



Eddie Law 51

Advantages of KLTs

- The kernel can simultaneously schedule multiple threads from the same process on multiple processors
- If one thread in a process is blocked, the kernel can schedule another thread of the same process
- Kernel routines can be multithreaded



Eddie Law 52

Disadvantage of KLTs

- The transfer of control from one thread to another within the same process requires a mode switch to the kernel

Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840

Null Fork: the time to create, schedule, execute, and complete a process/thread that invokes the null procedure (i.e., the overhead of forking a process/thread)

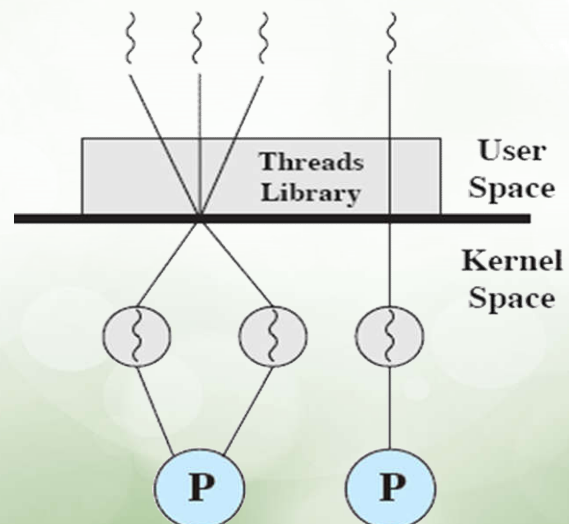
Signal-Wait: the time for a process/thread to signal a waiting process/thread and then wait on a condition (i.e., the overhead of synchronizing two processes/threads together)



Eddie Law 53

Combined Approaches

- Thread creation is done in the user space
- Bulk of scheduling and synchronization of threads is by the application
- Solaris is an example



Eddie Law 54

Relationship between Threads and Processes

Threads: Processes	Description
1:1	Each thread of execution is a unique process with its own address space and resources.
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.
1:M	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.
M:N	It combines attributes of M:1 and 1:M cases.

Eddie Law 55

Process, in Summary

- Have a virtual address space which holds the process image, which contains:
 - program and data
 - one stack per thread
 - PCB and one TCB per thread
- Protected access to memory, files, and I/O resources
- A process has one or more threads

Eddie Law 56

Thread, in Summary

- Has an execution state
 - Running, ready, blocked, ...
- Has an execution stack
- Thread context is saved and restored in thread switching
 - Read the following slides for how to change the steps of process switching to the steps of thread switching
- Has access to the memory and resources of its process
- Has some per-thread static storage for local variables

Eddie Law 57

Thread Switching, with Steps

1. Save processor state (incl. program counter and other registers) in the TCB: processor state information
2. Update the TCB with the new state and any accounting information
3. Move the TCB to appropriate queue – ready, blocked

Eddie Law 58

Thread Switching, Steps (cont'd)

4. Select another thread for execution
5. Update the TCB of the thread selected (new state and accounting information)
6. Update memory-management data structures
7. Restore context of the selected thread
 - Restore the previous value of the program counter and other registers from the TCB

Eddie Law 59

Thread Switching, Steps (Simplified)

- Thread switching from thread *A* to thread *B*:
 1. Save thread context of thread *A*: CPU \Rightarrow TCB *A*
 2. Update state in TCB *A*
 3. Move TCB *A* to appropriate queue
 4. Select another thread (thread *B*)
 5. Update state in TCB *B*
 6. Update memory-management data structures
 7. Restore thread context of thread *B*: TCB *B* \Rightarrow CPU

Eddie Law 60

Benefits of Threads (1)

- Less time to create / terminate a new thread than a process
- Less time to switch between two threads within the same process
- Communication among threads of the same process is easier

Eddie Law 61

Benefits of Threads (2)

- Less time to create / terminate a new thread than a process
 - When the OS creates a process, it has to allocate memory for the process image, load the program and data, and initialize other resources
 - When the OS creates a thread within an existing process, it does not need to initialize the process image. It only needs to set up a stack and TCB for the new thread.

Eddie Law 62

Benefits of Threads (3)

- Less time to switch between two threads within the same process
 - The “update memory management data structure” step is only necessary in case of thread switching between different process. (Why?)

Eddie Law 63

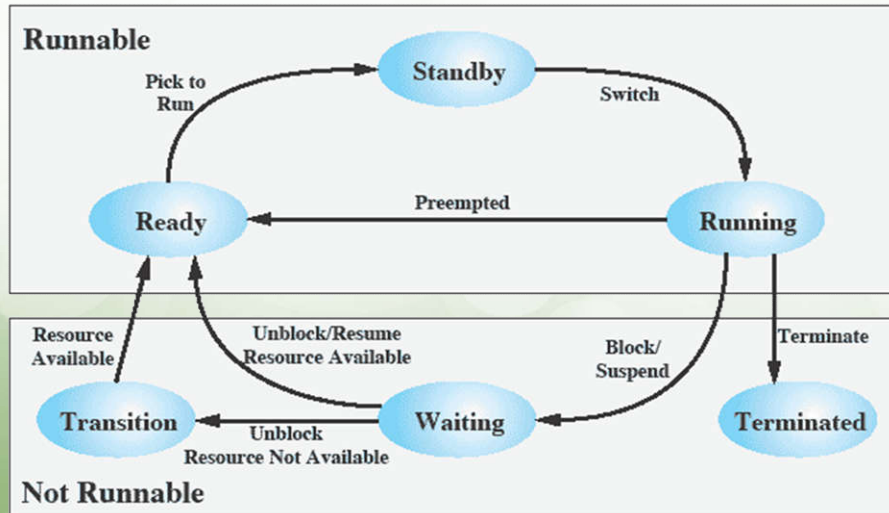
Benefits of Threads (4)

- Communication among threads of the same process is easier
 - Threads within a process share memory and files
 - Different processes have to communicate with the help of the kernel

Kernel is the core of OS

Eddie Law 64

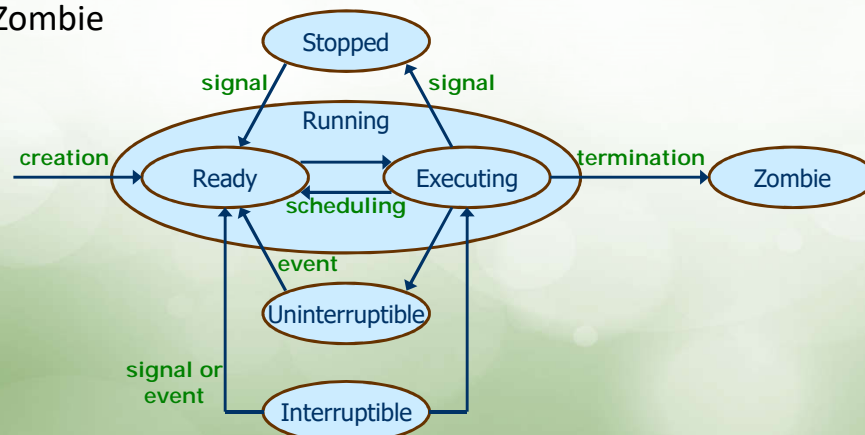
Example: Windows Thread States



Eddie Law 65

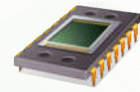
Linux States of a Process

- Process/thread model: Ready, Running, Interruptible, Uninterruptible, Stopped, Zombie



Eddie Law

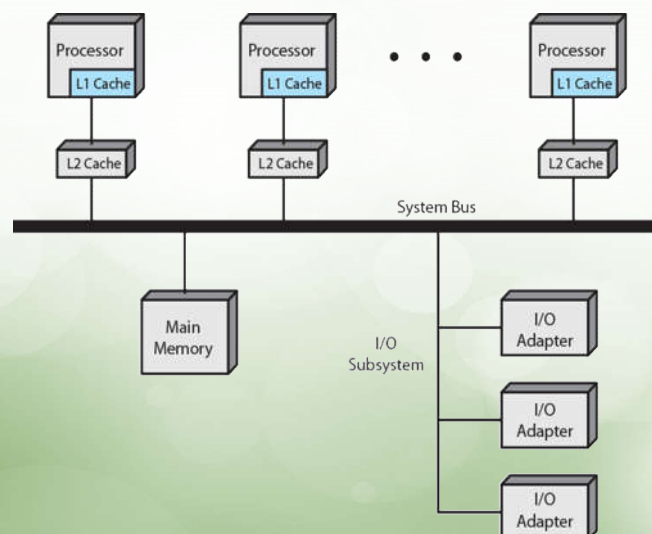
Symmetric Multiprocessors (SMP) and Multi-core



- A stand-alone computer system with the following characteristics:
 - Two or more similar processors of comparable capability
 - Processors share the same main memory and are interconnected by a bus or other internal connection scheme
 - Processors share access to I/O devices
 - All processors can perform the same functions
 - The system is controlled by an integrated operating system that provides interaction between processors and their programs at the job, task, file, and data element levels
- Multiprocessing - more than one processors in a machine

Eddie Law 67

SMP Organization



Eddie Law 68

SMP Advantages

Performance

- a system with multiple processors will yield greater performance if work can be done in parallel

Scaling

- vendors can offer a range of products with different price and performance characteristics

Availability

- the failure of a single processor does not halt the machine

Incremental Growth

- an additional processor can be added to enhance performance

Eddie Law 69

Multicore Computer



- Also known as a multiprocessor chip
- Combines two or more processors (cores) on a single piece of silicon (die)
 - Each core consists of all of the components of an independent processor
- In addition, multi-core chips also include L1 cache and in some cases L2 cache

Eddie Law 70

Performance on Multicore

- Amdahl's Law

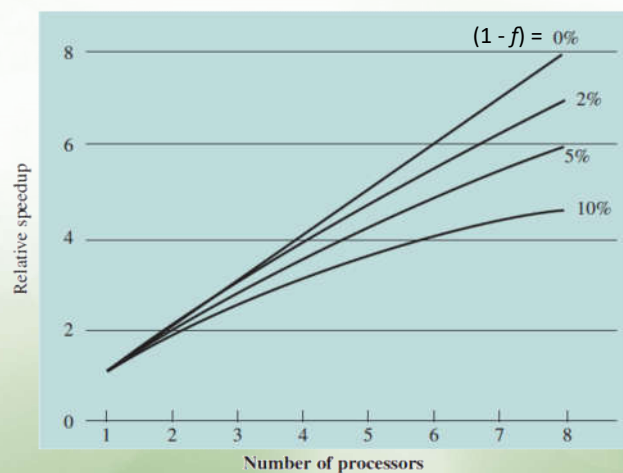
$$\text{Speedup} = \frac{\text{time to execute program on a single processor}}{\text{time to execute program on } N \text{ parallel processors}} = \frac{1}{(1 - f) + \frac{f}{N}}$$

- A program in which a fraction $(1 - f)$ of the execution time involves code that is inherently serial, and a fraction f that involves code that is infinitely parallelizable with no scheduling overhead

Eddie Law 71

Performance (1)

- If only 10% of the code is inherently serial ($f = 0.9$), running the program on a multicore system with 8 cores yields a performance gain of a factor of only 4.7

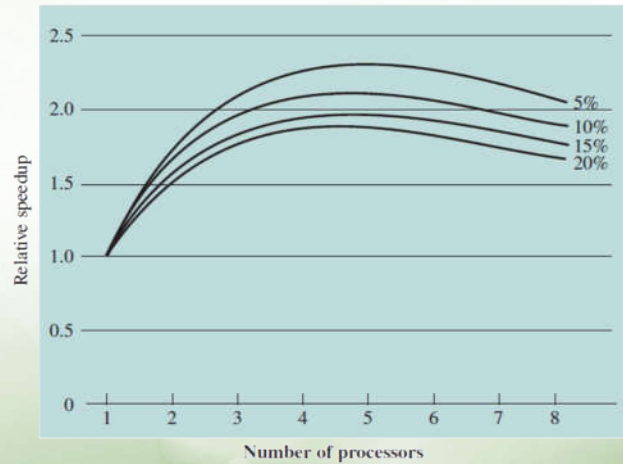


(a) Speedup with 0%, 2%, 5%, and 10% sequential portions

Eddie Law 72

Performance (2)

- Software typically incurs overhead as a result of communication and distribution of work to multiple processors and cache coherence overhead
- This results in a curve where performance peaks and then begins to degrade because of the increased burden of the overhead of using multiple processors



(b) Speedup with overheads

Eddie Law 73

Summary



- Process/related to resource ownership
- Thread/related to program execution
- Types of threads: User-level threads and kernel-level threads
- Symmetric Multiprocessors (SMP) and multicore

Eddie Law 74

Next Topic

- Concurrency: Mutual Exclusion and Synchronization
- Read Ch. 5