

## Notes #7: Virtual Memory

COMP 213  
 (211/212)  
 Operating Systems  
 2019-2020 1st Semester

### Notice Two Characteristics (1)

- Memory references are translated by the CPU dynamically
  - logical/physical address translation is done at each memory access (for data and program)
  - OS can put the pages of a process anywhere in RAM, and it can move the pages without interfering the process

Eddie Law 2


### Notice Two Characteristics (2)

- Each page is mapped to a frame **independently**
  - A process may be broken up into pieces that do not need to be located contiguously in RAM
  - the OS may partially swap a process. E.g. page 0 and page 2 are in RAM, while page 1 is swapped out to hard disk

Eddie Law 3

### Address Translated Dynamically (1)


Suppose "a=0" is compiled to the assembly "mov [00000100], 0". At the moment that the CPU executes this instruction, the logical address 0x00000100 will be translated to the physical address 0x00026100 using the page table.



Eddie Law 4


## Address Translated Dynamically (2)

After swapping, the process image moves to a different location in RAM. When the CPU executes the instruction "mov [00000100], 0" (a=0) again, logical address 00000100 will be translated to new location automatically (physical address 00022100).



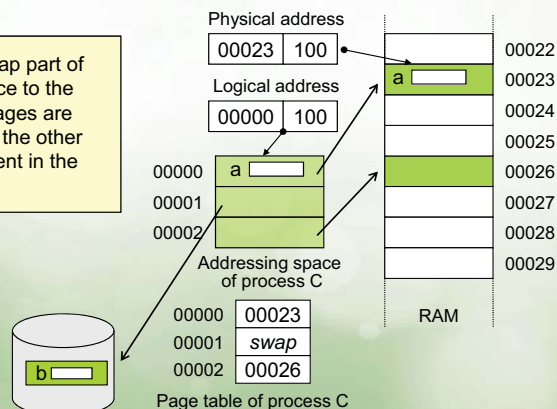
## Each Page is Mapped Independently (1)

The three pages 00000, 00001 and 00002 (the first 12k in the addressing space) appear to be contiguous to process C. But, in fact, they are mapped to separate frames in RAM.



## Each Page is Mapped Independently (2)

And the OS can swap part of the addressing space to the hard disk. Some pages are swapped out, while the other pages are still present in the RAM.



## What We'll Learn

- Virtual memory
  - what, advantages, how, page fault handling, efficiency
- Hardware support: page table, page table point, TLB
- OS support
  - esp. Page replacement algorithm
- Requirements of memory management revisited

## Virtual Memory

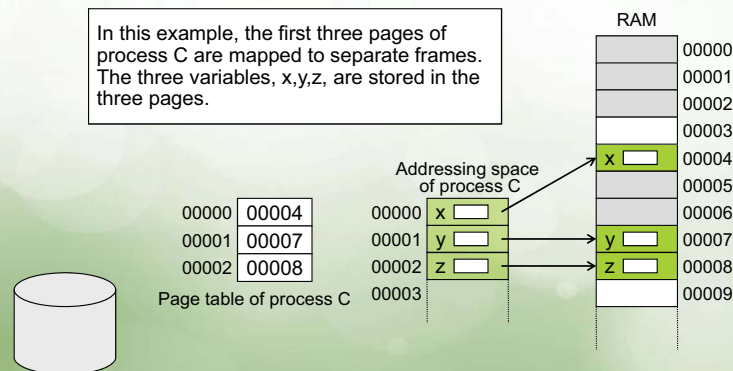
- “Simulate more RAM using hard disk space”
- The program/data used by a process do not have to be present in the RAM all the time. Pages not currently used are swapped out to disk.
- When needed, these pages are swapped in, back to some frames in RAM. This operation is invisible to the process

## Advantages of Virtual Memory

- More processes can be kept in main memory
  - only some pieces of each process are loaded
  - better utilization of the CPU
- It is possible for a process to be larger than all the main memory
  - easier life for programmer

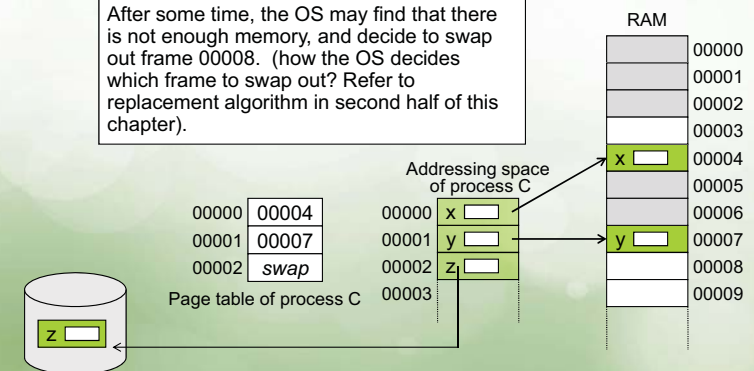
## How Virtual Memory Works (1)

In this example, the first three pages of process C are mapped to separate frames. The three variables, x,y,z, are stored in the three pages.

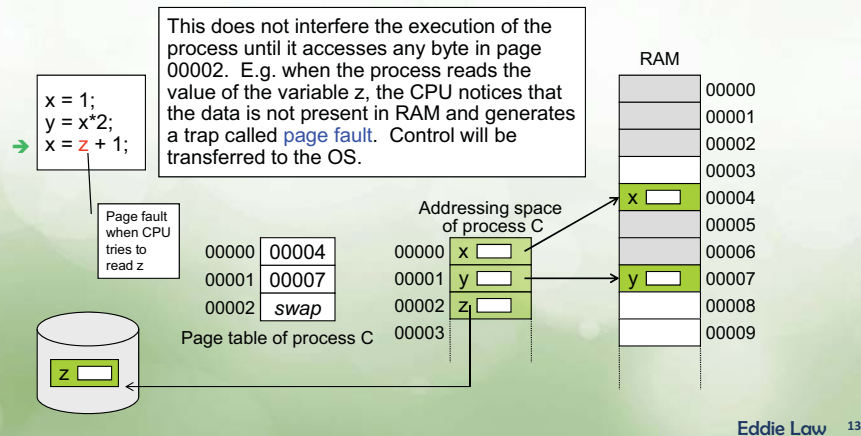


## How Virtual Memory Works (2)

After some time, the OS may find that there is not enough memory, and decide to swap out frame 00008. (how the OS decides which frame to swap out? Refer to replacement algorithm in second half of this chapter).



## How Virtual Memory Works (3)

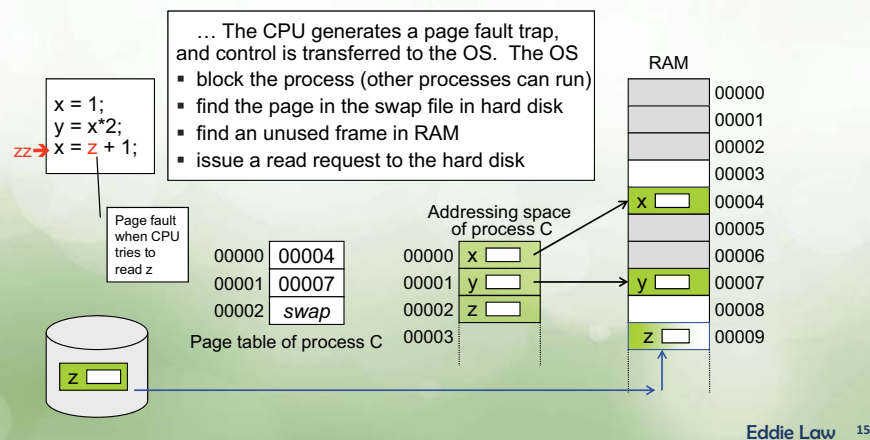


## How Virtual Memory Works (4)

This does not interfere the execution of the process until it accesses any byte in page 00002. E.g. when the process reads the value of the variable z, the CPU notices that the data is not present in RAM and generates a trap called page fault. Control will be transferred to the OS.

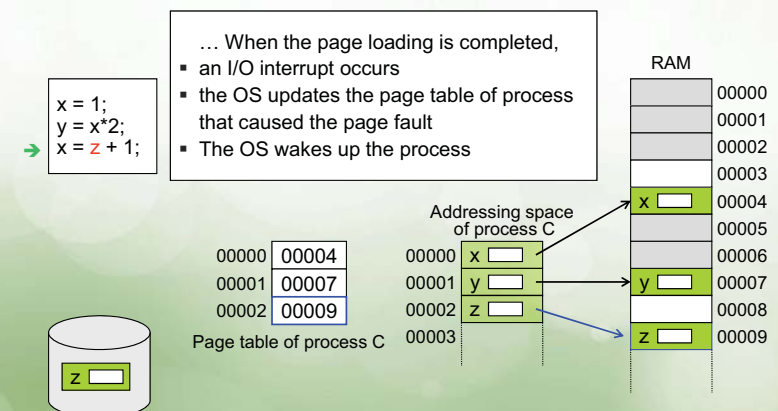
How? We will learn later in this chapter that the CPU detects page fault by checking the P bit in the page table entry of the page in address translation.

## How Virtual Memory Works (5)



## How Virtual Memory Works (6)

No suspended state?  
- When swap space used up



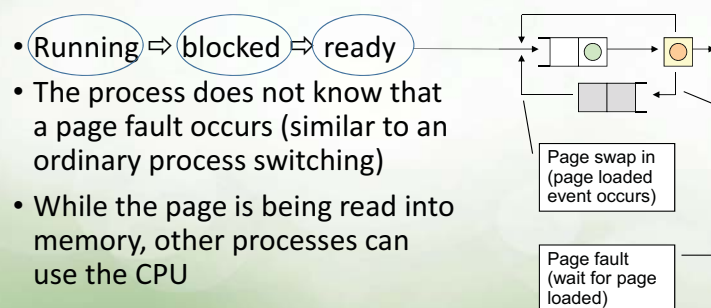
## Page Fault Handling, the Steps (1)

- A page fault occurs when a reference memory address is not present in RAM
- The CPU issues a trap, and the OS handles it:
  - block the process (while other processes can run)
  - find the page in the swap file in hard disk
  - find an unused frame in RAM
  - issue a read request to the hard disk

## Page Fault Handling, the Steps (2)

- When page loading is completed (disk I/O finishes)
  - an interrupt occurs
  - OS updates the page table of the blocked process that causes the page fault
  - OS wakes up the process

## Process State Changes in Page Fault Handling



## Efficiency of Virtual Memory

- Hard disk I/O is much slower than RAM
- If the OS swaps out a page just before using it, it has to get that page almost immediately. This is known as *thrashing*
- Thrashing is a state in which the system spends most of its time swapping pieces rather than executing instructions


## Efficiency of Virtual Memory (cont'd)

- To avoid thrashing, the OS tries to guess (based on recent history) which pieces are least likely to be used in the near future.
- In other words, can we just keep a few pages in RAM for a process without thrashing?

## Principle of Locality

- Program and data references within a process tend to cluster
  - Over a long period of time, the clusters in use change
  - Over a short period of time, the processor is primarily working with fixed clusters of memory references

## Locality in a Memory-Reference Pattern



- Note that during the lifetime of the process, references are confined to a subset of pages

## Principle of Locality (cont'd)

- **Temporal locality** – tendency to access data that are accessed recently
  - Pieces used frequently will likely be used in the near future
  - The OS can make intelligent guess of which page to swap out (page replacement algorithm)
- **Spatial locality** – tendency to access data near data that are recently accessed
  - Only a few pages of a process are needed within a short period of time



## Hardware Support to Virtual Memory

- Detail structure of the page table
  - How the CPU detects page fault
- How the CPU accesses the page table of the running process, the **page table pointer register**
- How the TLB caches most recently used entries of the running process

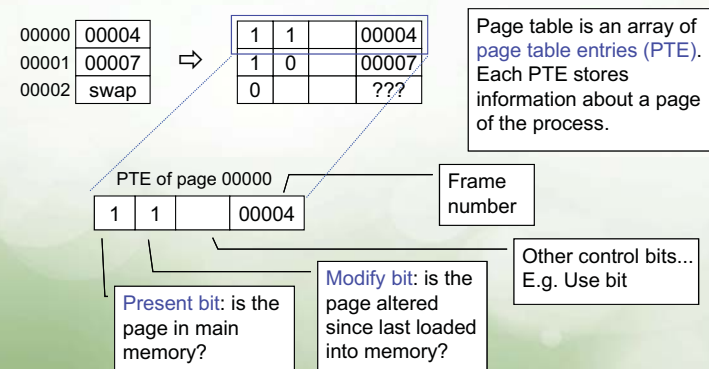
	P	M	
00000	1	1	00002
00001	1	0	00000
00002	0		???

Page table

TLB	
00001	00000
Page table ptr	
0123A000	

CPU

## Page Table Entry, PTE



## Page Table

- Each process has its own page table
  - What about threads within this process?
- Each page table entry contains
  - Frame number – where this page maps to
  - Present bit – the page is in RAM?
  - Modify bit – any byte modified in the page since last loaded?
  - Some other control bits...

## Present Bit

- “Is the page in RAM?”
- In address translation, the CPU first checks the Present bit of the page
  - If P=1, the page is present, and address translation continues. The process continues without blocking.
  - If P=0, a page fault occurs, and the OS is invoked to handle this

For detail, refer to “page fault handling” in the beginning of this chapter.

## Modify Bit

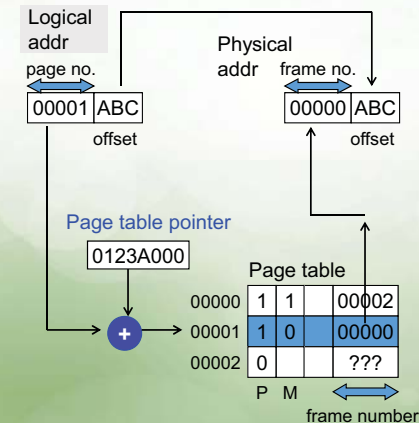
- “Any byte modified in the page since last loaded?”

- Example usage:
- If no change has been made since last loaded into memory, the page content is the same as the copy in the swap file (hard disk)
- Therefore, when it is swapped out, the OS does not need to write the page to the swap file

Other usage of the M bit is discussed in “page replacement algorithm” later in this chapter.

## Page Table Pointer

For virtual memory management  
Virtual address maps to logical address  
Logical address = Page number + Offset

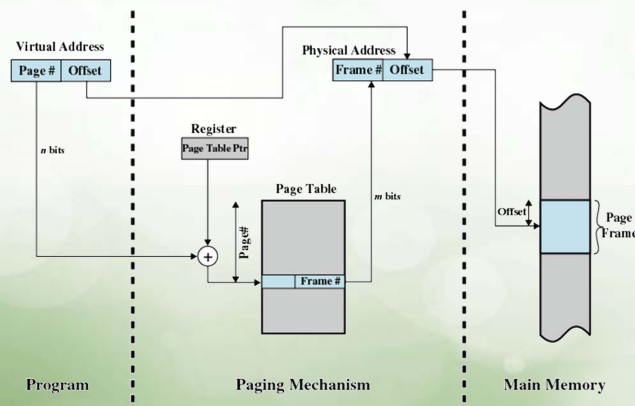


The **page table pointer** is a register that holds the starting address of the page table of the running process.

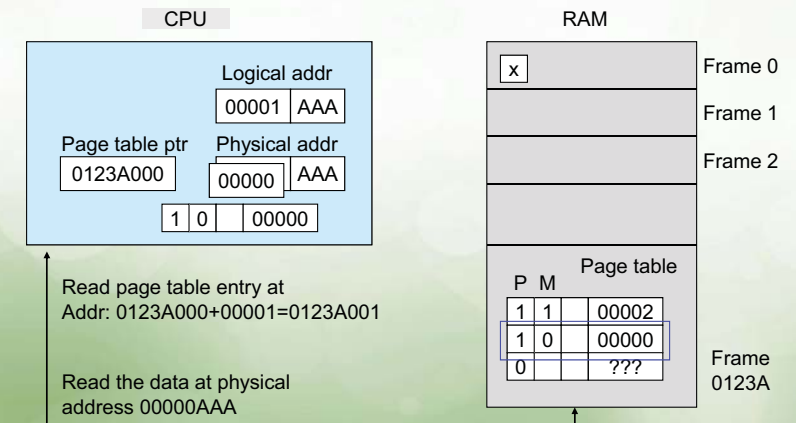
In address translation, CPU computes the address of the relevant PTE by adding the page number to the page table pointer.

What happens to page table pointer in process switching? In thread switching?

## Address Translation in a Paging System



## How CPU reads PTE

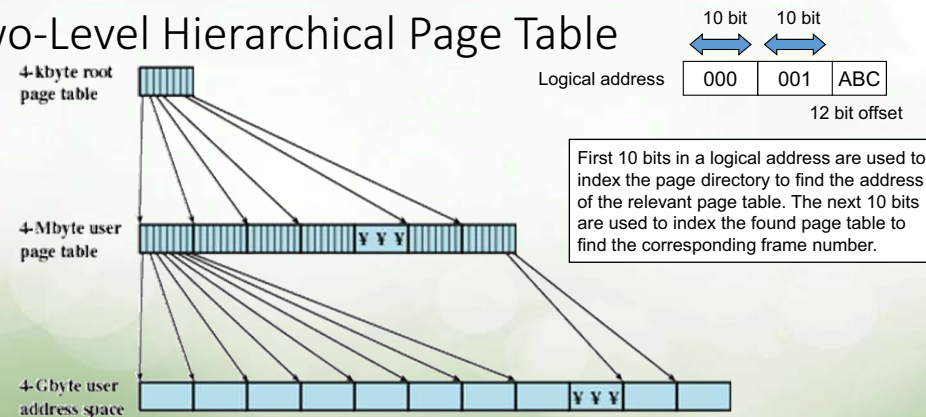




## Page Table can be Big...

- The entire page table may take up too much memory, e.g.,
  - A process in Pentium can have at most  $2^{20}$  pages
  - Each page table entry takes 32-bit (4-byte)
  - So a page table can occupy up to 4MB!
- Solutions:
  - Page table stored in virtual memory
  - Two-level page table

## Two-Level Hierarchical Page Table




## Virtual Memory Access can be Very Slow...

- Each virtual memory reference can cause two physical memory accesses
  - One to fetch the page table entry
  - One to fetch the data

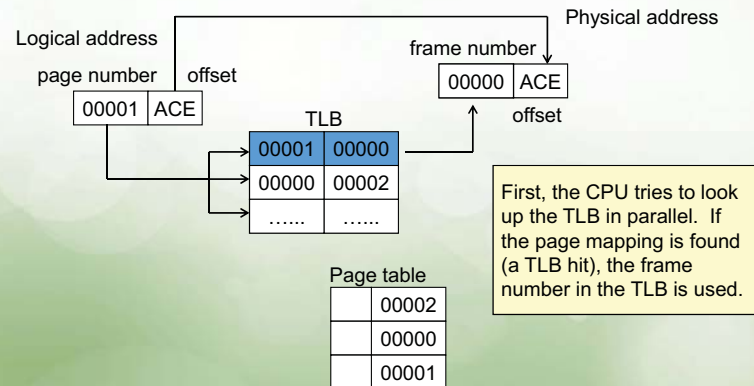
## Translation Lookaside Buffer

- **Translation Lookaside Buffer, TLB** is an associative cache that contains page table entries most recently used
- TLB is fast (cache inside the CPU)
- TLB is small (16\*4 entries in Pentium)
- Why such a small cache is enough?

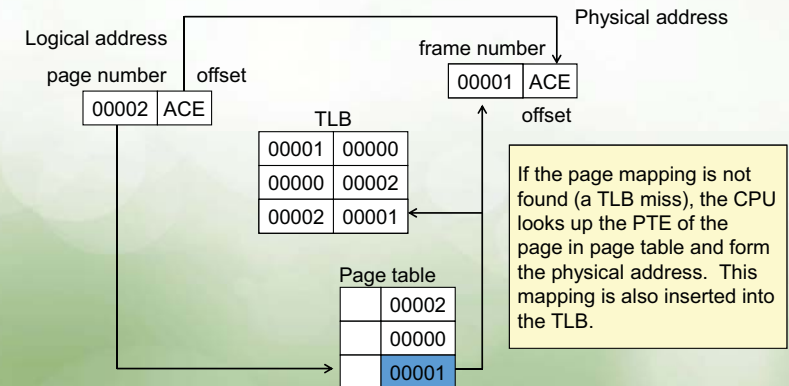


What happens to TLB in process switching? In thread switching?


## How TLB Works, TLB Hit



## How TLB Works, TLB Miss



## Use of a Translation Lookaside Buffer



## How TLB works (1)

- Given a logical address, processor first checks the TLB
- If the page table entry is found there (a hit)
  - the frame number is retrieved and the physical address is formed
  - no memory reference to the page table

## How TLB works (2)

- If the page table entry is not found (a miss)
  - CPU uses the page number to index the page table of the process to obtain the corresponding frame number
    - Page fault might occur and is handled by the OS
  - The TLB is updated to include the new page entry
  - the frame number is used to complete the physical address

## At Each Memory Access (1)

At each virtual memory access, the MMU (memory management unit) performs the following:

- split the logical address into **page number** and **offset**
- try to look up the TLB for the page number
- TLB hit: use the frame number cached to complete the physical address
- TLB miss: ...

1

- add the page number to the page table pointer register to obtain the address of the relevant page table entry
- try to read the PTE
- add the mapping to TLB
- check P bit

2

- present: use the frame number in PTE to complete physical address

3

- absent: page fault ...

## At Each Memory Access (2)

3

...  
- absent: page fault ...

### Page fault handling:

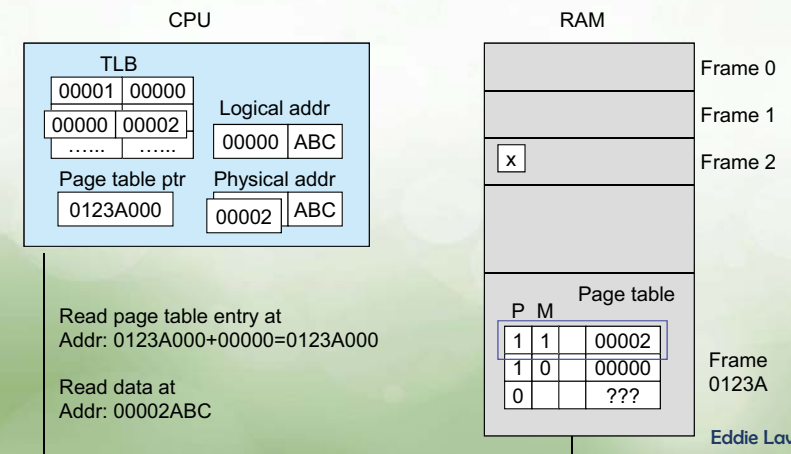
The CPU issues an interrupt/trap, and the OS handles it:

- block the process (other processes can run)
- find the page in the swap file in hard disk
- find an unused frame in RAM (w/ page replacement algorithm)
- issue a read request to the hard disk

### When the disk I/O finishes:

- an interrupt occurs
- OS updates the page table of the blocked, page fault process
- OS wakes up the process

## Example: Memory Access at TLB Miss (Page Present)



## Page Size

- Smaller page size means
  - Less amount of internal fragmentation
- But ... smaller page size also means
  - More pages required per process
  - More pages per process means larger page tables
  - Larger page tables means large portion of page tables in virtual memory
  - Secondary memory is designed to efficiently transfer large blocks of data so a large page size is better

## Page Size (cont'd)

- Small page size, large number of pages will be found in main memory
- As time goes on during execution, the pages in memory will all contain portions of the process near recent references
  - Page faults low
- Increased page size causes pages to contain locations further from any recent reference
  - Page faults rise

## Segmentation

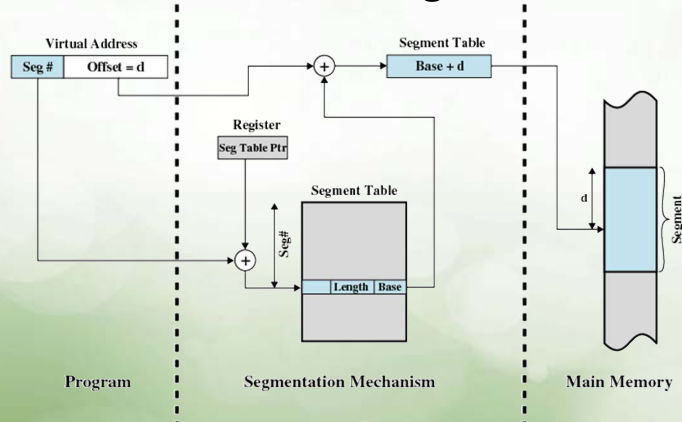
- May be unequal, dynamic size
- Simplifies handling of growing data structures
- Easier to sharing data among processes
- Easier to protect data

## Segment Table

- Each entry contains the length of the segment, P bit, M bit, and the base address of the segment

Segment table				
	P	M	length	base
0	1	0	40k	10101000
1	1	1	50k	20202000
2	1	1	20k	30303000

## Address Translation in a Segmentation System



Address Translation in a Segmentation System

Eddie Law 49

## Page Replacement Algorithms: Reference

Algorithm	Comment
Optimal	Not implementable, but useful as a benchmark
NRU (Not Recently Used)	Very crude
FIFO (First-In, First-Out)	Might throw out important pages
Second chance	Big improvement over FIFO
Clock	Realistic
LRU (Least Recently Used)	Excellent, but difficult to implement exactly
NFU (Not Frequently Used)	Fairly crude approximation to LRU
Aging	Efficient algorithm that approximates LRU well
Working set	Somewhat expensive to implement
WSClock	Good efficient algorithm

Eddie Law 50

## Next Topic

- I/O Management and Disk Scheduling

Eddie Law 51