

Informed Search

Searches

Uninformed searches

- Easy to implement
- Very inefficient in many situations
 - Huge search tree, time & space complexities

Informed searches (Heuristic)

- Use problem-specific information
- Reduce size of search tree
 - Resolve time & space complexities


Both give similar results

- Completeness & Optimality

Informed Search

Best-First Search

- Arrange nodes in the queue
- In the order of their goodness




Use evaluation function, $f(n)$

- Calculate goodness towards the goal

The order of expanding nodes is essential

- Affect size of the search tree
- Smaller size → less space, faster



Best-First Search

Applying evaluation function $f(n)$

- Every node is with a value stating its goodness

Based on values of $f(n)$

- Nodes in the queue are arranged
 - The best one is placed/expanded firstly

Do not guarantee

- The node to expand is really the best

The node only appears to be the best

- $f(n)$ is not omniscient

Evaluation Function $f(n)$

Path cost $g(n)$ is an example

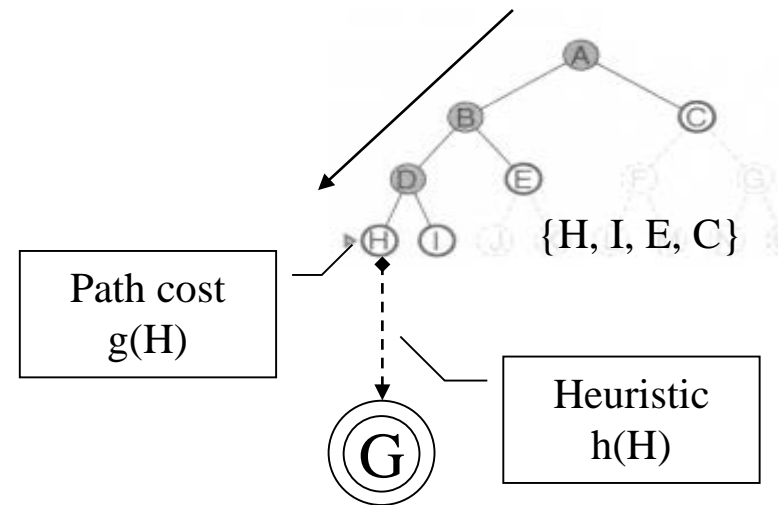
- $f(n) = g(n) \rightarrow$ uniform-cost search
- Do not direct search towards the goal

Heuristic function $h(n)$ is required


- Estimated cost of the cheapest path
 - From node n to a goal state
- Expand node closest to the goal
 - Node with least cost to goal
- If n is a goal state, $h(n) = 0$

$f(n) = h(n) \rightarrow$ greedy search

- Only try to expand the node closest to the goal
 - Likely lead to solution quickly



Example of $h(n) - h_{\text{SLD}}$



Values of h_{SLD} —straight-line distances to Bucharest.

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

h_{SLD} - Straight Line Distance


Cannot be computed from problem itself

Only obtainable from experience

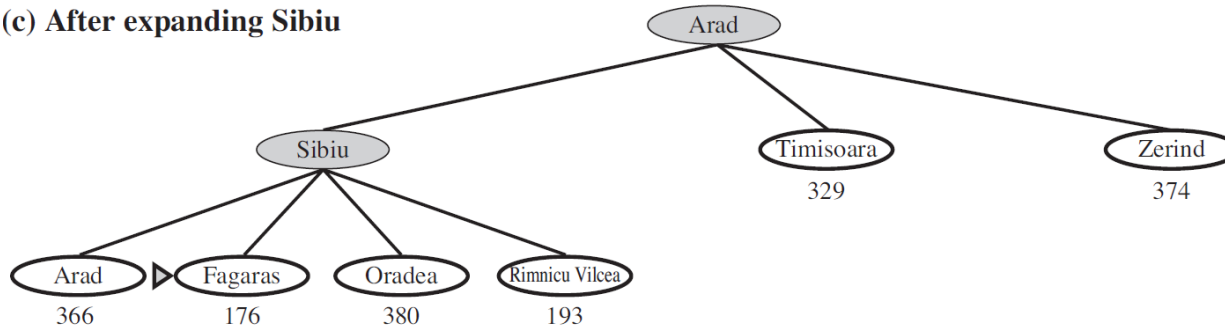
(a) The initial state



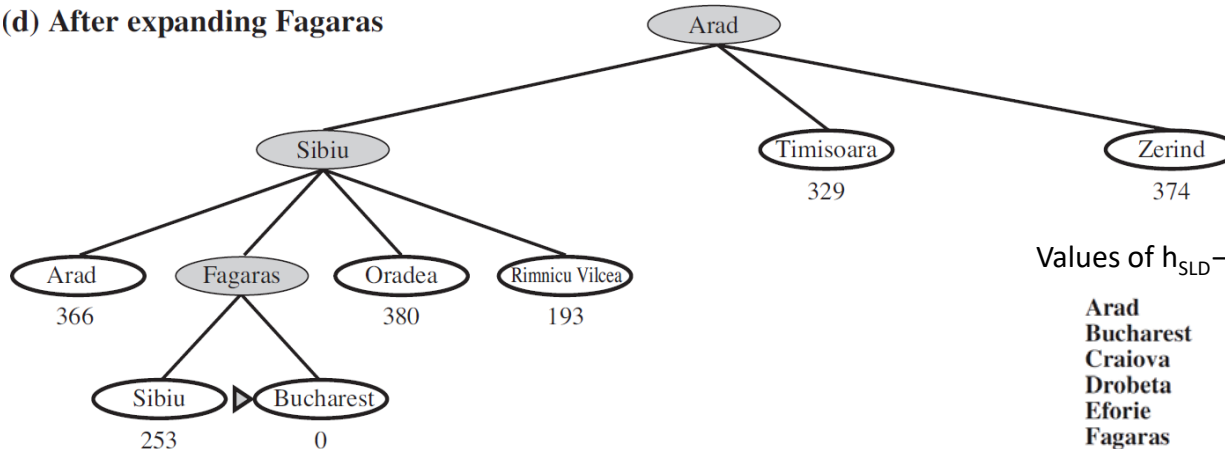
(b) After expanding Arad



(c) After expanding Sibiu



(d) After expanding Fagaras



Values of h_{SLD} —straight-line distances to Bucharest.

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Greedy Best-First Search

$$f(n) = h(n)$$

Good ideally

- Poor practically
- Cannot make sure
 - Heuristic $h(n)$ give correct estimation or not

Dangerous

- Just depending on the estimates $h(n)$

Analysis of Greedy Search

Similar to depth-first search

- Not optimal
 - Find the closest goal
- Incomplete
 - Repeated states may happen
 - Solution may never be found
- Time and space complexities
 - Depend on quality of heuristic function h

POOR

A* Search


A* Search

Most well-known best-first search

- Evaluate nodes by combining
 - Path cost $g(n)$ and heuristic $h(n)$
 - $f(n) = g(n) + h(n)$
 - $g(n)$ – cheapest known path
 - $h(n)$ – cheapest estimated path

Minimize total path cost

- Combine
 - Uniform-cost search
 - Greedy search




A* Search

Uniform-cost search


- Optimal and complete
- Minimize cost of the path so far, $g(n)$
- Can be very inefficient

Greedy search + Uniform-cost search


- Evaluation function $f(n) = g(n) + h(n)$
- [Evaluated so far + Estimated future]
- $f(n)$ = Estimated cost of the cheapest solution through n




(a) The initial state



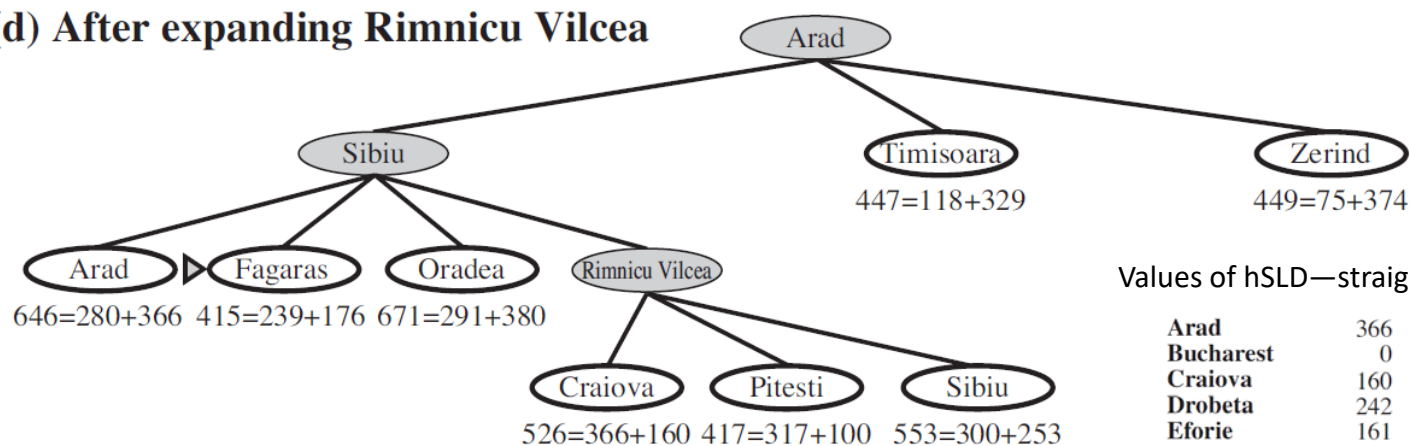
(b) After expanding Arad



(c) After expanding Sibiu



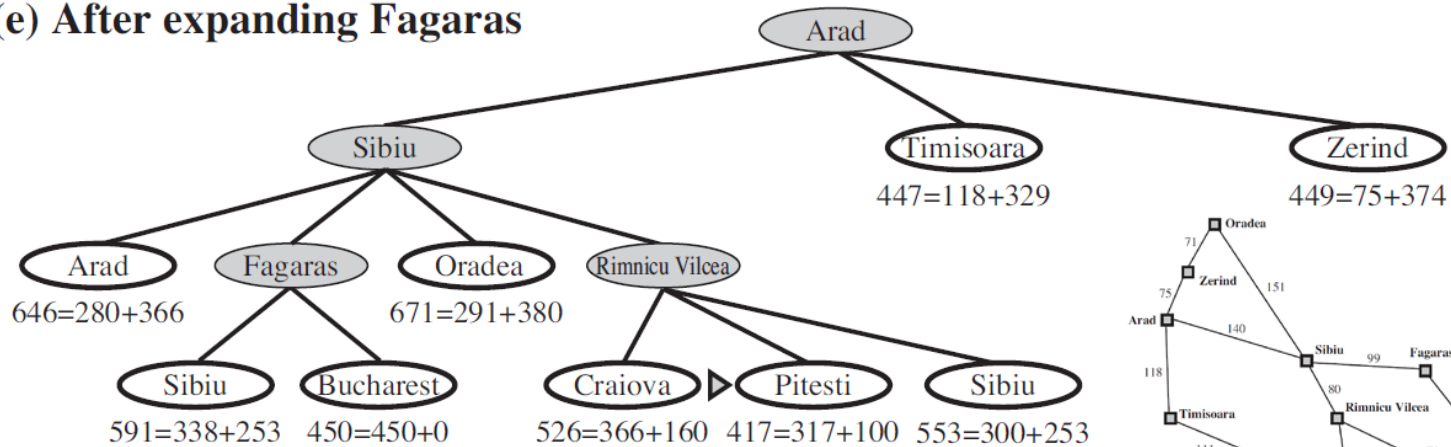
(d) After expanding Rimnicu Vilcea



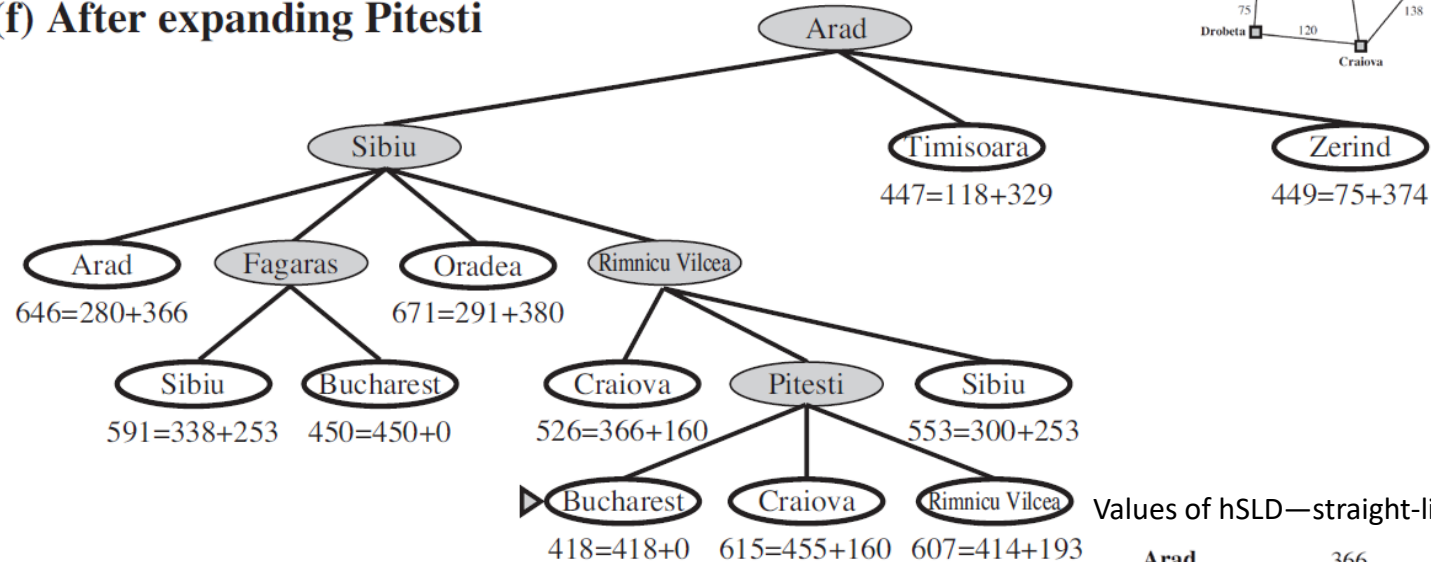
Values of hSLD—straight-line distances to Bucharest.

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

(e) After expanding Fagaras



(f) After expanding Pitesti



Values of hSLD—straight-line distances to Bucharest.

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Analysis of A* Search

Complete and optimal


- Uniform-cost search

Reasonable time & space complexities

- Greedy best-first search

Optimality can be assured

- Only $h(n)$ is admissible
 - Never overestimates cost from n to the goal
 - Can underestimate
- h_{SLD} , never overestimate
 - Straight Line Distance = Shortest distance between 2 places



Memory Bounded Search

Memory

- Another issue besides time constraint
- More important than time
 - Solution cannot be found
 - Not enough memory
- Solution can still be found
 - Even need a long time

Iterative Deepening A* Search

Iterative deepening (ID) + A*



- Similar to ID
- Reduces memory constraints effectively

Complete and optimal

- Similar to A*

Use $f = g + h$ for cutoff

- Instead of depth d or cost p
- Smallest f -cost of any nodes
 - Exceeded cutoff value in previous iteration



Heuristic Function

Heuristic Functions

Problem of 8-puzzle


- Two heuristic functions
 - Cut down search tree
- h_1 = number of misplaced tiles
- h_1 is admissible, never overestimates
 - At least h_1 steps to reach the goal

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State



Heuristic Functions

- h_2 = sum of distances of tiles from their goal positions
- City block distance or Manhattan distance
 - Count horizontally and vertically
- h_2 is admissible
 - $h_2 = 3 + 1 + 2 + 2 + 2 + 3 + 3 + 2 = 18$
 - Actual cost = 26
- Choose h_1 or h_2
 - Quality of heuristic

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8


Goal State

Branching Factor

Represent quality of a heuristic

Assume

- N = Total number of nodes expanded by A^*
- d = Solution depth
- b^* = Average branching factor of the tree
 - $N = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$
- If b^* tends to 1
 - N tends to be smallest
 - Tree becomes minimized



Branching Factor

h_1 = number of misplaced tiles

h_2 = sum of distances of tiles from their goal positions

	Search Cost (nodes generated)			Effective Branching Factor		
d	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035	227	73	2.78	1.42	1.24
14	—	539	113	—	1.44	1.23
16	—	1301	211	—	1.45	1.25
18	Tree size N	3056	363	—	1.46	1.26
20	—	7276	676	—	1.47	1.27
22	—	18094	1219	—	1.48	1.28
24	—	39135	1641	—	1.48	1.26

Branching Factor

h_2 is better than h_1

- h_1 = number of misplaced tiles
- h_2 = sum of distances of tiles from their goal positions

h_2 dominates h_1

- Any node n
- $h_2(n) \geq h_1(n)$
- Do not overestimate

7	2	4
5		6
8	3	1

Conclusion

- Better use heuristic function with higher values
- As long as it does not overestimate

Admissible Heuristic Functions

Relaxed problem

- Problem with less restriction on operators

Cost of exact solution to relaxed problem

- Often good heuristic for original problem

h_2 = sum of distances of tiles
from their goal positions

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Admissible Heuristic Functions

Original problem

- Tile can move from square A to square B
 - A is horizontally or vertically adjacent to B
 - B is blank

7	2	4
5		6
8	3	1

Relaxed problem

1. Tile can move from square A to square B
 - A is horizontally or vertically adjacent to B
2. Tile can move from square A to square B
 - B is blank
3. Tile can move from square A to square B

Admissible Heuristic Functions

Given 3 heuristics

- Choose the easiest one
- Do not waste too much time on $h(n)$


Do not know best heuristic

- Among h_1, \dots, h_m heuristics
- Set $h(n) = \max(h_1(n), \dots, h_m(n))$
- Let computer run it
 - Determine at run time


Admissible Heuristic Functions

Can also be derived

- From solution cost of a subproblem of given problem
- Get only 4 tiles into their positions
- Cost of optimal solution of this subproblem
 - Used as a lower bound, as a heuristic value



Start State



Goal State

Local Search

Local Search

Previously, find solution by searching paths



Initial state \rightarrow goal state

Many problems

- Solution is not a path to goal
 - 8-queens problem
- Solution
 - Final configuration
 - Not the order they are added or modified

Other kinds of method

- Local search




Local Search

A group of searches work in search space

Landscape has two axes

- Location (defined by states, a vector)
- Elevation (defined by objective function)



Local Search

Operate on a single current state


- Instead of a tree of multiple paths

Only move to

- Neighbors of current state

Paths in the search

- Not retained
 - Only retain current state
- Method is not systematic




Local Search

Two advantages

- Little memory – constant amount
 - Current state and some information
- Find reasonable solutions
 - Large or infinite (continuous) state spaces
 - Systematic algorithms are unsuitable

Suitable

- Optimization problems
- Find the best state with
 - Objective function



Analysis of Local Search

Completeness

- Complete local search algorithm
 - Find a goal if exists
- Most local search methods
 - Incomplete

Optimality

- Optimal algorithm
 - Find a global maximum or minimum (optimum)
- No local search methods are optimal

Theoretically, POOR

Practically, work well

- At least reasonable

Hill-climbing Search

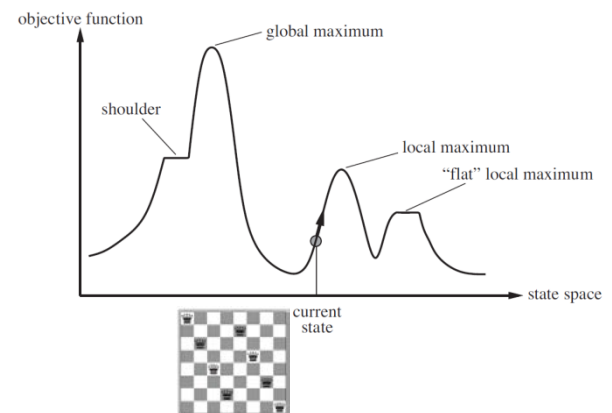
One type of local search

- Loop
 - Current state \rightarrow multiple successors (neighbors)
 - Evaluate each neighbors \rightarrow evaluation value
- Continually move in direction of increasing value
 - Current state \leftarrow best neighbor
 - Uphill (Climbing Hill)

No search tree is maintained

Node

- State
- Its evaluation
 - Objective value, real number



Hill-climbing Search

Multiple best neighbors

- Select among them randomly

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum


current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

neighbor \leftarrow a highest-valued successor of *current*

if neighbor.VALUE \leq current.VALUE **then return** current.STATE


current \leftarrow neighbor



Drawbacks of Hill-climbing Search

Hill-climbing is also called

- Greedy local search
- Grab a good neighbor state
 - Without thinking about where to go next





Local maxima

- Peak lower than the highest peak in state space
- Algorithm stops
 - Even though solution is far from satisfactory

Drawbacks of Hill-climbing Search

Ridges


- Grid of states is overlapped on a ridge
- Rise from left to right, right to left
- Unless move directly along top of ridge
- Search may oscillate from side to side
- Make little progress



Drawbacks of Hill-climbing Search

Plateaux

- Area of the state space landscape
- Evaluation function is flat
- Shoulder
- Impossible to make progress
- May be unable to find way off the plateau



function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

neighbor \leftarrow a highest-valued successor of *current*

if *neighbor*.VALUE \leq *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*



Solution

Previous problems happen

- Poor initial states

Random-restart hill-climbing (RRHC)

- Conduct a series of hill-climbing searches
 - Random generated initial states
 - Save best result found
- Fixed number of RRHC
 - Continue until best saved result has not been improved
 - For a certain number of iterations
- Cannot ensure optimality
 - Reasonably good solution




Simulated Annealing

Annealing

- Cooling down liquid until it freezes

Search

- Allow downhill steps
- Leave the local maximum
- Instead of start again randomly




Simulated Annealing

Every iteration

- Do not choose the best move
- Choose a random one

If the move results better

- Always execute
- Otherwise, take the move with a probability less than 1



Simulated Annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

inputs: *problem*, a problem

schedule, a mapping from time to “temperature”

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

for $t = 1$ **to** ∞ **do**

$T \leftarrow \text{schedule}(t)$



if $T = 0$ **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow \text{next.VALUE} - \text{current.VALUE}$

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$



T , temperature also affects the probability

- Since $\Delta E \leq 0$, $T > 0$, $\Delta E/T \leq 0$
- Probability $0 < e^{\Delta E/T} \leq 1$

Probability decreases exponentially


- “Badness” of the move = ΔE

Simulated Annealing

Higher T

- More likely bad move is allowed
- Usually happen in early schedule
- When T is large and $|\Delta E|$ is small (≤ 0)
 - $\Delta E/T$ is a negative small value $\rightarrow e^{\Delta E/T}$ is close to 1

$$0 < e^{\Delta E/T} \leq 1$$



T becomes smaller and smaller



- Until $T = 0$
 - Normal hill-climbing

Local Beam Search

No good in keeping only one current state

Local beam search

- k current states simultaneously
- All k states are generated randomly initially
- Generate successors
 - All successors of k states are generated
 - $m*k$ successors together
- Halt if any one is a goal
- Select k best successors
 - From complete successor list ($m*k$) and repeat




Local Beam Search

Different from random-restart hill-climbing

- Do not make k independent searches


Work together

- Collaboration
- Choose the best successors
 - Among those generated together by k states



Stochastic beam search

- Choose k successors at random
- Instead of k best successors




Genetic Algorithms

Variant of beam search

Successor states are generated by

- Combining two parent states
- Instead of modifying a single state


Successor state \rightarrow offspring



Genetic Algorithms

GA firstly make a population


- A set of k randomly generated states
- Individuals



Genetic Algorithms

Each state or individual

- Represent as a string over a finite alphabet
 - Binary or 1 to 8, etc.



State Representation

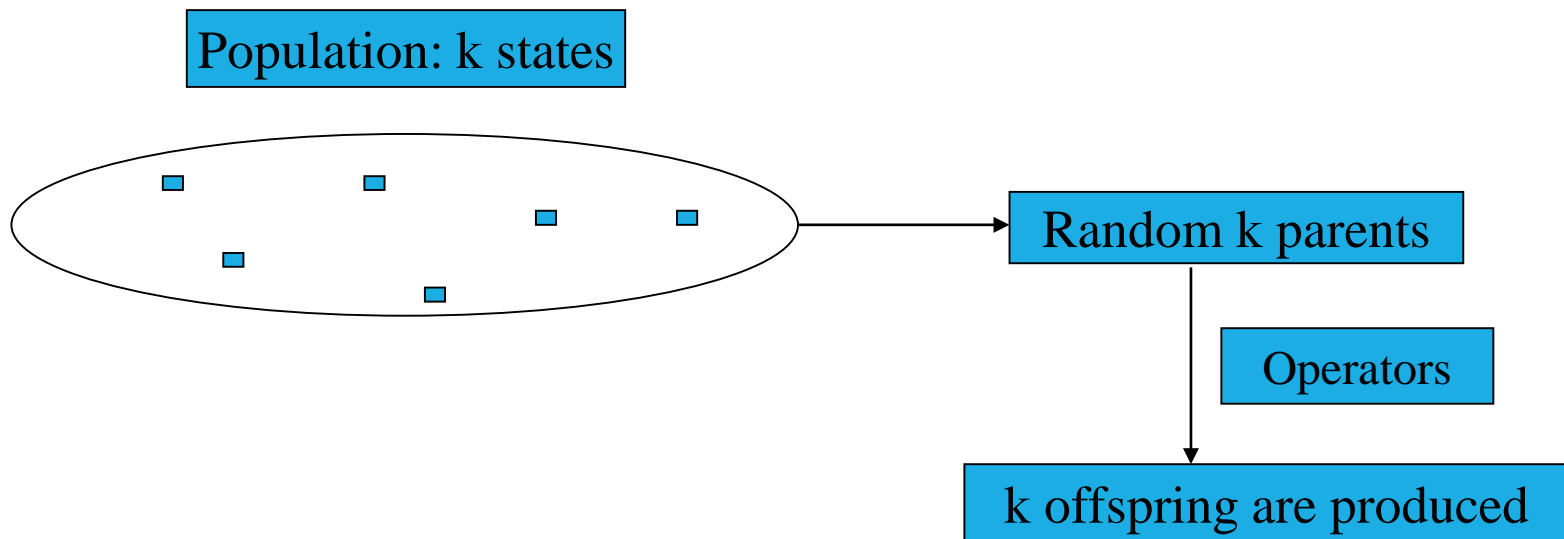
24748552

Genetic Algorithms

Choose k parents randomly

- Parents may occur more than once

Produce next generation of k individuals




Genetic Algorithms

Population (k) + Next gen (k) = 2k

Choose k best individuals as the next population

- Based on some probabilities over the fitness values



Genetic Algorithms

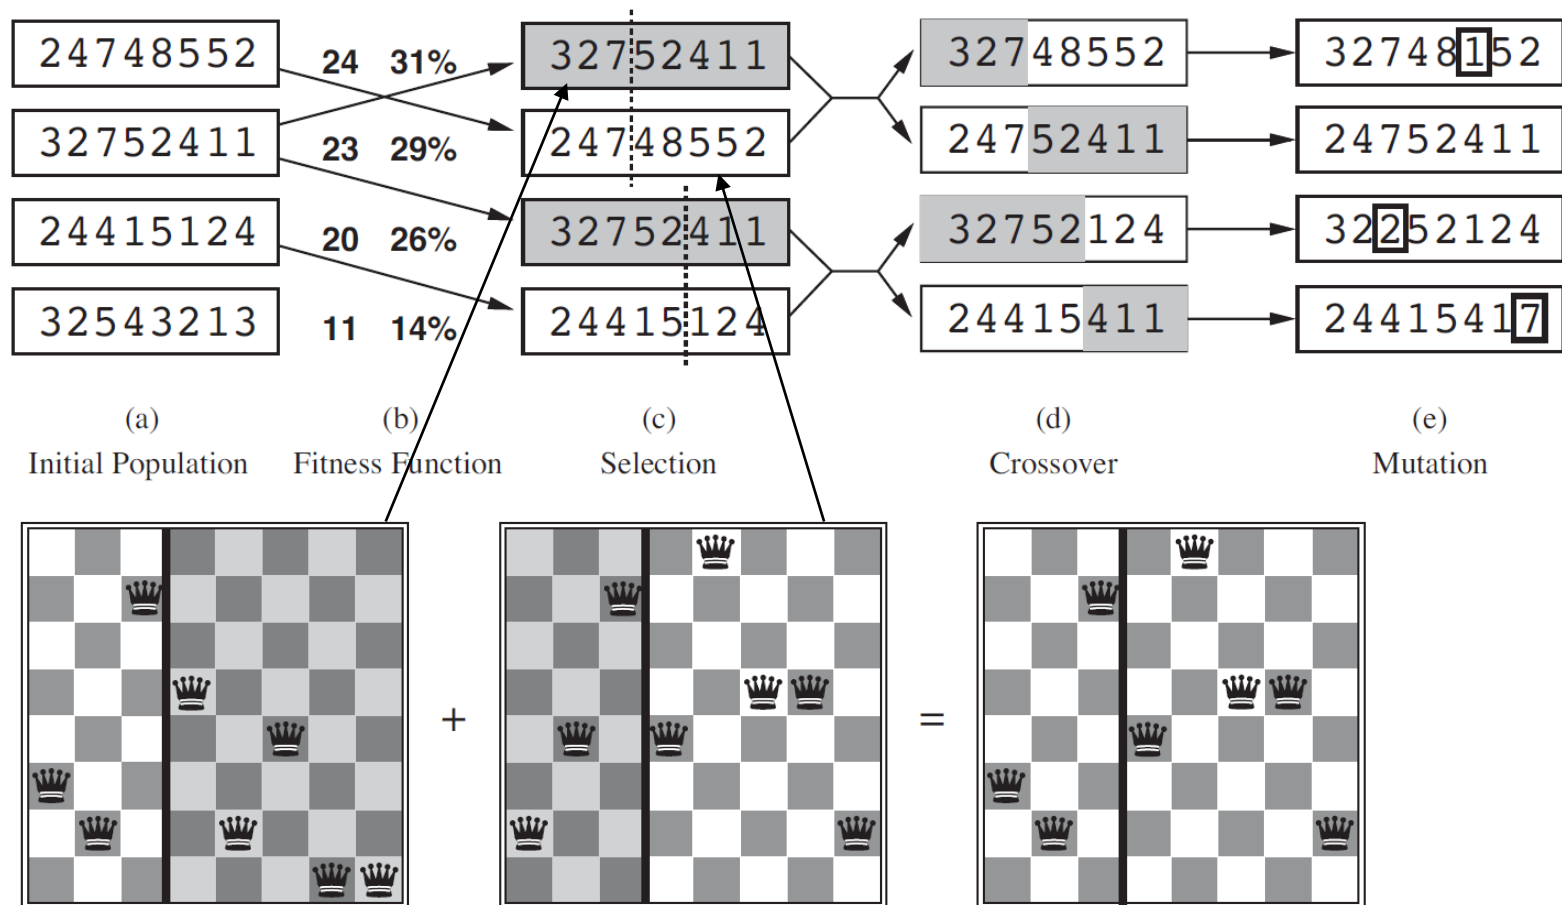
Operations for reproduction

- Crossover
 - Combine two parent states randomly
 - Crossover point is randomly chosen from positions in string
- Mutation
 - Modify the state randomly
 - A small independent probability

Efficiency and effectiveness

- State representation
- Algorithms for operations

Genetic Algorithms



Genetic Algorithms

function GENETIC-ALGORITHM(*population*, FITNESS-FN) **returns** an individual

inputs: *population*, a set of individuals

FITNESS-FN, a function that measures the fitness of an individual

repeat

new_population \leftarrow empty set

for $i = 1$ **to** SIZE(*population*) **do**

x \leftarrow RANDOM-SELECTION(*population*, FITNESS-FN)

y \leftarrow RANDOM-SELECTION(*population*, FITNESS-FN)

child \leftarrow REPRODUCE(*x*, *y*)

if (small random probability) **then** *child* \leftarrow MUTATE(*child*)

add *child* to *new_population*

population \leftarrow *new_population*

until some individual is fit enough, or enough time has elapsed

return the best individual in *population*, according to FITNESS-FN

function REPRODUCE(*x*, *y*) **returns** an individual

inputs: *x*, *y*, parent individuals

$n \leftarrow$ LENGTH(*x*); $c \leftarrow$ random number from 1 to n

return APPEND(SUBSTRING(*x*, 1, c), SUBSTRING(*y*, $c + 1$, n))