

Machine Learning

Machine Learning

Machine? Learning?

- A computer. Improve through experience

Build computer programs

- Improve itself at some task
- Through experience

Experience

- Through training examples

Machine Learning

Subfield of computer science

Give computers ability to learn

- Without being explicitly programmed

Study and construction of algorithms

- Learn from and make predictions on data

Employed in a range of computing tasks

- Programming explicit algorithms is infeasible
- Spam filtering, detection of network intruders, optical character recognition (OCR) and search engines

Relations to Other Subjects

Artificial Intelligence

- Provide a way for implementing A.I.
- A.I. – methods for human-level cognitive tasks
 - Most cognitive tasks – classification / prediction

Statistics

- Learning with mathematics
- Mostly aimed for prediction tasks
- Most ML methods are from statistics
 - Training data / experience E

Relations to Other Subjects

Pattern Recognition

- One of applications of A.I.
 - An application of Machine Learning Methods
 - For signals, graphics, and multimedia
- Mostly for classification tasks

Data Mining

- Machine learning with very large or distributed data sets
 - Traditional methods only work for limited-size data set
 - New efficient algorithms are necessary

Classification of Machine Learning Tasks

Classification of Machine Learning Tasks

Supervised learning

- Present computer with examples
 - Inputs and their desired outputs
- Given by a “teacher”
- Learn a general rule that maps inputs to outputs

Unsupervised learning

- No labels are given
- Find structure of input its own
- Discover hidden patterns in data

Classification of Machine Learning Tasks

Semi-supervised learning

- Between supervised and unsupervised learning
- Teacher gives an incomplete training signal
 - Training set with some or many missing target outputs

Reinforcement learning

- Computer program interacts with a dynamic environment
- Perform a certain goal
 - Drive vehicle or play game against an opponent
- Feedback are provided
 - In terms of rewards and punishments
 - When it navigates its problem space

Supervised Learning

Infer a function from labeled training data

- Analyze training data
- Produce an inferred function
- Use for mapping new examples

Training data

- A set of training examples
- Each is a pair
 - Input object (typically a vector)
 - Desired output value (also called supervisory signal)

Optimal scenario

- Determine class labels for unseen instances correctly

Supervised Learning

Given data $\{(x_1, y_1), \dots, (x_n, y_n)\}$

Seek a function that explains relationship between

- Input attribute x
- Output attribute y
- $y = f(x) + \epsilon$

Algorithms

- Naive Bayes classifier
- Neural network
- Support vector machines
- Nearest Neighbor Algorithm

Unsupervised Learning

Infer a function

- Describe hidden structure from unlabeled data

Distinguishable from other learning schemes

- Examples given are unlabeled
- No error or reward signal
- Cannot evaluate a potential solution
- No objective evaluation
 - Accuracy of the structure output

Approaches


- Clustering

Supervised Learning

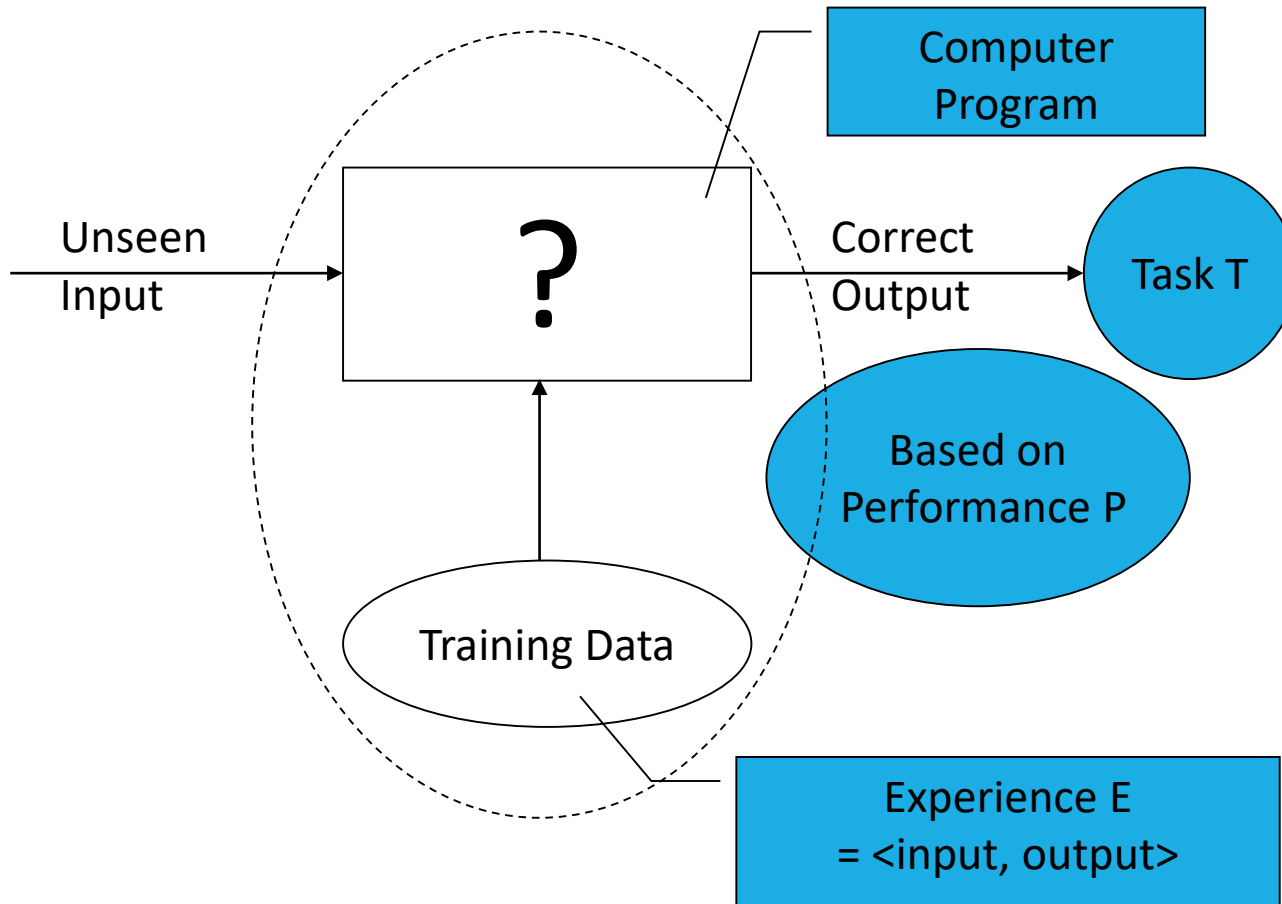
Supervised Learning

A computer program

- Learn from experience E
 - In some class of tasks T with performance P
- Its performance at tasks in T improves with experience E
 - Performance is measured by P



Supervised Learning



Tasks for Supervised Learning

Classification / Prediction

By making a target function

- Estimate a mathematical model
- Through a set of training examples

Only outputs are different

- Classification – Discrete values
 - True or false, Yes or no, Class A to class F
- Prediction – Continuous values

Examples of Learning Tasks

Classification

- Recognize spoken words
- Classify new astronomical structures

Prediction

- Stock trend
- Robot control


Summary of Supervised Learning

A job of estimating a function (program, ?)

- Take E (training data) as inputs
- Output as accurate as possible
 - Solve problems in task T

Learning job may be iterative

- Improve the function accuracy
 - Not just one time but many times



Designing a Supervised Learning System

Designing a Supervised Learning System

1. Exact type of knowledge to be learned, i.e., the '?'
 - A program? A decision tree? Logical rules? A function?
 - Usually a mathematical function
2. Representation for this target knowledge, or function
 - Format of $f(x)$? Sine wave? Exponential?
 - Usually $f(x) = wx + b$ or polynomial
3. A function approximation algorithm

Example Problem Domain

Chess playing program

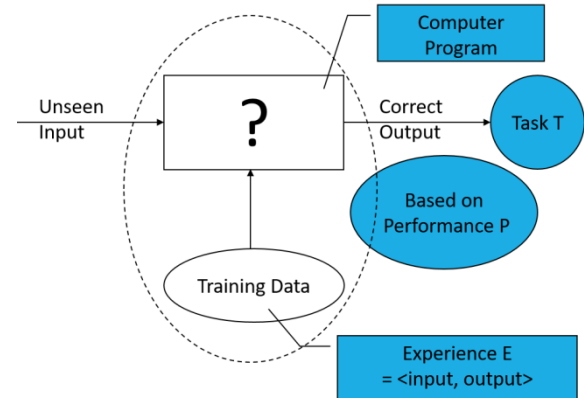
- Input: Any board state
- Output: Legal moves

Experience $E = \langle \text{input}, \text{output} \rangle$

- $E = (\text{State}, \text{Legal moves})$

Learn = Make the function ‘?’

- Under any state
 - Return the *best* move among legal moves
- *Best* is judged by performance P
 - Rules of the game




Example Problem Domain

Function '?' = *ChooseMove*

- *ChooseMove*: State \rightarrow Move
- How to compare resulting Move?

Easier alternative

- y : State \rightarrow R
 - A function evaluating any state into a real number
 - The higher the score, the better the state
- State S_0
 - Under legal moves M_1, M_2, M_3, \dots
 - Generates S_1, S_2, S_3, \dots
 - $y: S_1 \rightarrow R_1, y: S_2 \rightarrow R_2, \dots$
 - Choose the best state (i.e. the move)



Example Problem Domain

Learned function is changed

- From *symbolic* to *numerical*

Function y is usually not efficiently computable

- Hard to find it exactly

Our target

- Estimate an approximate function \hat{y} to replace y
 - Representation of \hat{y}

Example Problem Domain

Representations for description of \hat{y}

- A look up table
- A collection of logical rules
- A polynomial function
- A simple linear combination

$$\hat{y}(S) = w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$$

where $S = \langle x_1, x_2, \dots, x_n \rangle$

A Basic Function Approximation Algorithm

Data for Training Algorithm

A set of training examples

- Each is a pair
 - $\langle S, \text{Value}_S \rangle = \langle \text{input}, \text{output} \rangle$
 - $S = \text{State}$, $\text{Value}_S = \text{Value of State } S$
- E.g. $\langle (x_1=1, x_2=3, x_3=0, \dots), 50 \rangle$
 - Record in a database
 - Many records in the database

1	$(x_1=1, x_2=3, x_3=0, \dots)$	50
2	...	40
\vdots	\vdots	\vdots
m	...	35

Evaluating Function Estimation

Approximating \hat{y} using training data


- $E = (S, \text{Value}_S)$
- A score $\hat{y}(S)$ can be obtained for state S
 - Evaluation of S

$$\hat{y}(S) = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n$$

where $S = \langle x_1, x_2, \dots, x_n \rangle$

Error can be measured

- Between True and Estimated values
- $\text{Error}_S = (\text{Value}_S - \hat{y}(S))^2$



$$\text{Error} = \sum_{\substack{\langle S, \text{Value}_S \rangle \\ \in \text{Training Data}}} (\text{Value}_S - \hat{y}(S))^2$$

1	$(x_1=1, x_2=3, x_3=0, \dots)$	50
2	...	40
... m	...	35

Evaluating Function Estimation

Find \hat{y}

- Such that *error* can be minimized

Since $\hat{y} = \sum w_i x_i$, $i = 0$ to n

- x_i are constants (input data)
- Adjust $w_i \rightarrow$ Adjust \hat{y}
 - Minimize the error
- Algorithm is called LMS (Least Mean Square)
 - A very rational and easy training rule

$$\hat{y}(S) = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n$$

where $S = \langle x_1, x_2, \dots, x_n \rangle$

Adjust Weights

For each training example $\langle S, \text{Value}_S \rangle$

- Use current weights w_i
 - Calculate $\hat{y}(S)$
- For each weight w_i , update it as

$$\hat{y} = \sum w_i x_i$$

$$w_i \leftarrow w_i + \eta (\text{Value}_S - \hat{y}(S)) x_i$$

$$w_2 \leftarrow 1 + 0.05 (50 - 48) 3 = 1.3$$

A small constant controlling
learning rate, “eta”

Neural Network


Neural Network

A famous and powerful learning method

- Linear and nonlinear target function
- Single output
 - Real-valued, discrete-valued
- Multiple output
 - Vector-valued
- Other features
 - Robust (insensitive to noise)
 - Noise = some misleading / incorrect input values
 - Easy to implement
 - Fast and efficient
 - But hard to interpret

Neural Network Representation

Consist of at least three layers



Neural Network Representation

Input layer


- Accept values of training examples $\langle \underline{\mathbf{x}}, \mathbf{y} \rangle$
 - If \mathbf{x} is 7-tuple, then 7 input neurons

Output layer

- Similar to input layer
 - If \mathbf{y} is triple, then 3 output neurons

Hidden layer

- Handle unknowns
- Handle nonlinearity within data
- Number of hidden neurons
 - Usually $>$ input neurons




Neural Network Representation

Between any two layers

- Connected by some arcs
- Associated with a weight


Train neural network

- Update the weights in layers
- Fixed representation for target function



$$y = \sum_i w_{3i} \varphi_i \left\{ \sum_j w_{2ij} [\sigma_j (\sum_k w_{1jk} x_k)] \right\}$$

Example of Steering Control



Output Layer

A binary-vector

$[0 \ 0 \ 0 \ 1 \ 0 \ \dots \ 0]$

$$y_i = \sum_j w_{2ij} [\sigma_j(\sum_k w_{1jk} x_k)]$$

Input Layer

= 30 x 32

= 960 neurons

Appropriate Problem

Problems have

- Instances (data records)
 - Many attribute-value pairs
 - Input = matrix or vector
- Target function may be any data type
 - Discrete-valued, real-valued, or a vector of attributes
- Noisy training data

Properties of NN learning

Training takes long time

- From seconds to hours
- Or even days or weeks

Trained target function

- Execution takes very short time
 - Several seconds
- NN structure
 - Linear combination of weights and inputs
 - $y = wx + b = \mathbf{w}\mathbf{x}$

Support Vector Machines

Support Vector Machines

Abbreviated as SVM

Main application

- Classification
- Regression
 - Model/Function estimation
 - Approximation

Support Vector Machines

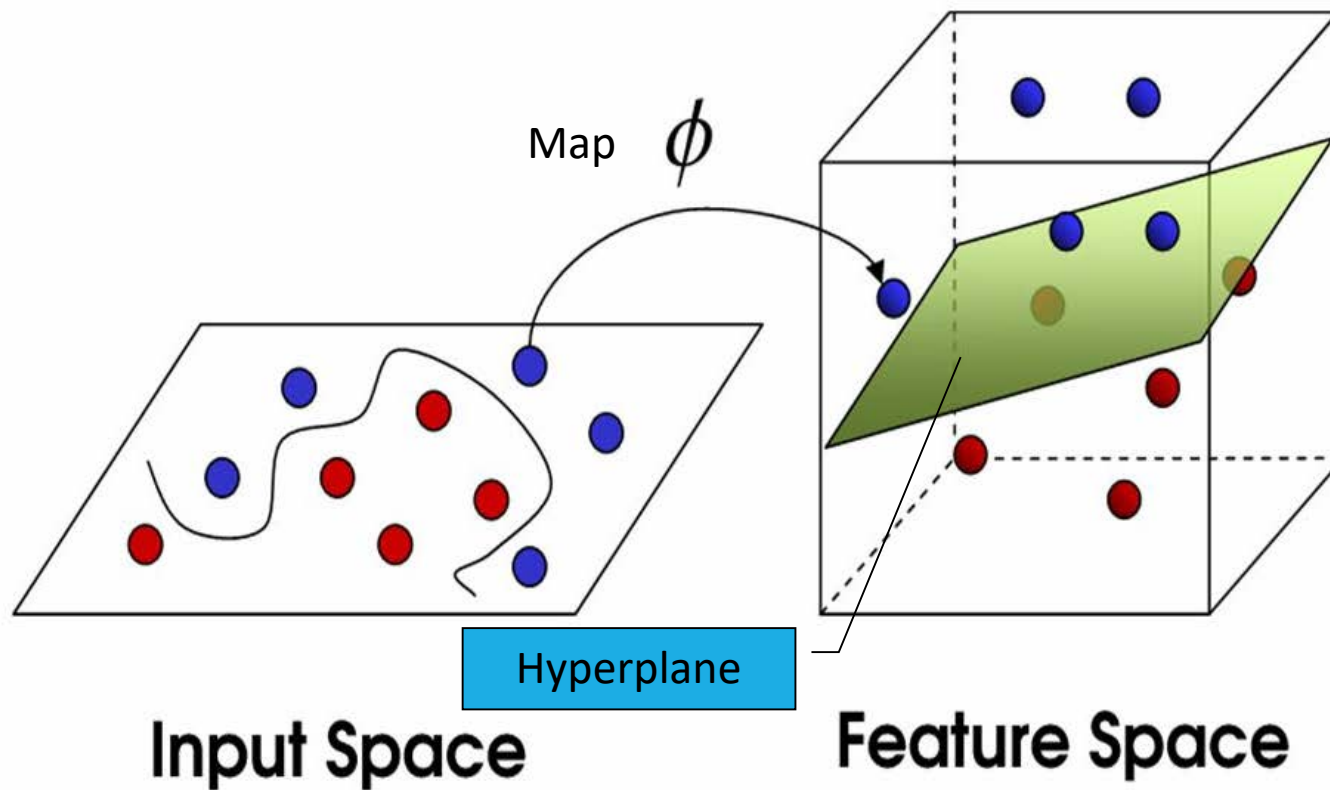
Can only build linear functions

- Goal: nonlinear functions
 - Done by some nonlinear transformation ϕ
 - Nonlinear input spaces \rightarrow High dimensional linear feature spaces
 - Build linear functions

Special property

- Required hidden units are automatically determined

Support Vector Machines



k -Nearest Neighbor Method

k -Nearest Neighbor


Nearest neighbors of an instance

- Defined on Euclidean distance

Euclidean distance

- An instance $\mathbf{x} = \langle x_1, x_2, \dots, x_n \rangle$
- Distance between two instances \mathbf{x}, \mathbf{y}

$$d(\mathbf{x}, \mathbf{y}) \equiv \sqrt{\sum_{r=1}^n (x_r - y_r)^2}$$




k -NN for Discrete Values (Classes)

Compare query \mathbf{x}

- With every training example $\langle \mathbf{y}, f(\mathbf{y}) \rangle$
- For $k = 1$
 - Return most similar object $\hat{\mathbf{y}}$
 - Based on Euclidean equation
 - Assign $f(\hat{\mathbf{y}})$ to $f(\mathbf{x})$
- For $k > 1$
 - Retrieve a set of k similar instances $\hat{\mathbf{y}}_i$

$$f(\mathbf{x}) \leftarrow \operatorname{argmax}_{v \in V} \sum_{i=1}^k \delta(v, f(\hat{\mathbf{y}}_i))$$

V = all possible classes {square, tri}



$$\delta(a, b) = 1 \text{ if } a = b$$

$$\delta(a, b) = 0 \text{ if } a \neq b$$

k -NN for Continuous Target Values


Replace function of estimating class

$$f(\mathbf{x}) \leftarrow \frac{1}{k} \sum_{i=1}^k f(\hat{\mathbf{y}}_i)$$

Continuous target values

- Return the mean of
 - Target values of k similar nearest neighbors

Distance-Weighted NN



Refinement to k NN

- k similar instances are weighted
 - Based on their distance to \mathbf{x}

Suggested weights for instance

$$w_i \equiv \frac{1}{d(\mathbf{x}, \mathbf{y}_i)^2}$$

If $d(\mathbf{x}, \mathbf{y}_i)$ is zero
 $f(\mathbf{x}) = f(\mathbf{y}_i)$

For discrete value: target class

$$f(\mathbf{x}) \leftarrow \operatorname{argmax}_{v \in V} \sum_{i=1}^k w_i \delta(v, f(\hat{\mathbf{y}}_i))$$

Distance-Weighted NN

For continuous value: target value

$$f(\mathbf{x}) \leftarrow \frac{\sum_{i=1}^k w_i f(\hat{\mathbf{y}}_i)}{\sum_{i=1}^k w_i}$$

$$w_i \equiv \frac{1}{d(\mathbf{x}, \mathbf{y}_i)^2}$$

If $w_i = 1$ for all i , this term = k
i.e. kNN

Remarks on k NN

Pros

- Distance-weighted k NN is highly effective
- Robust to noisy training data
 - Provided a large set of training examples

Cons

- Similarity metric depends on all attributes
 - Some attributes may be irrelevant = noise
 - Misleading / wrong
- Distance metric = Euclidean space
 - Do not guarantee that it can represent similarity

Unsupervised Learning – Clustering



Clustering

Divide a set of objects into groups

- Objects in same group are similar
- Objects in different groups are not similar

Input


- A set of points P
- A set of centers C (optional)




Clustering

- Assign every point of P to the nearest center in C

Clustering Problem



Clustering Problem



Applications of Clustering

Image Processing

- Cluster images
 - Based on visual content

Web

- Cluster groups of users
 - Based on access patterns on webpages
- Cluster webpages
 - Based on content

Bioinformatics

- Cluster similar proteins together
 - Similarity in chemical structure or functionality

Clustering Problem

Discrete vs continuous clustering

Discrete clustering

- Restrict centers of clustering to a subset of input

Continuous clustering

- Centers might be placed anywhere in given metric space

Clustering Algorithm

K-Centers clustering

K-Median clustering

K-Means clustering

Hierarchical clustering

K-Means Clustering

Minimize function:

$$E(\Gamma, V) = \sum_{i=1}^k \sum_{j=1}^n \gamma_{ij} \|\bar{x}_j - \bar{v}_i\|^2$$

Data points: $X = \{\bar{x}_1, \bar{x}_2, \dots, \bar{x}_n\}$

Clusters: C_1, C_2, \dots, C_k

Centers: $V = \{\bar{v}_1, \bar{v}_2, \dots, \bar{v}_k\}$

Partition matrix: $\Gamma = \{\gamma_{ij}\}$

$$\gamma_{ij} = \begin{cases} 1 & \text{if } \bar{x}_j \in C_i \\ 0 & \text{otherwise} \end{cases}$$

Iterative algorithm

- Initialize the centers V (the position of k centers)
 - By randomly picking points from X
- Assign each data point to the nearest center
 - Recalculate partition matrix and $E(\Gamma, V)$
- Adjust the position of each center
- Repeat above two steps until convergence

K-Means Clustering

Disadvantages

- Dependent on initialization

